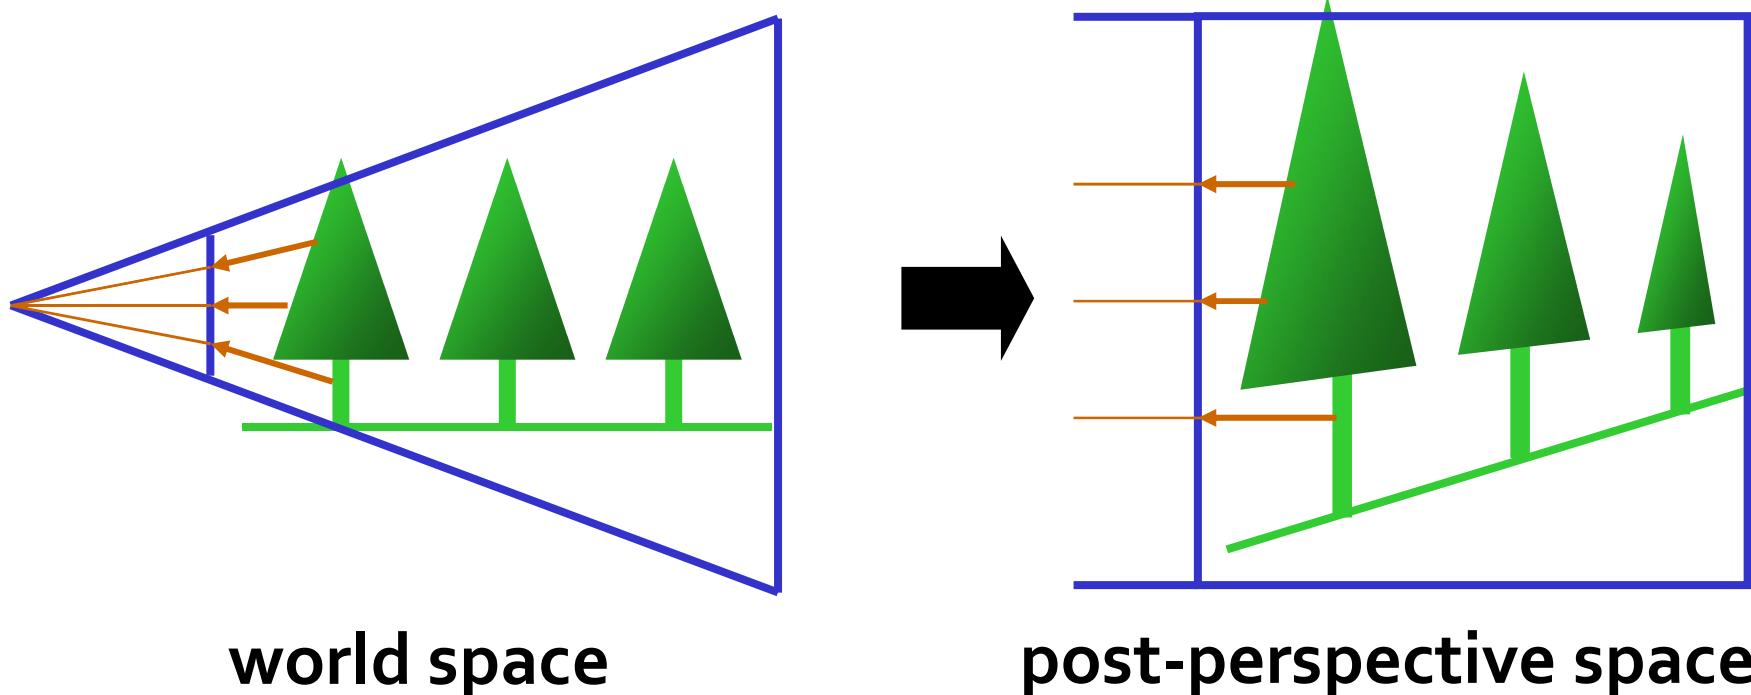


Perspective Shadow Maps

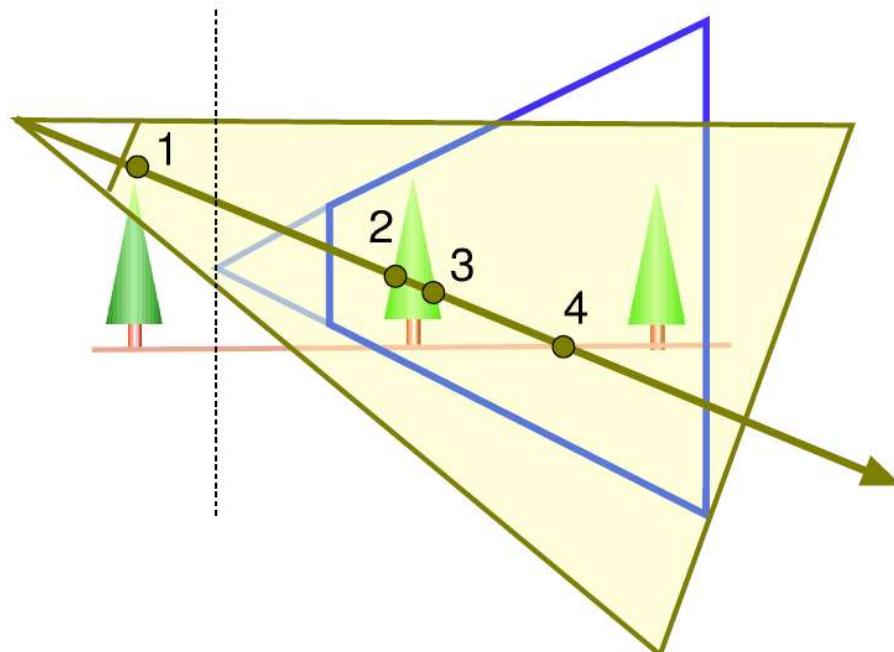
[Stamminger & Drettakis 2002]

- Want to account for perspective warping of eye
- Shadow map in **post-perspective** eye space
- Reduce perspective aliasing

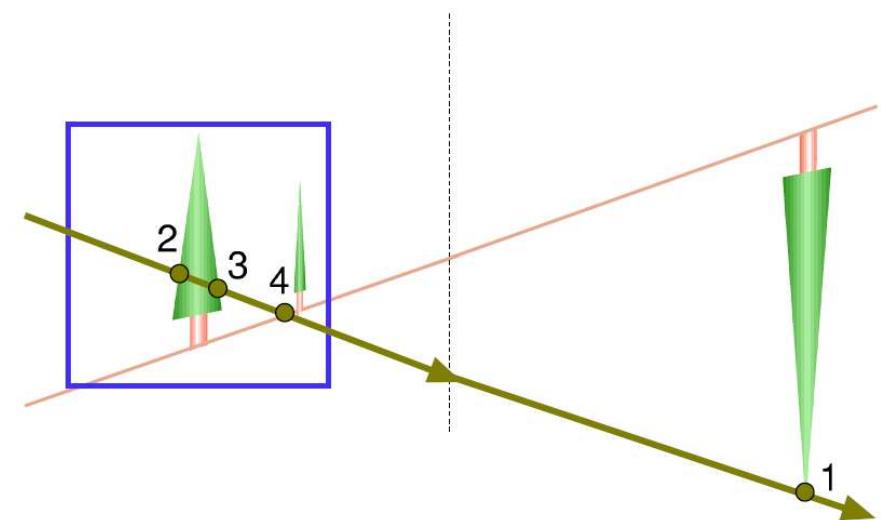


PSM Problems

- Shadows from behind



world space

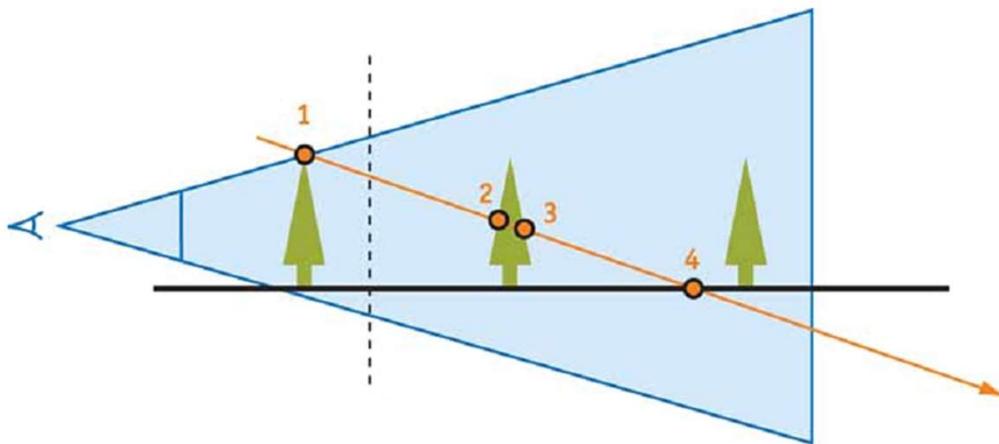


post-perspective space

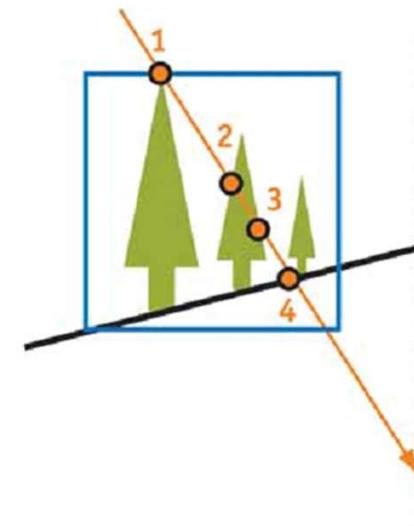
PSM Problems

■ Shadows from behind

- Require move back due to singularity
- Changes resulting quality of shadow

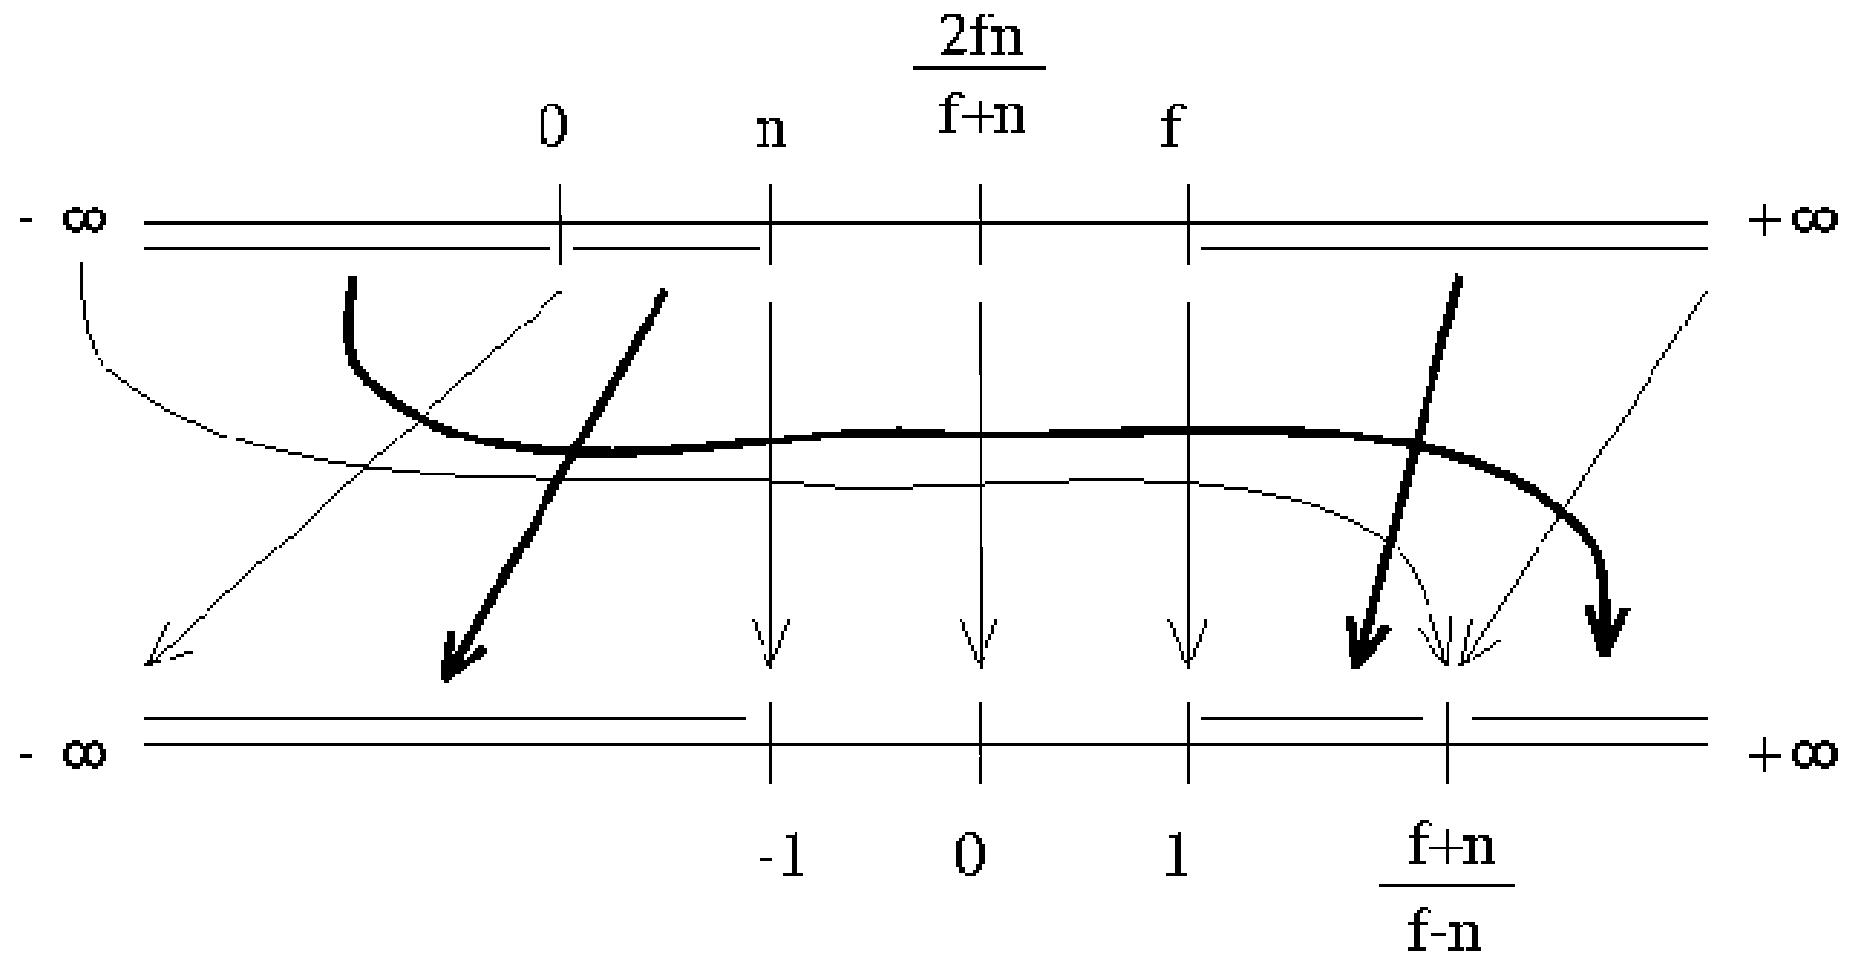


world space

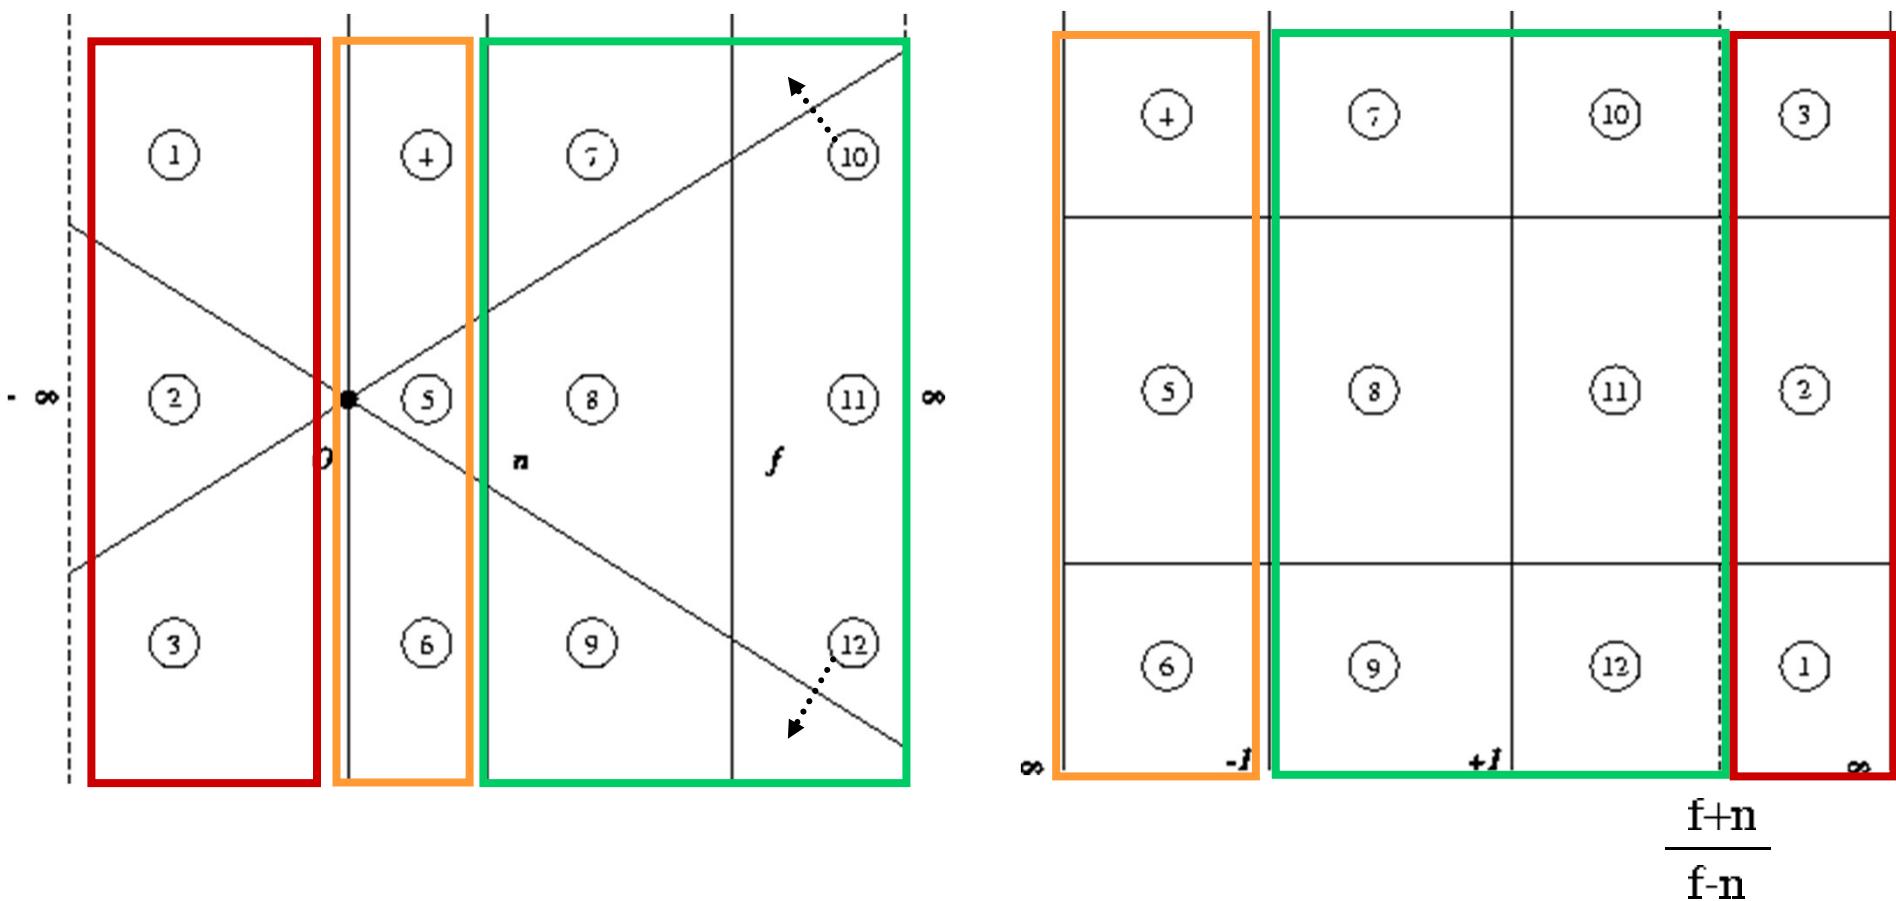


post-perspective space

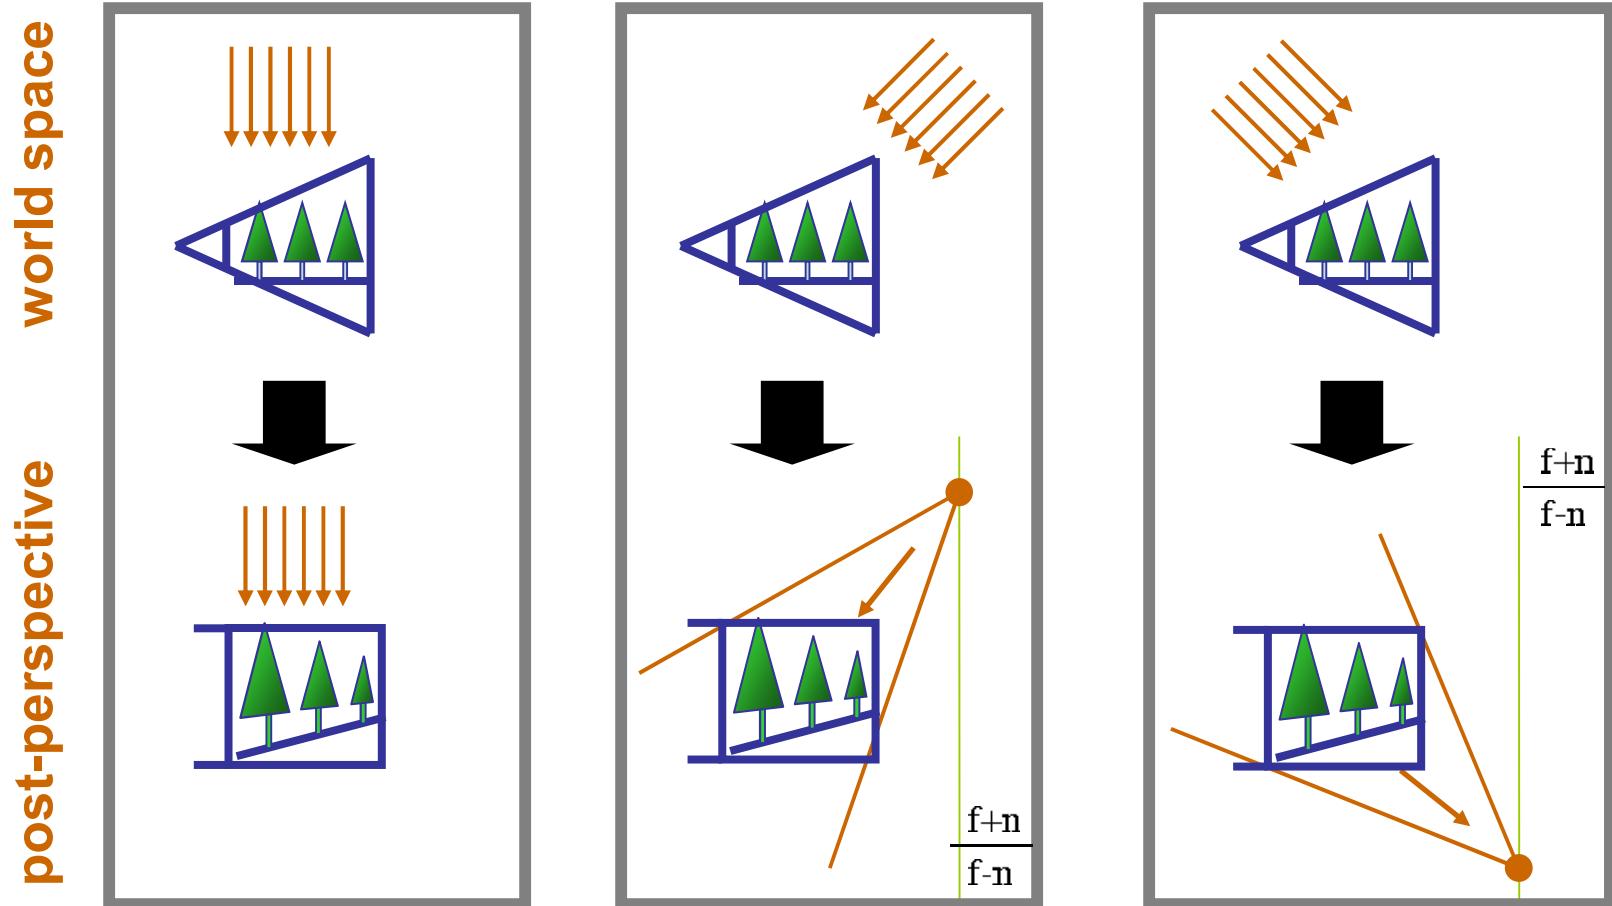
Post Perspective Space - Depth



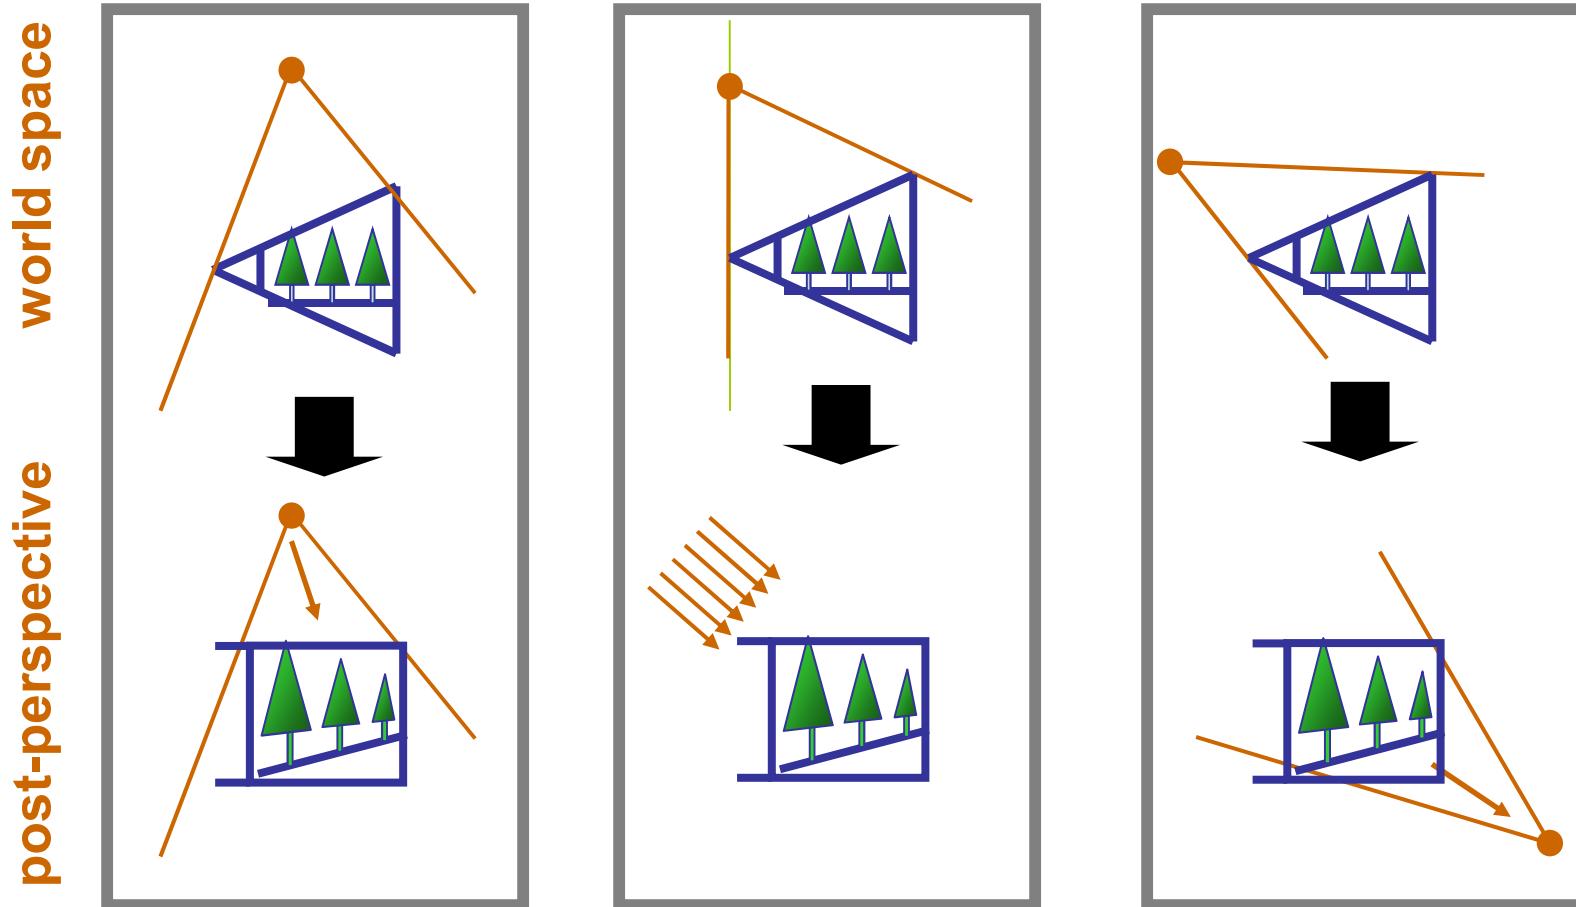
Post Perspective Space – x or y



Parallel Light Transformation



Point Light Transformation

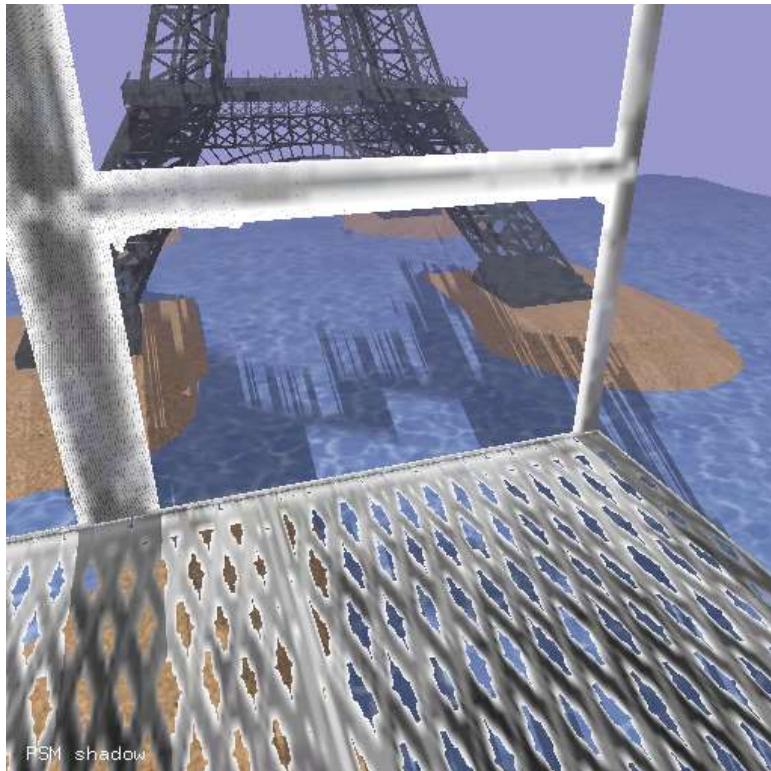


PSM Problems

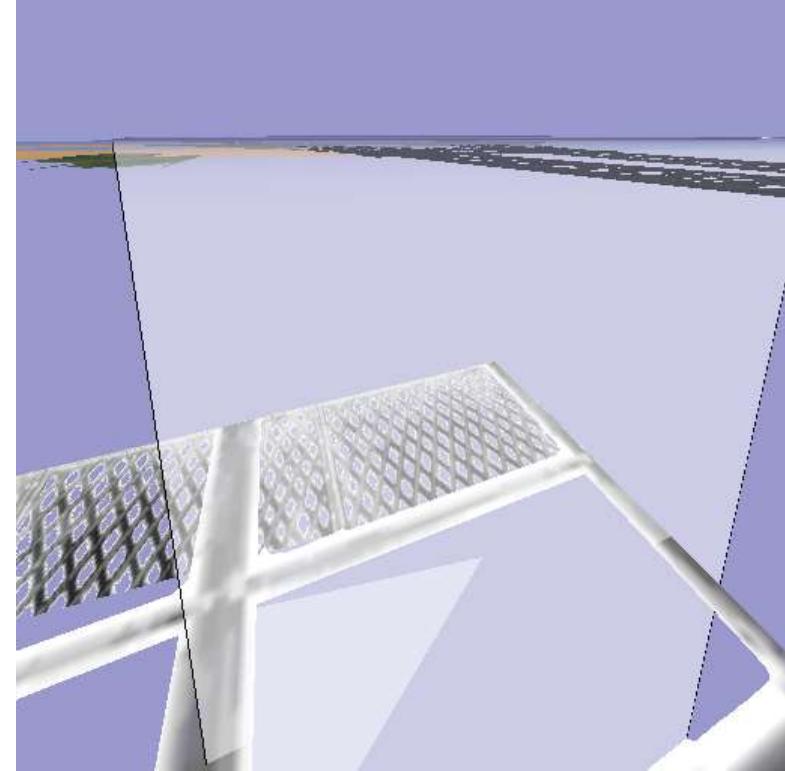
- Shadows from behind
 - Require move back due to singularity
 - Changes resulting quality of shadow
- Lights change their type (point/directional)
 - Many special cases
 - Non-intuitive post-perspective space
- View-dependent shadow quality

PSM Problems

- Most severe: uneven z-distribution
 - Good near viewer, very bad far away!
 - Can be reduced by pushing near plane away



world space



post-perspective space

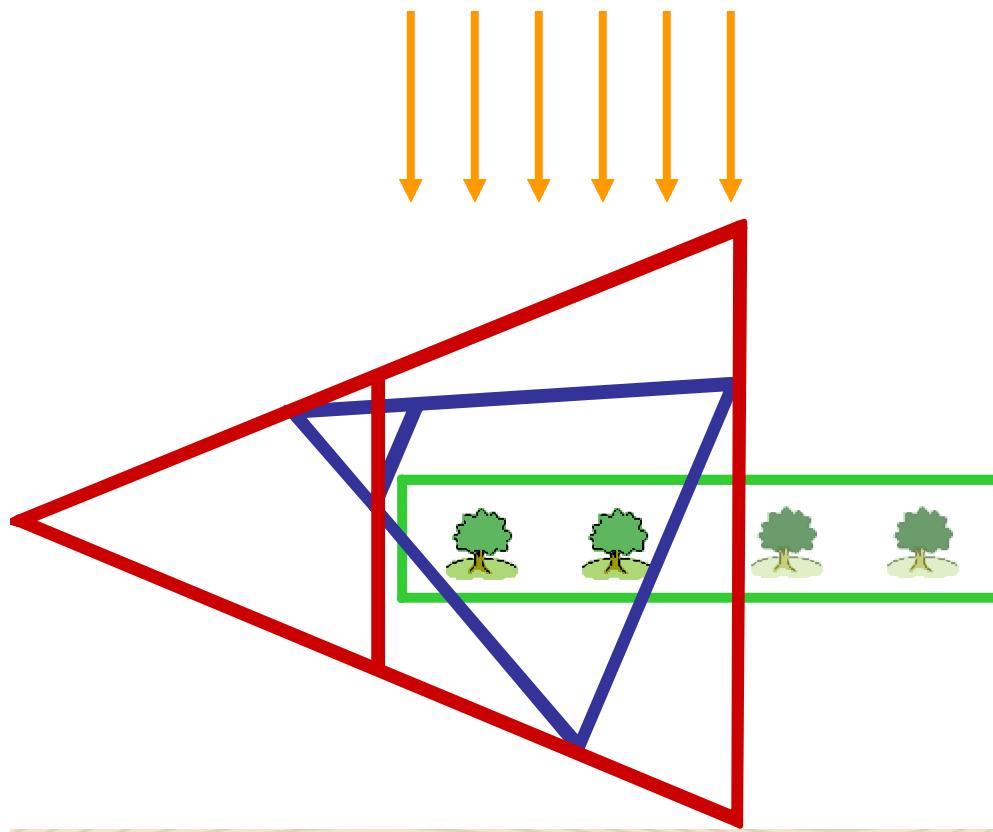
Observation

- Want to redistribute samples in shadow map
- Perspective transform is a good option
 - Supported by hardware
- BUT: why choose perspective based on observer transform???

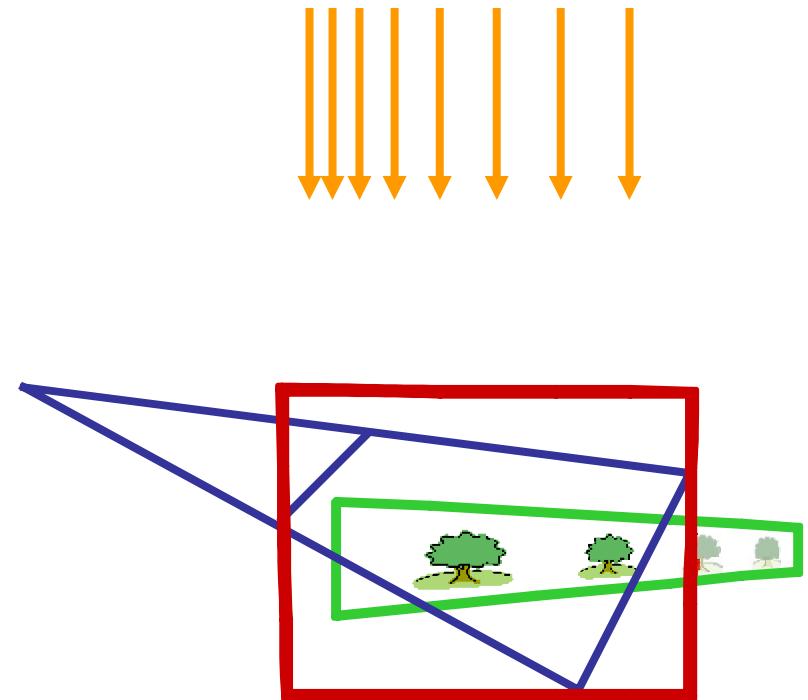
Light Space Perspective Shadow Maps

- ◆ Use an **additional perspective frustum** in **light space**

uniform shadow map



LiSPSM shadow map



Light Space



light space

Advantages of LiSPSM

- No singularities in relevant part of the scene
 - ... by construction
- Directional lights remain directional lights
 - ... and direction remains the same
- Points lights are converted to directional lights
 - ... perspective point light transformation already done in light space
- Easy construction, general and robust
- Smooth quality changes

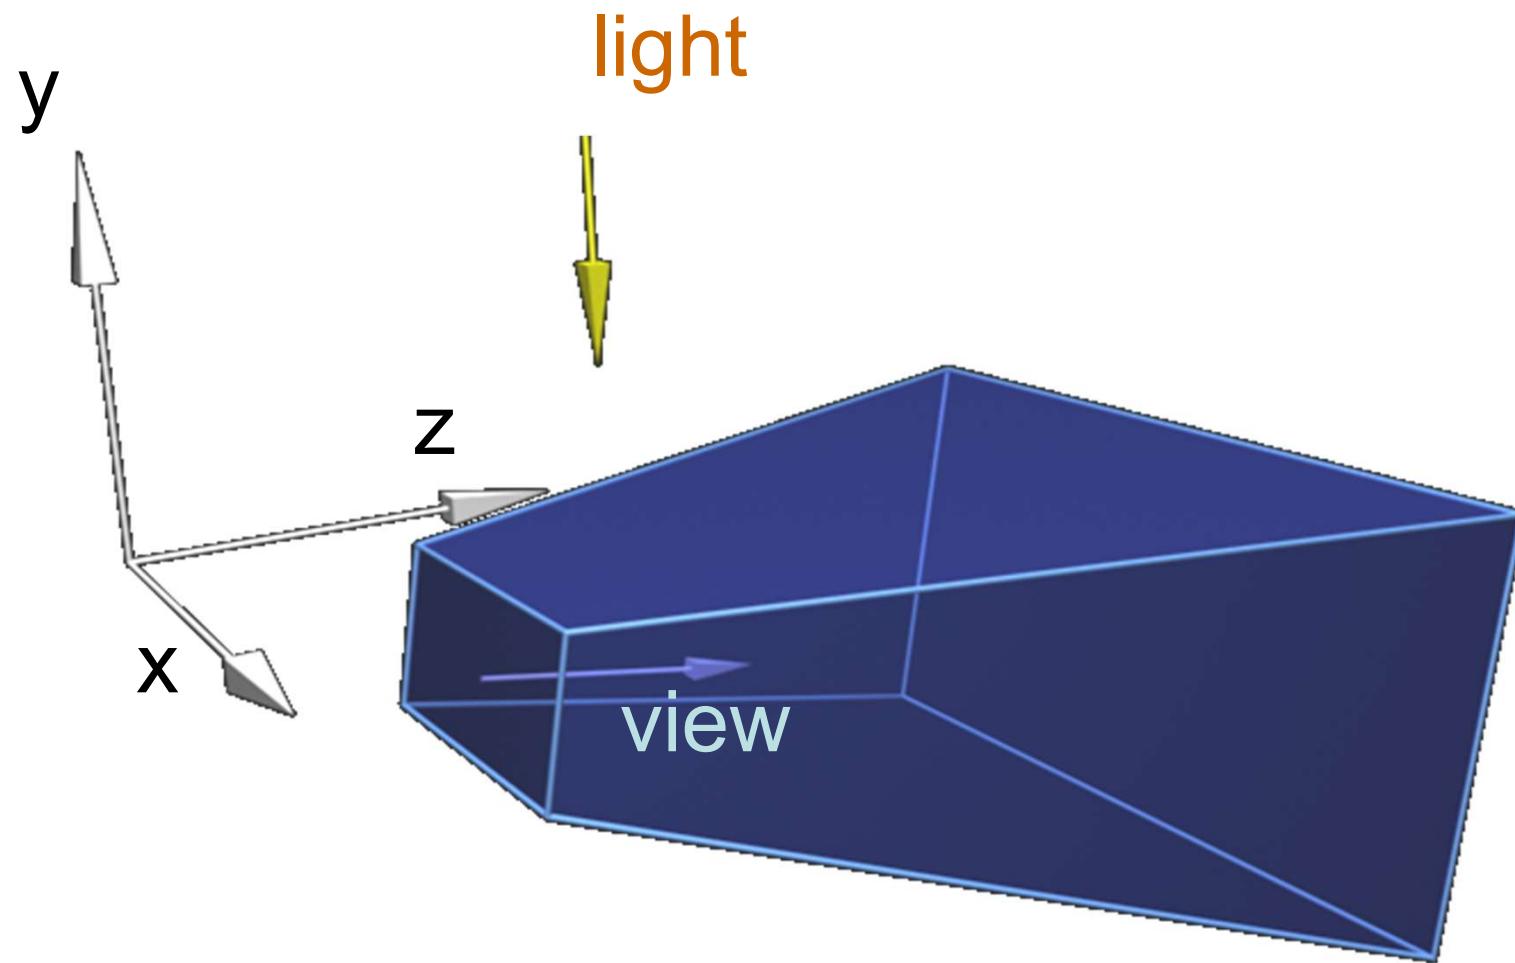
- Most important: **control over shadow quality**
 - ... by choice of free parameter

LiSPSM Overview

- Focus shadow map
- Construct light space
- Create perspective transform
- Apply

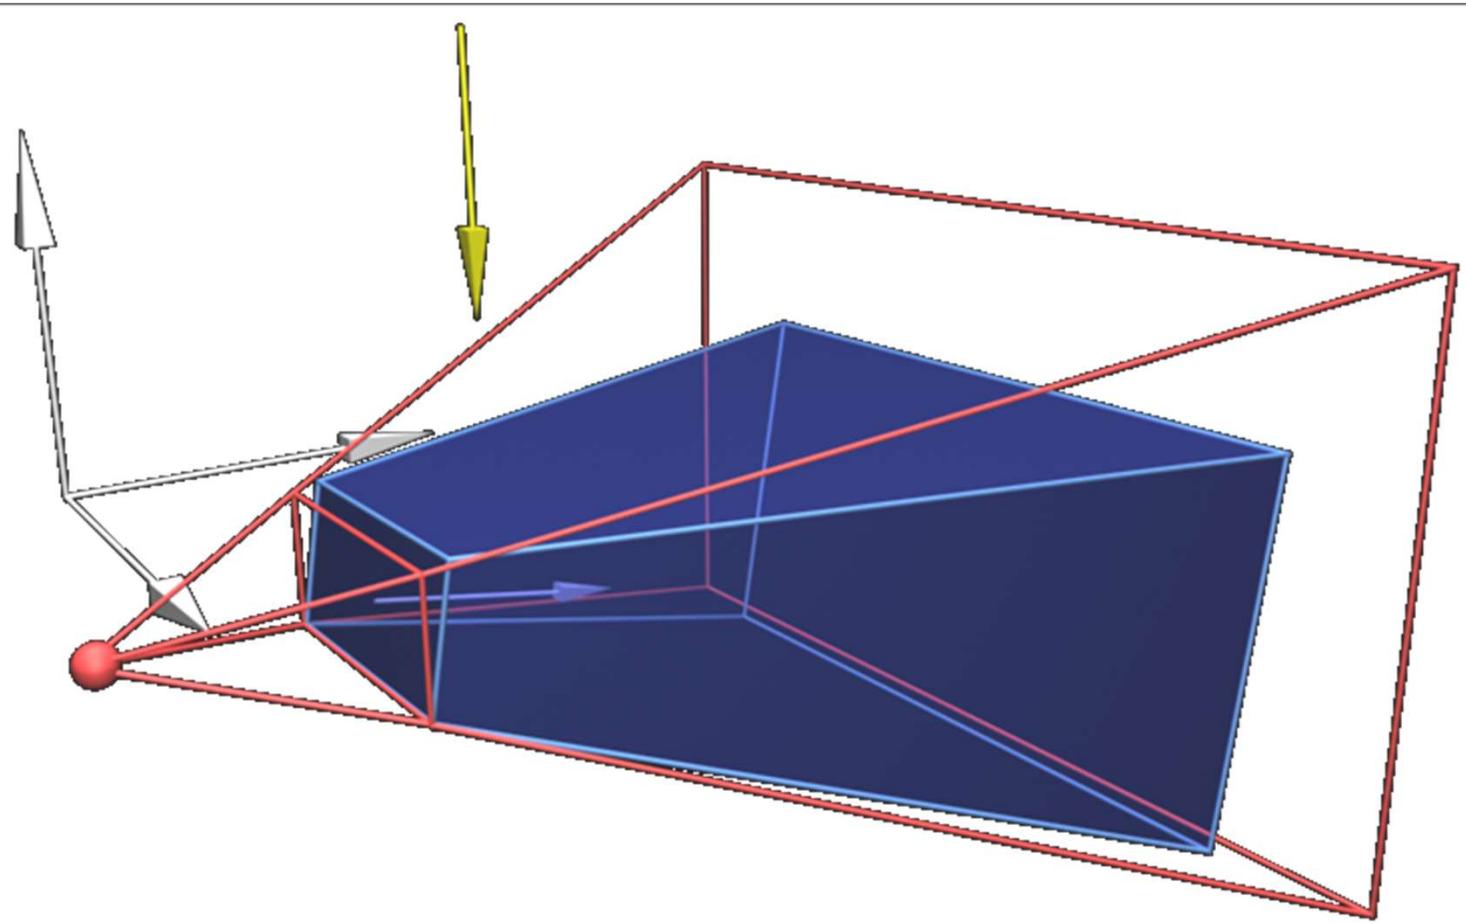
LiSPSM Fitting and Warp Construction

- Light and **view** vector define yz-plane

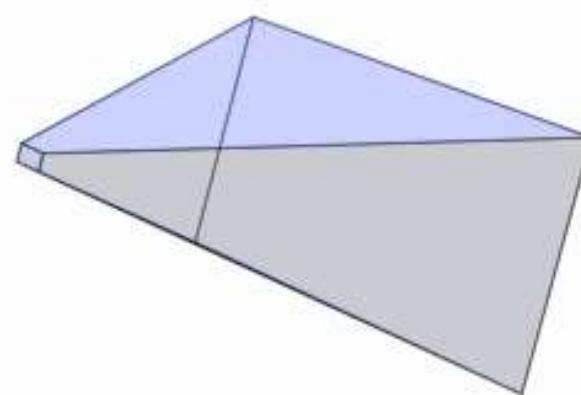
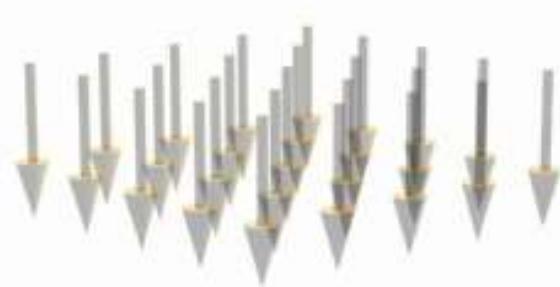


LiSPSM Fitting and Warp Construction

- Find a **tight perspective frustum** on focused region
 - In light space!

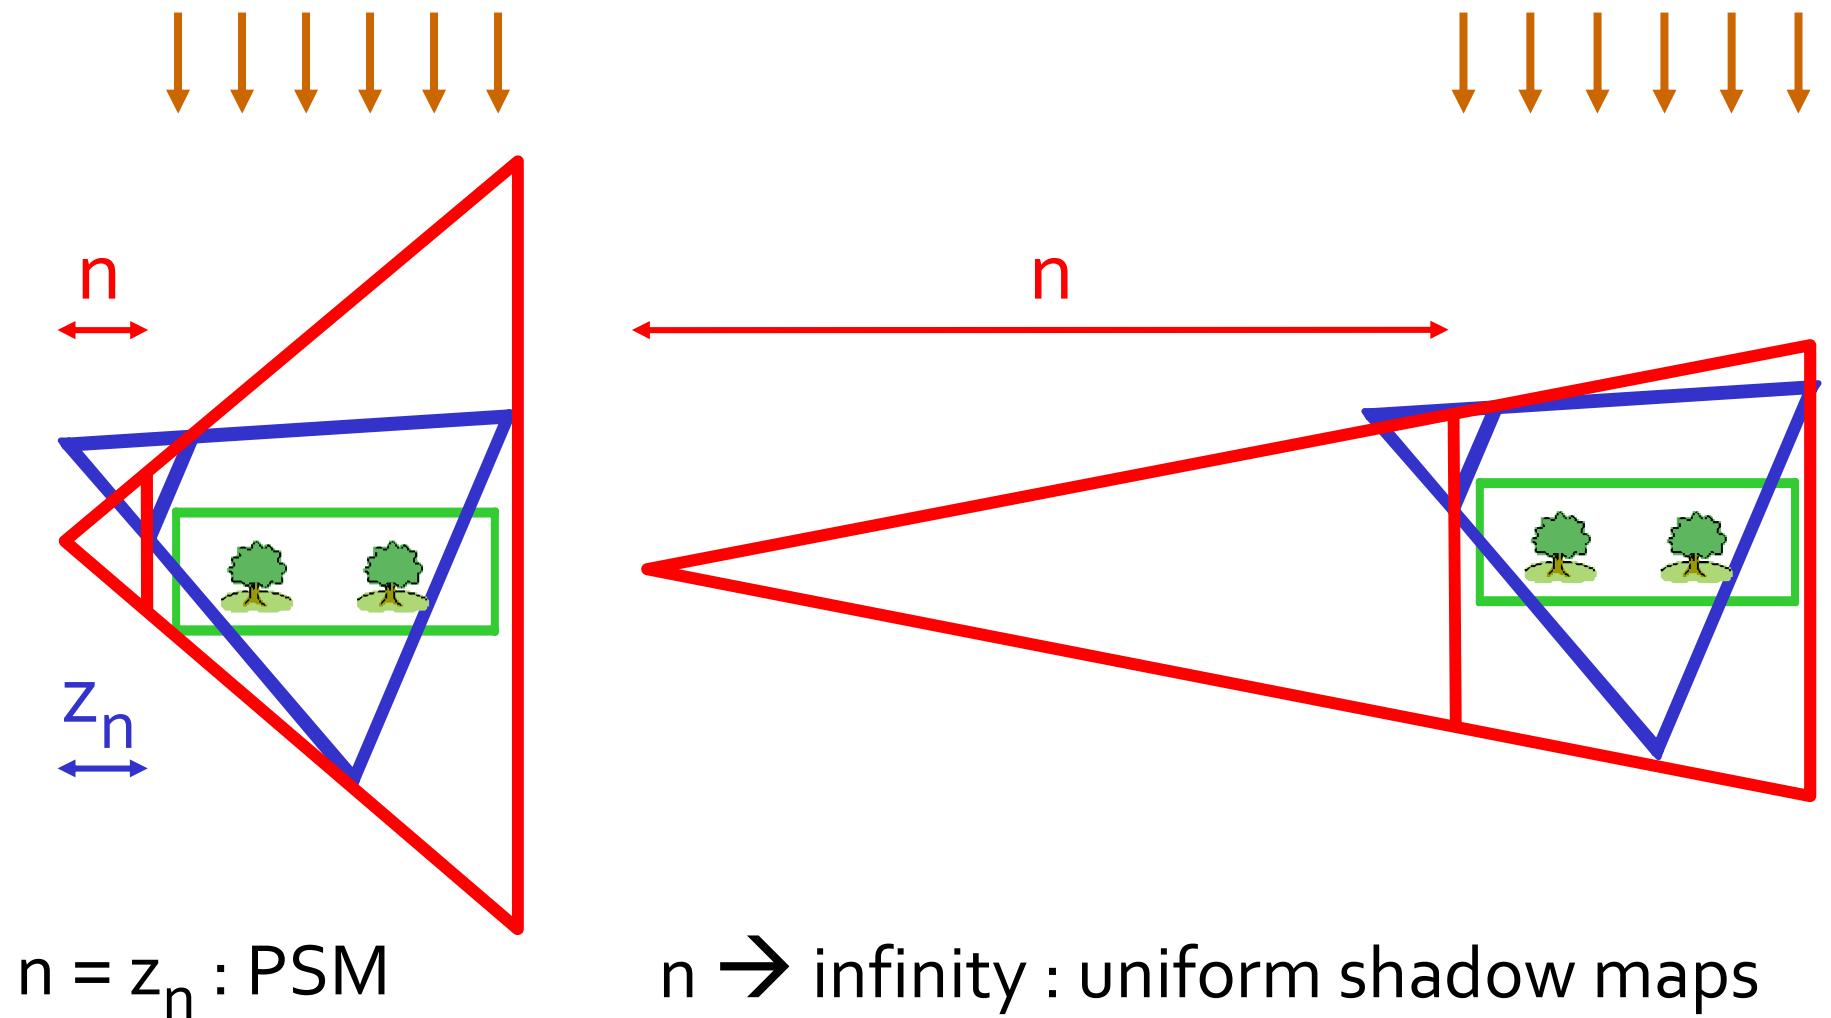


LiSPSM Construction

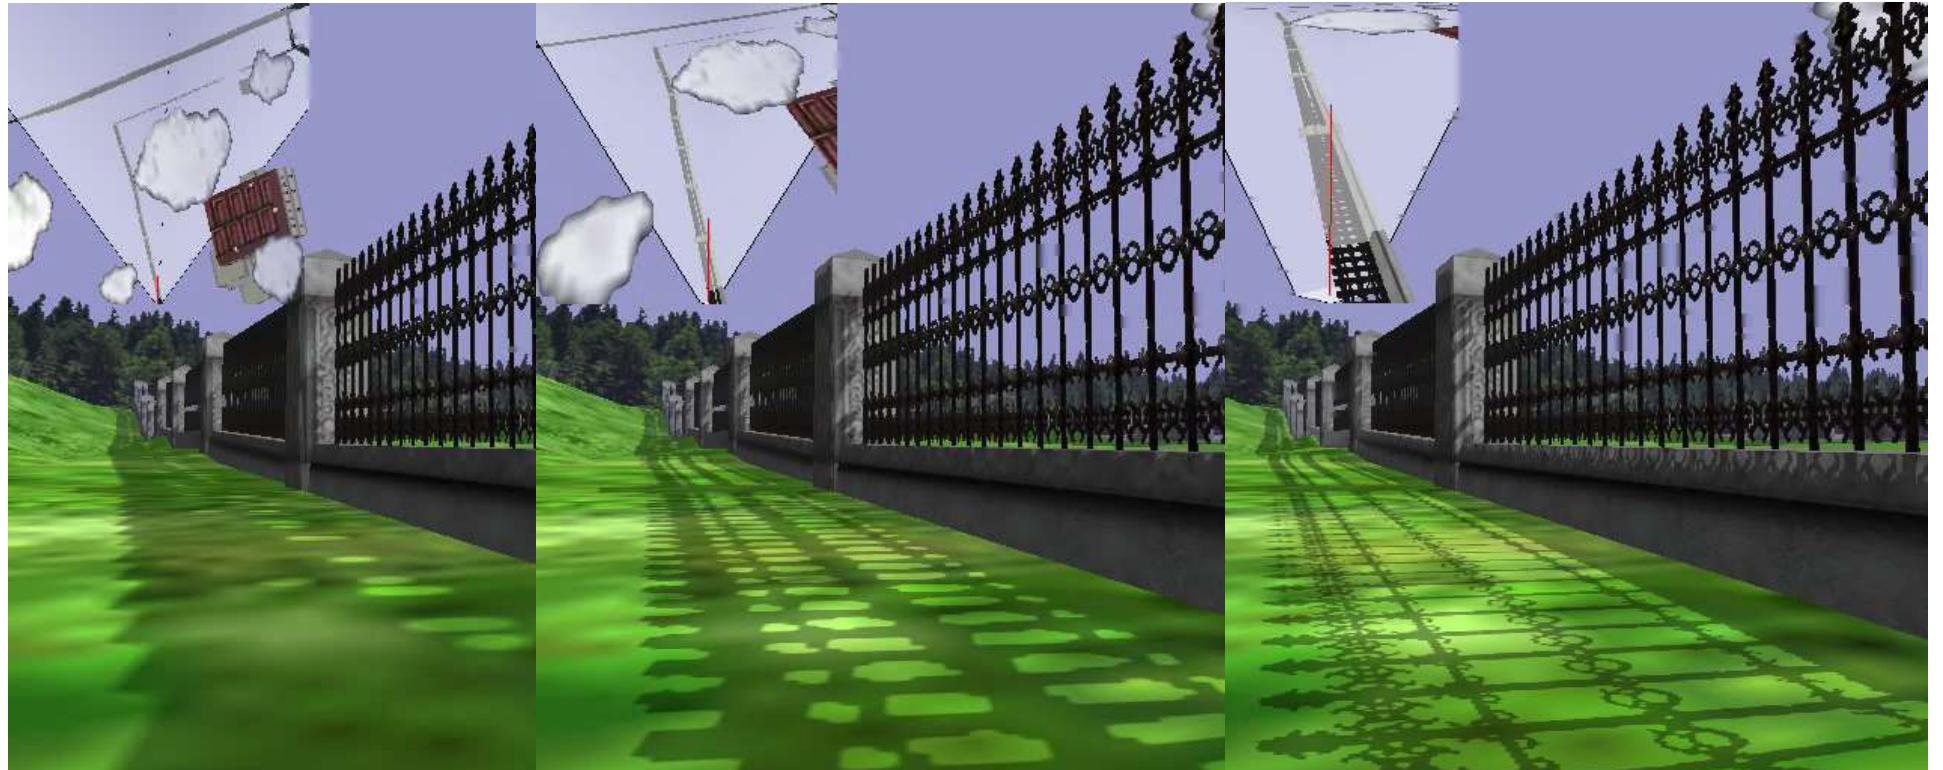


Free Parameter n

- Controls warping effect



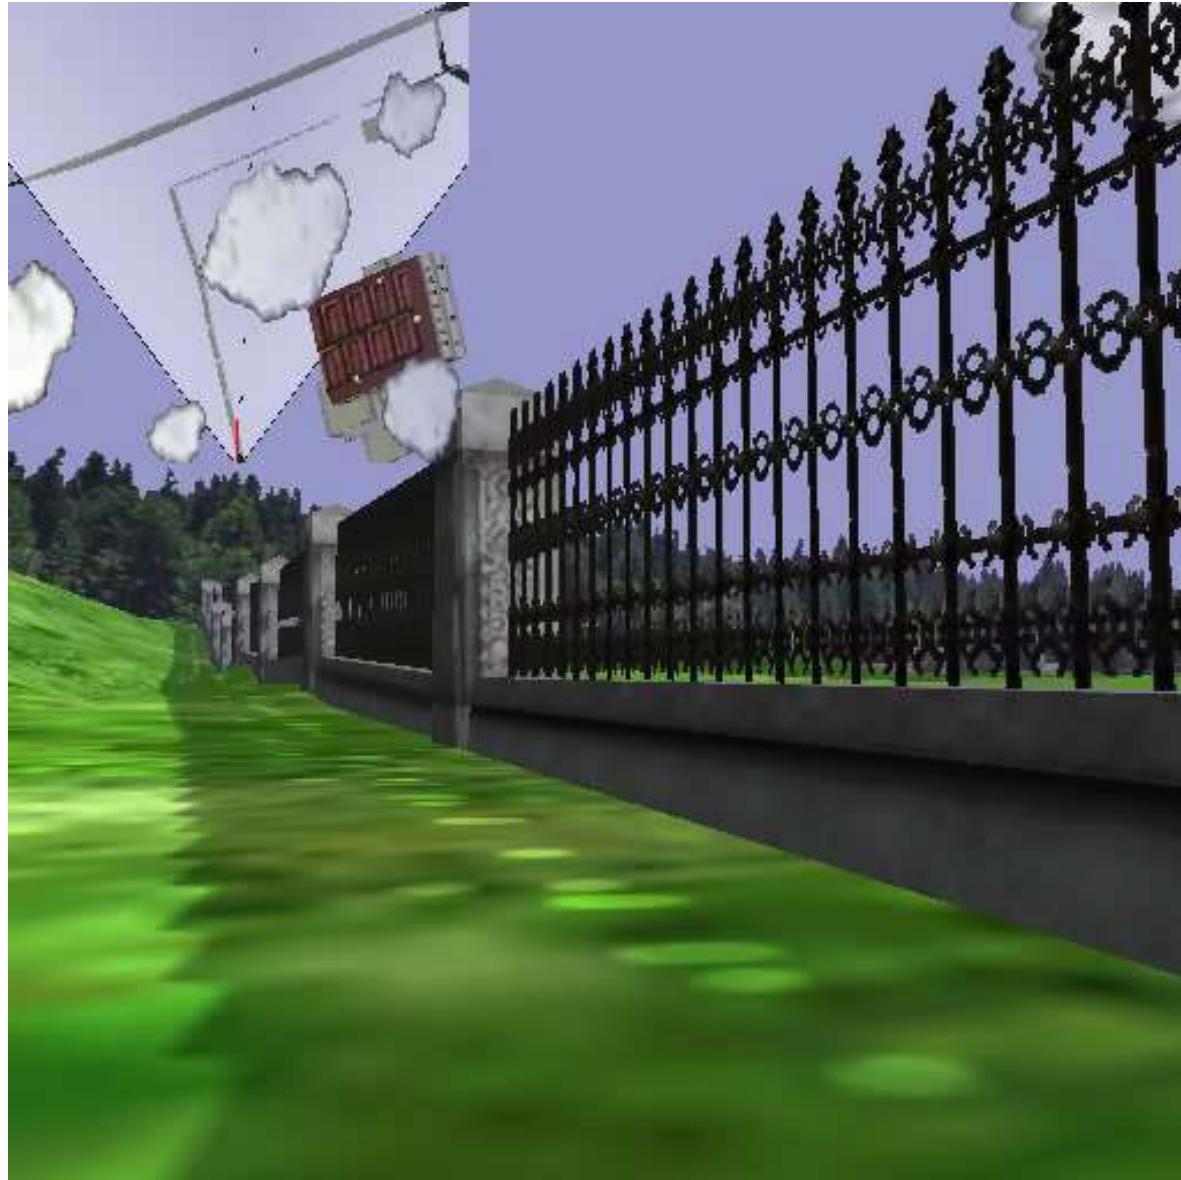
Free Parameter n



very big n

very small n

Free Parameter n



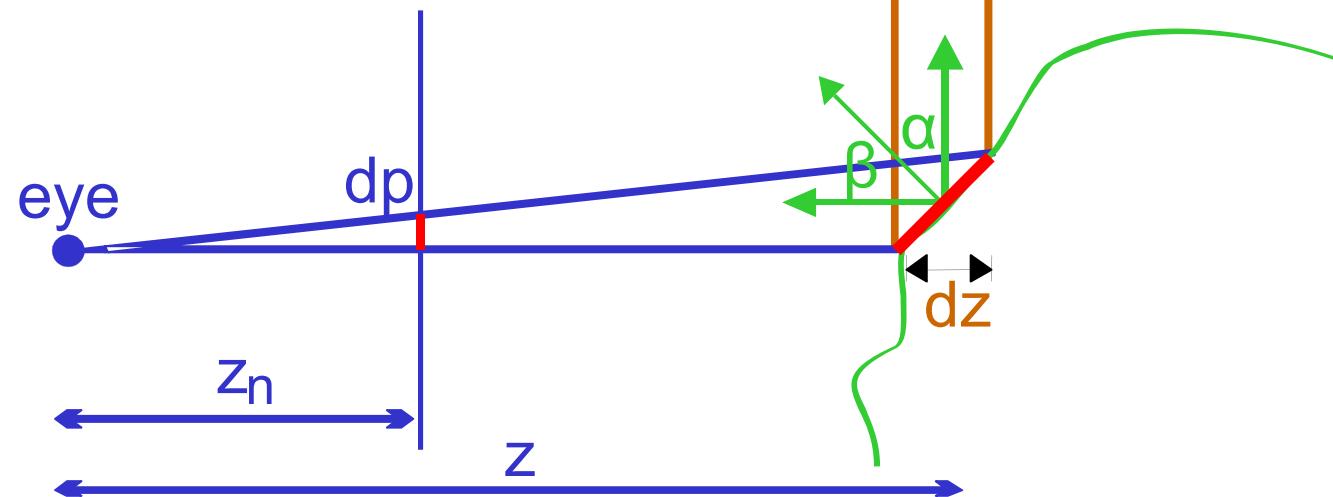
Sampling Error Analysis

- Aliasing error:
 - Parameterization
 - Perspective
 - Projection

$$\frac{dp}{ds} = \frac{z_n dz \cos \alpha}{z ds \cos \beta}$$

↓ ↓ ↓ ↓

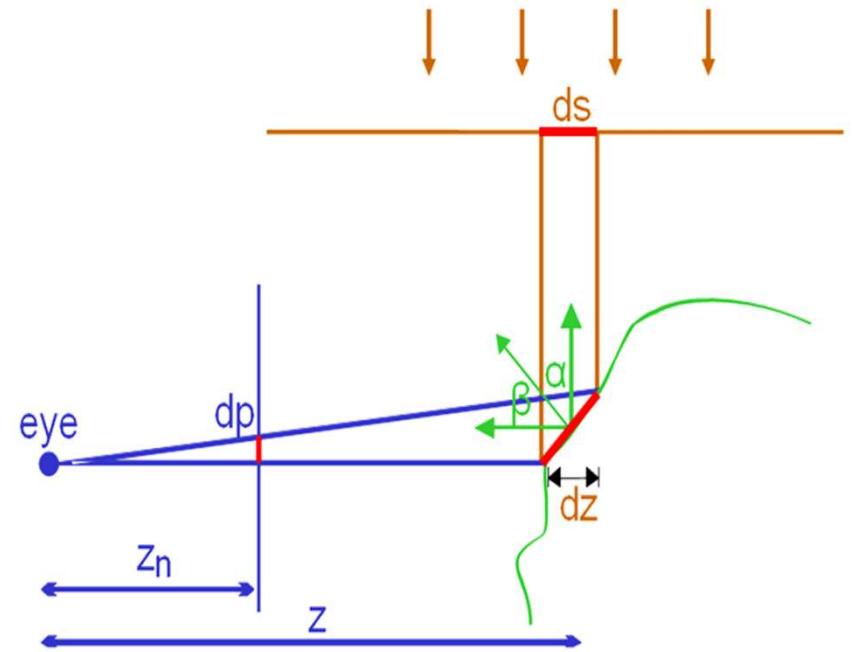
ds



How to Choose the Free Parameter n ?

- Recall error analysis
- $\frac{dp}{ds} > 1 \rightarrow$ shadow map undersampling
- **Projection aliasing** cannot be changed
- Counter **perspective aliasing** with new shadow map **parameterization** s
- Goal: $\frac{dp}{ds} \sim 1$

$$\frac{dp}{ds} = \frac{z_n}{z} \frac{dz}{ds} \frac{\cos \alpha}{\cos \beta}$$



Error Analysis

$$\frac{dp}{ds} = \boxed{\frac{1}{z} \frac{dz}{ds}}$$

- Uniform shadow maps

$$s \sim z \quad \xrightarrow{s' \atop d} \quad \frac{ds}{dz} \sim 1 \quad \Rightarrow \quad \frac{dp}{ds} \sim \frac{1}{z}$$

- Perspective shadow maps

$$s \sim \frac{1}{z} \quad \xrightarrow{s' \atop d} \quad \frac{ds}{dz} \sim \frac{1}{z^2} \quad \Rightarrow \quad \frac{dp}{ds} \sim z$$

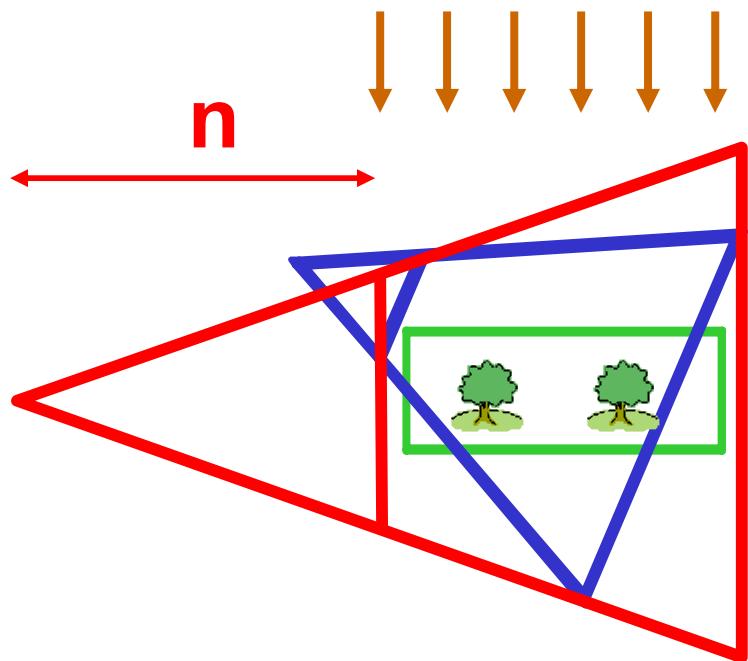
- Linear increase in error!
- Perfect: logarithmic re-parameterization

$$s \sim \log z \quad \xrightarrow{s' \atop d} \quad \frac{ds}{dz} \sim \frac{1}{z} \quad \Rightarrow \quad \frac{dp}{ds} \sim 1$$

- Hardware support? [Lloyd 2007, 2008]

Error Analysis: LiSPSM Optimal Choice

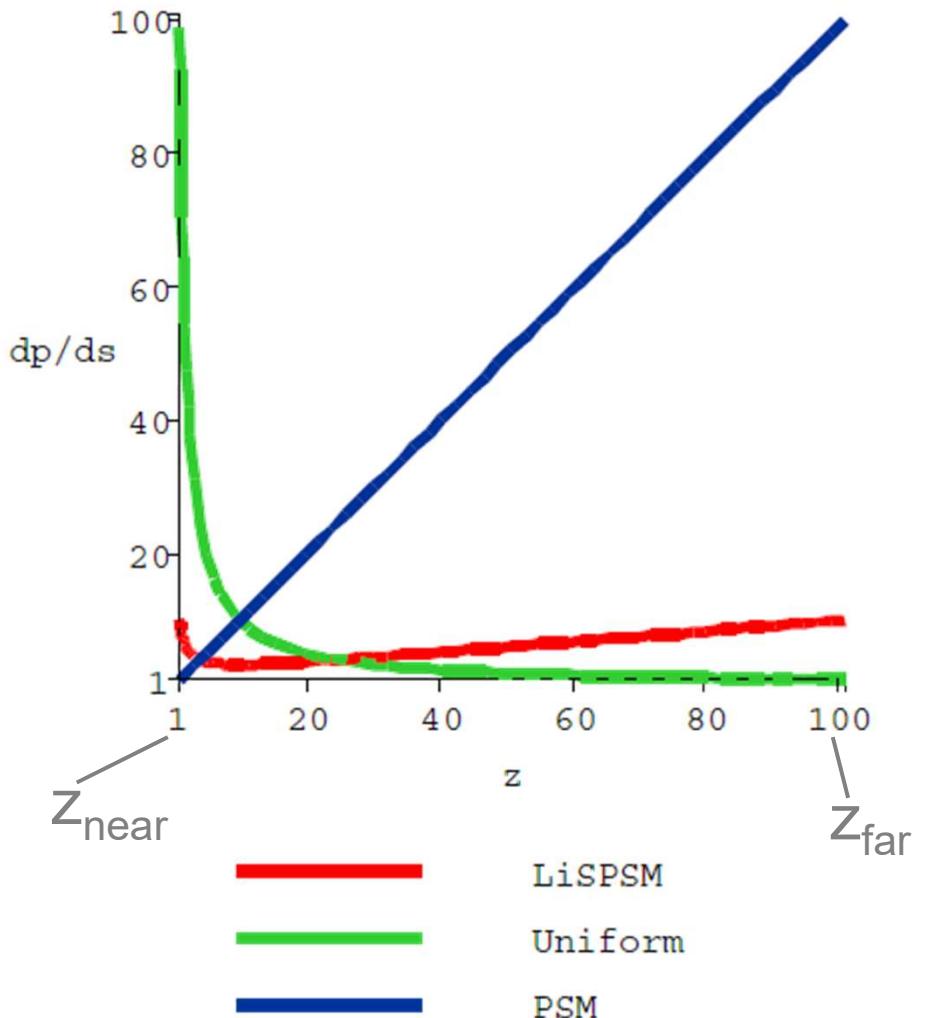
- For LiSPSM, $\frac{dp}{ds}$ depends on **n**
 - Gives $\frac{dp}{ds}$ between uniform and perspective
- Optimal choice:



$$n_{opt} = z_n + \sqrt{z_f z_n}$$

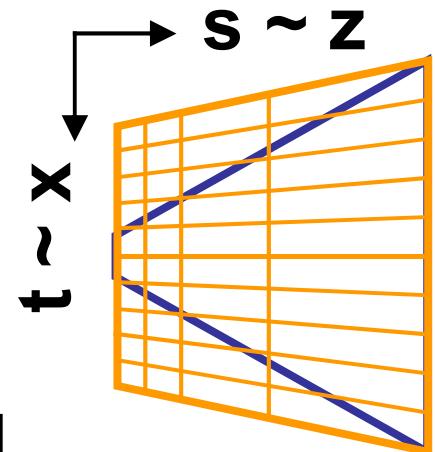
Error Comparison

- LiSPSM optimal choice
- Measured along view dir
- LiSPSM vs PSM
 - for same depth range, LiSPSM error much lower
- More advanced analysis was done in [Lloyd 2006]

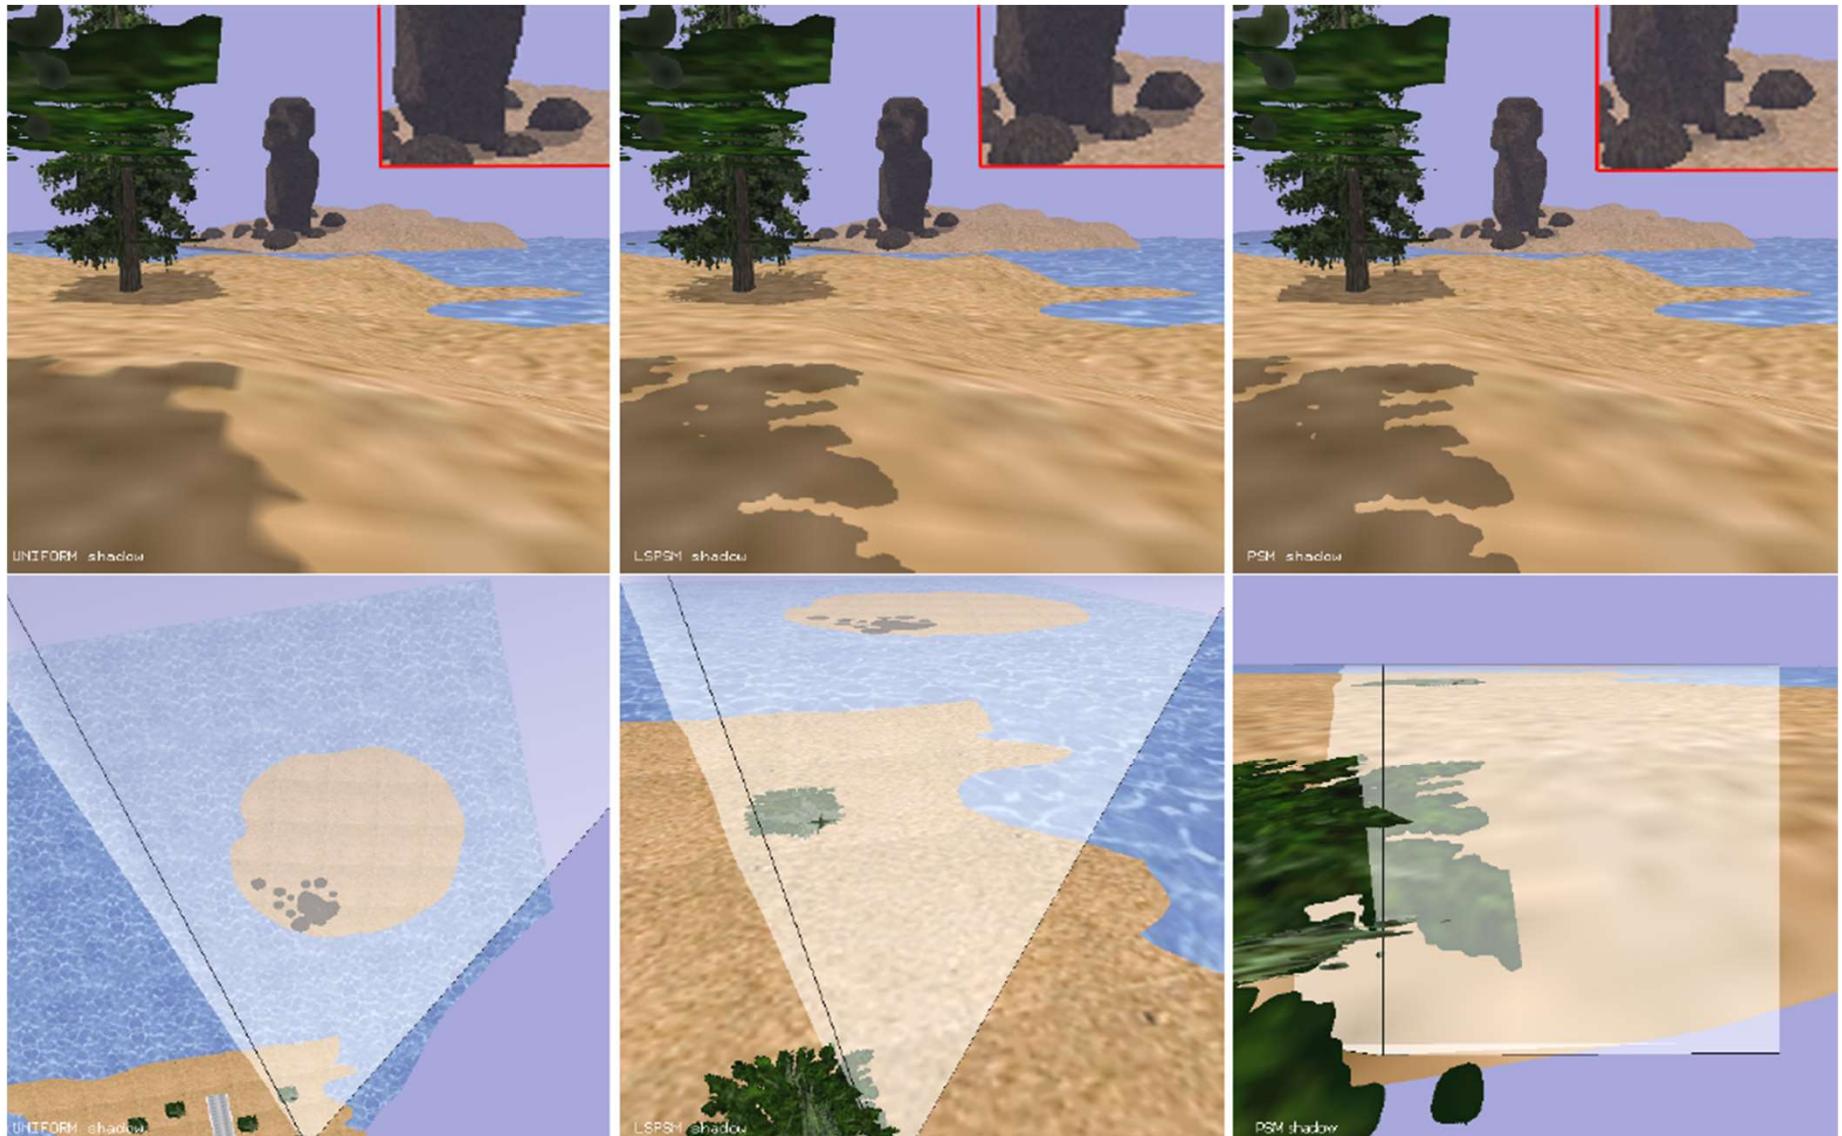


Error Comparison

- Caveat: only measured along view direction
- What about $\frac{dp}{dt}$?
 - $\frac{dp}{dt} \sim 1$ for PSM, slightly worse for LiSPSM
- More advanced analysis was done in [Lloyd 2006]
 - Result: “storage factor” **constant** for n in $[1, n_{opt}]$
 - But: best error distribution for n_{opt}



Comparison



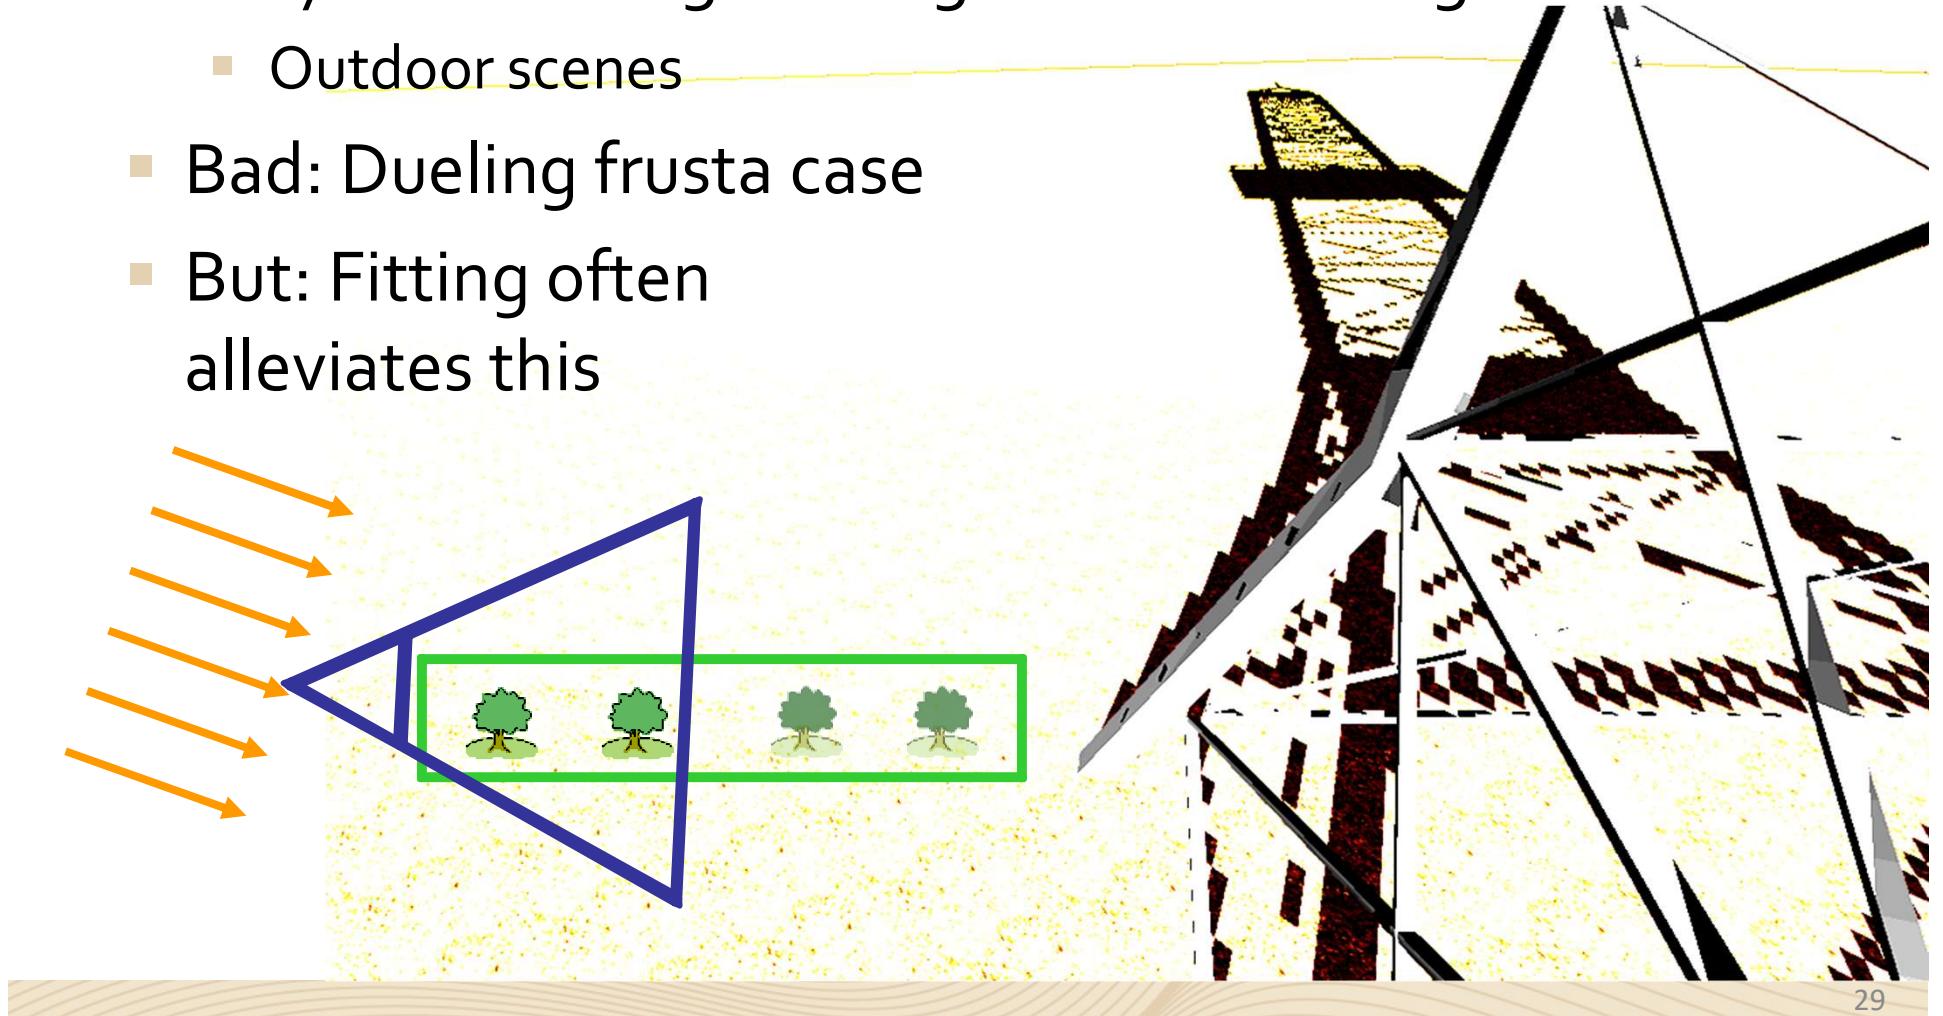
Uniform

LiSPSM

PSM

Warping: Problems

- Only works if large z-range visible from light
 - Outdoor scenes
- Bad: Dueling frusta case
- But: Fitting often alleviates this



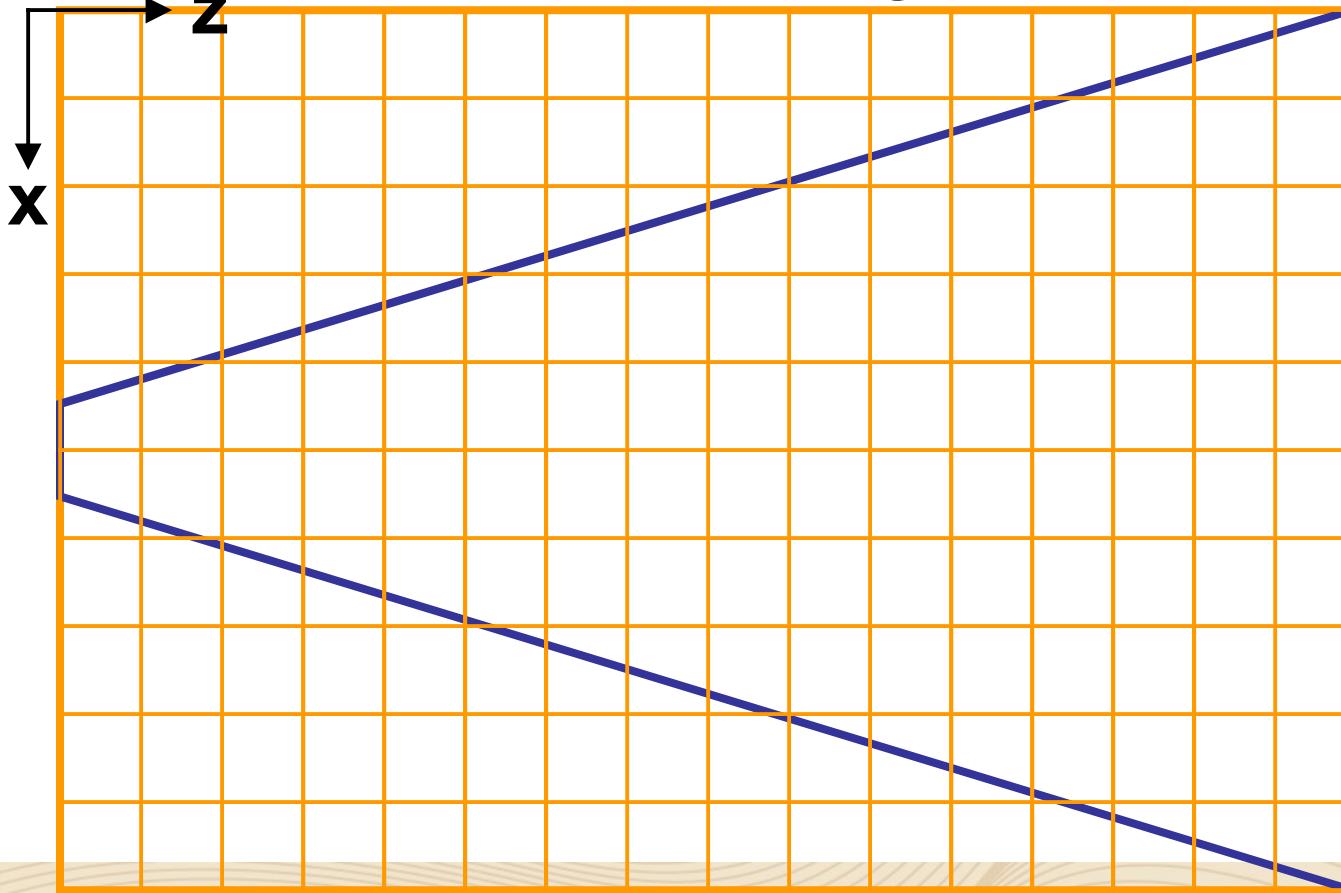
Hard Shadows

Fighting Undersampling - Partitioning

The background features a series of fine, light brown wavy lines that create a sense of depth and motion. These lines are more concentrated in the lower half of the slide, forming a base for the text area. In the upper half, they fan out and become less dense, creating a lighter, more open feel.

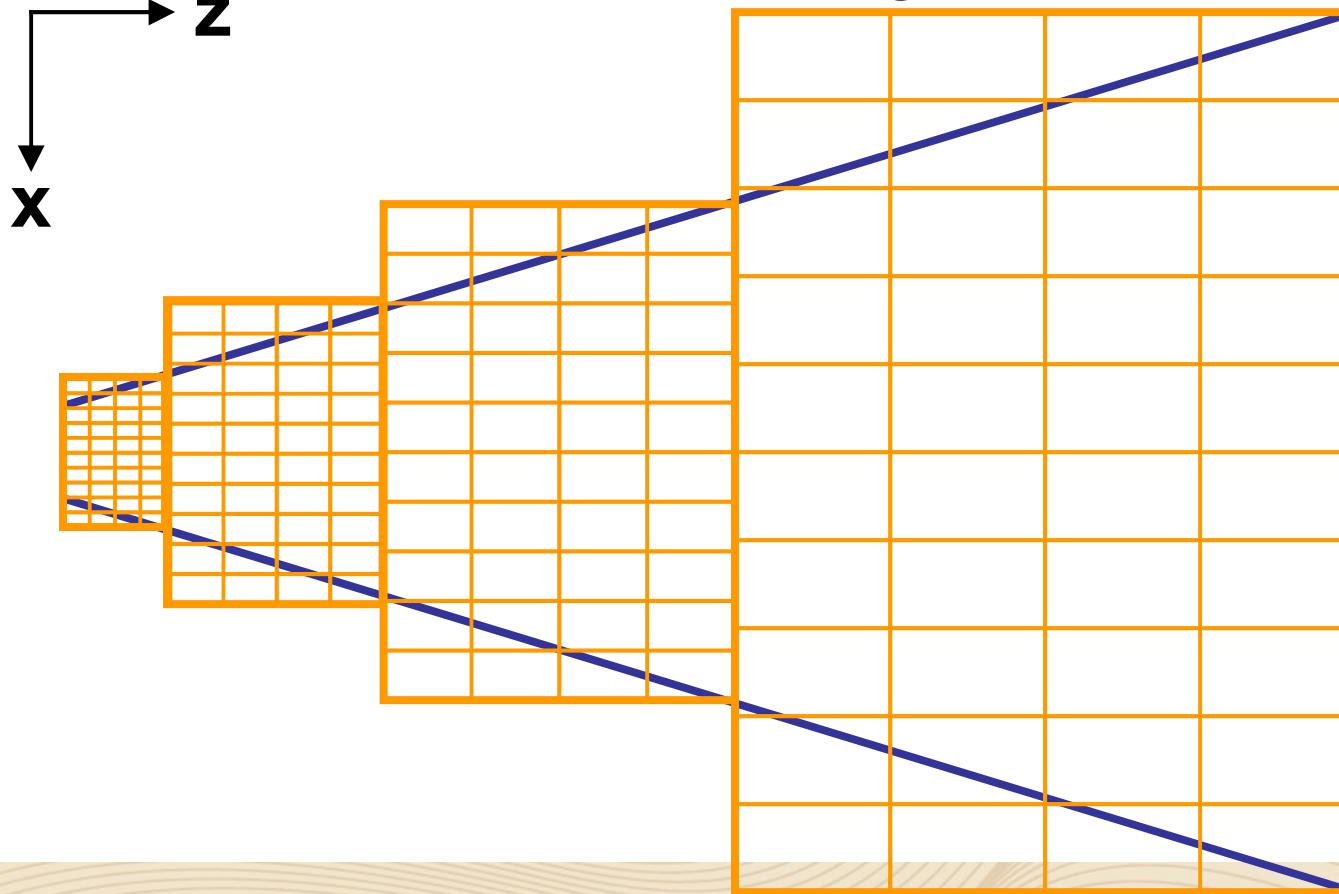
Z-Partitioning: Idea

- Parallel Split Shadow Maps [Zhang 2007]
- Cascaded Shadow Maps [Engel 2007][Zhang 2009]



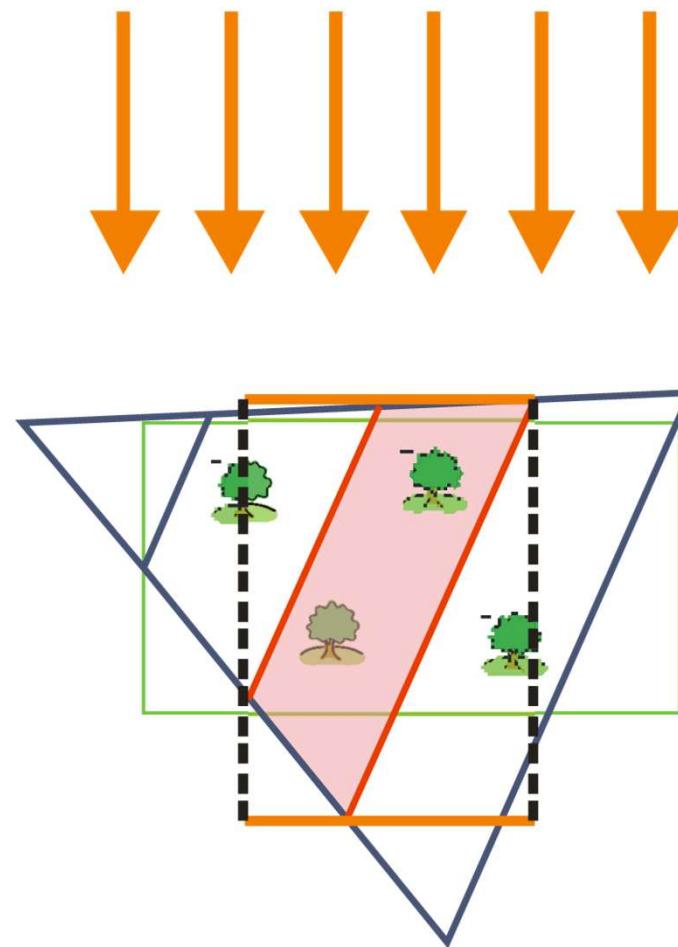
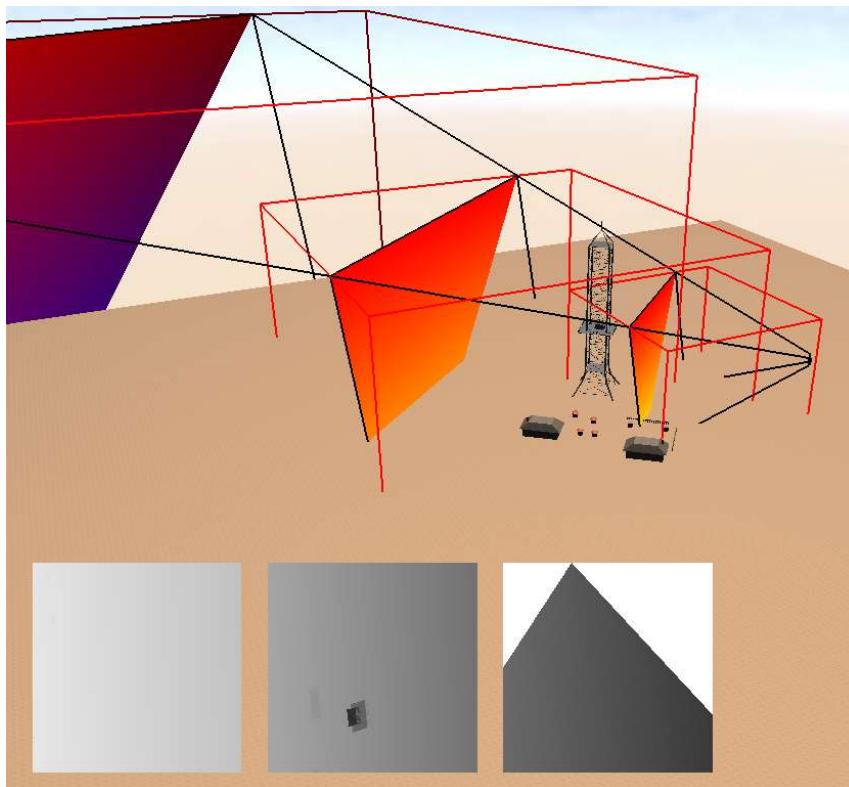
Z-Partitioning: Idea

- Parallel Split Shadow Maps [Zhang 2007]
- Cascaded Shadow Maps [Engel 2007][Zhang 2009]



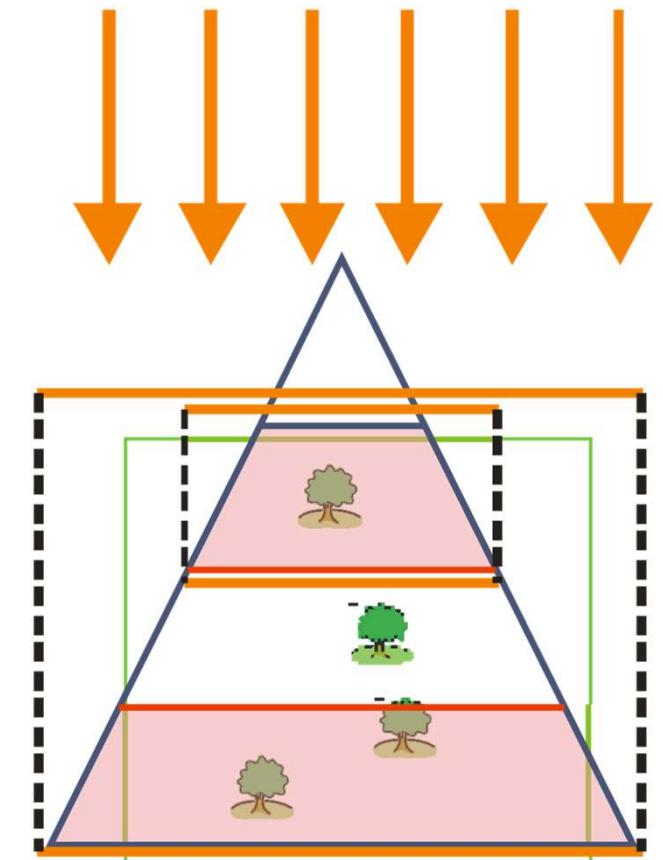
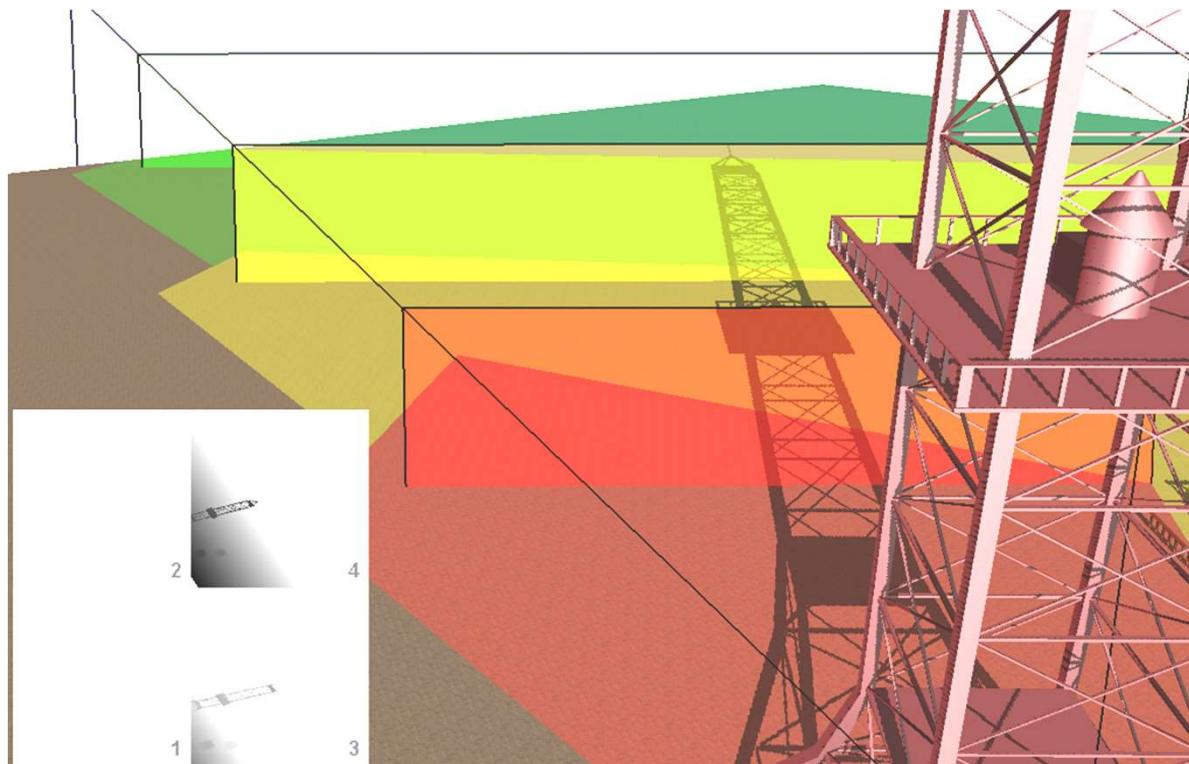
Z-Partitioning

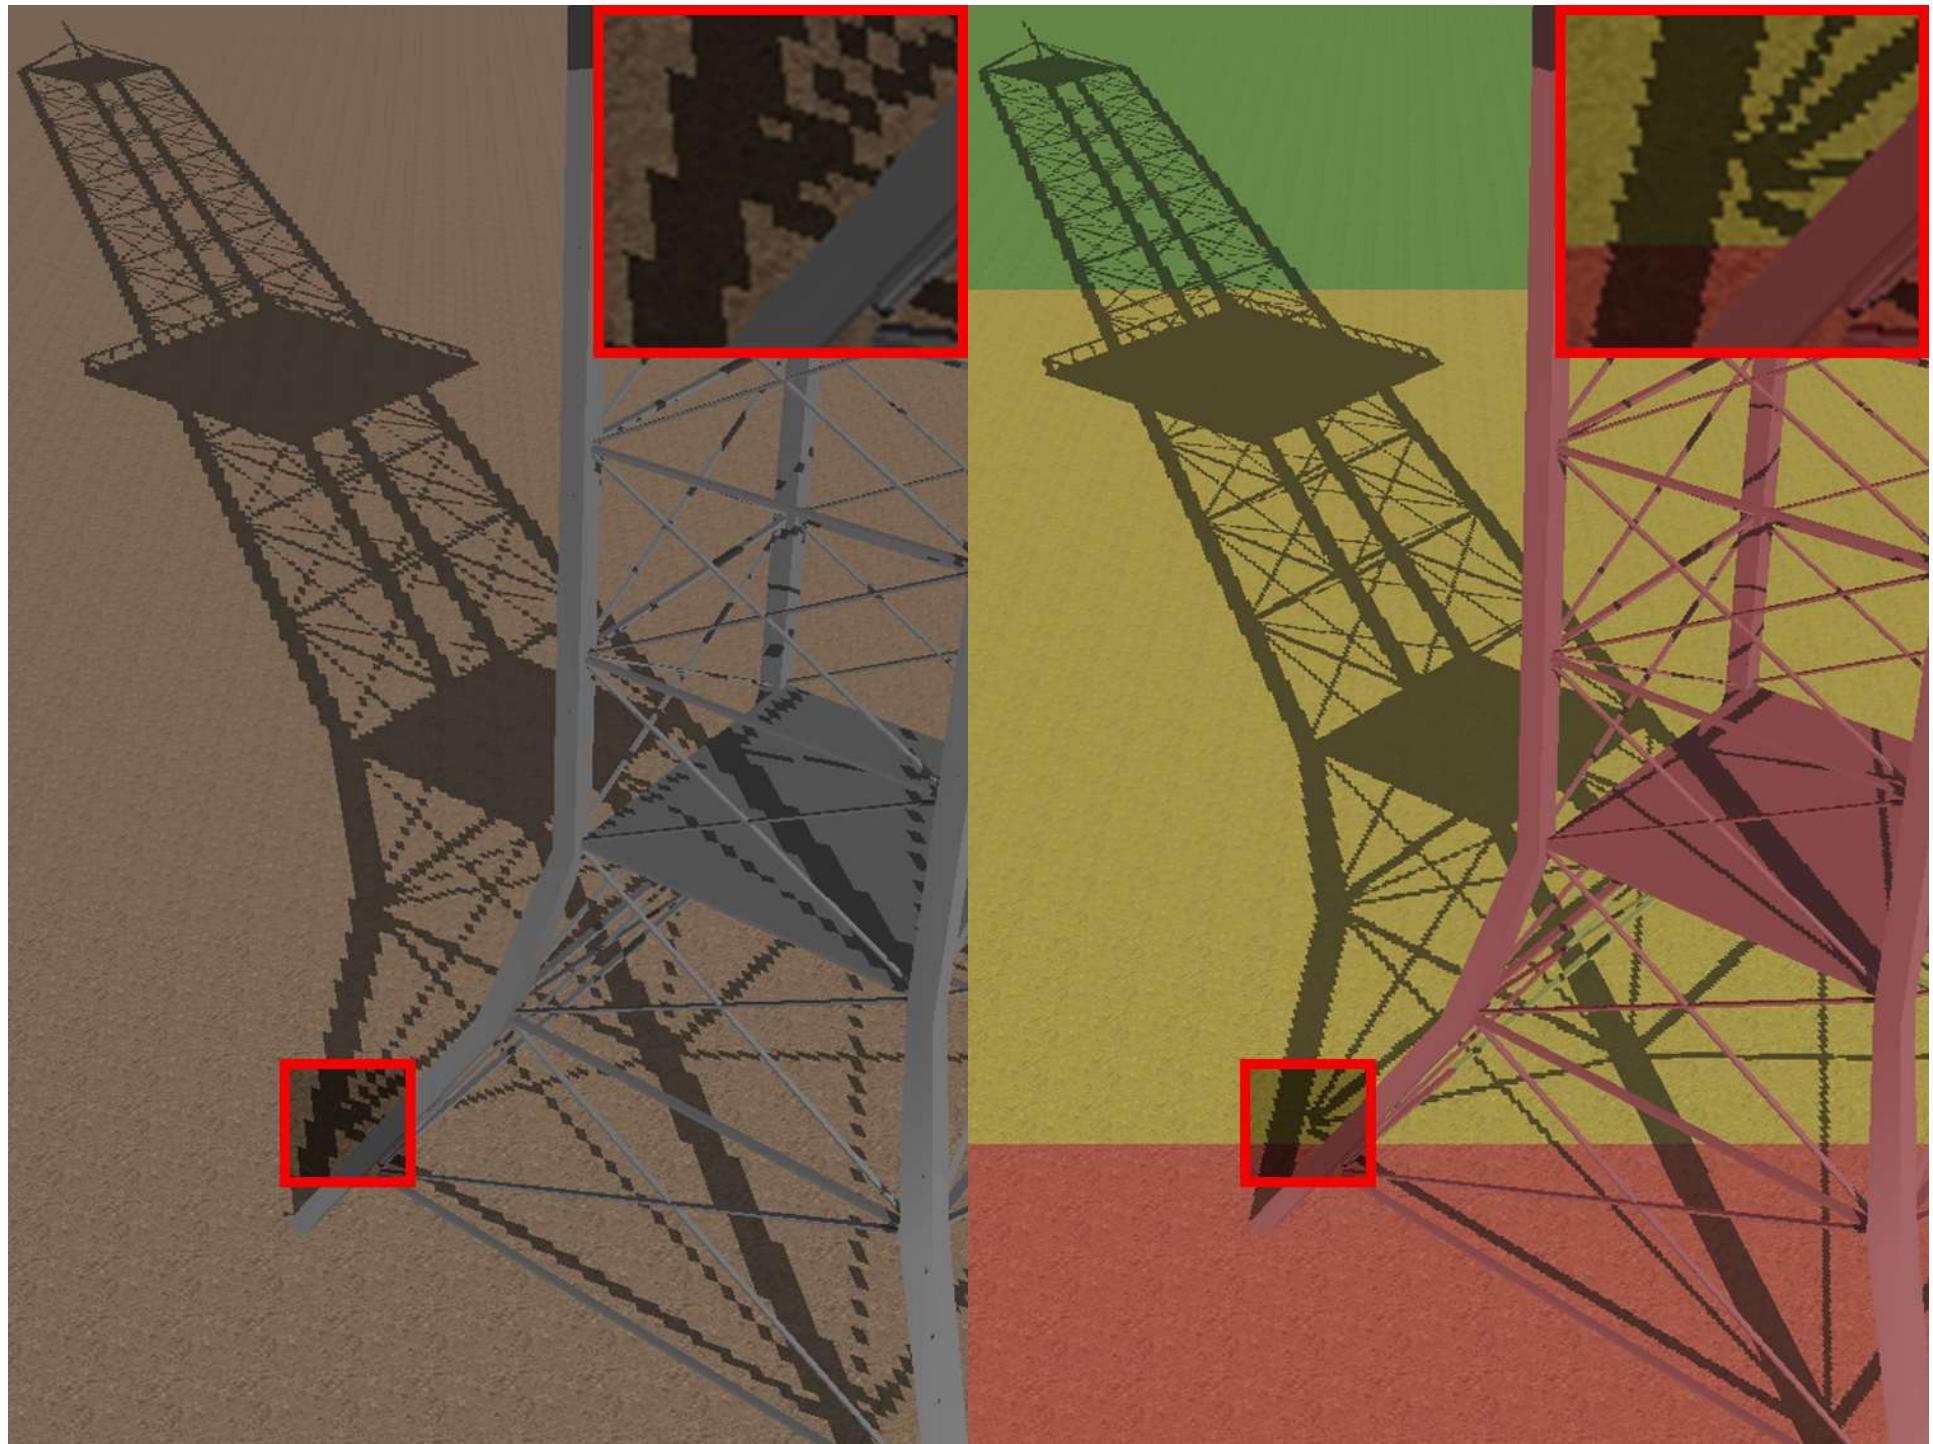
- Partition view frustum into n sub-frusta
- Calculate separate shadow map for each



Z-Partitioning

- Works even in cases where warping fails
 - Light from behind, dueling frusta

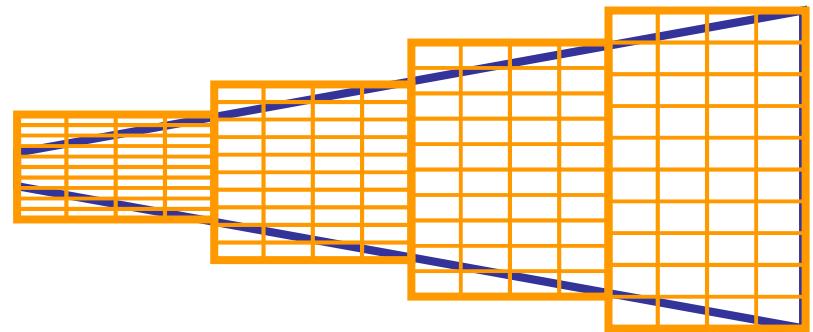




Z-Partitioning

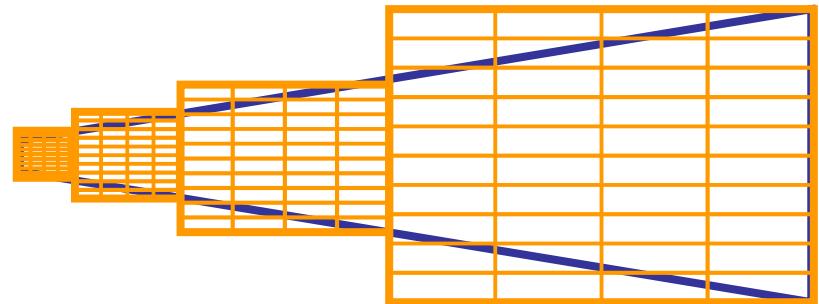
- How to choose partition sizes?

- Uniform

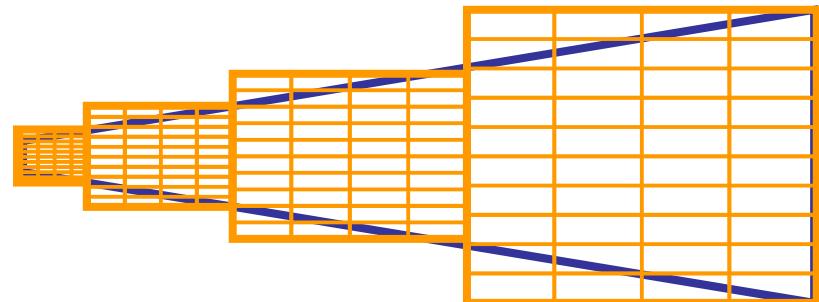


- Logarithmic/self-similar

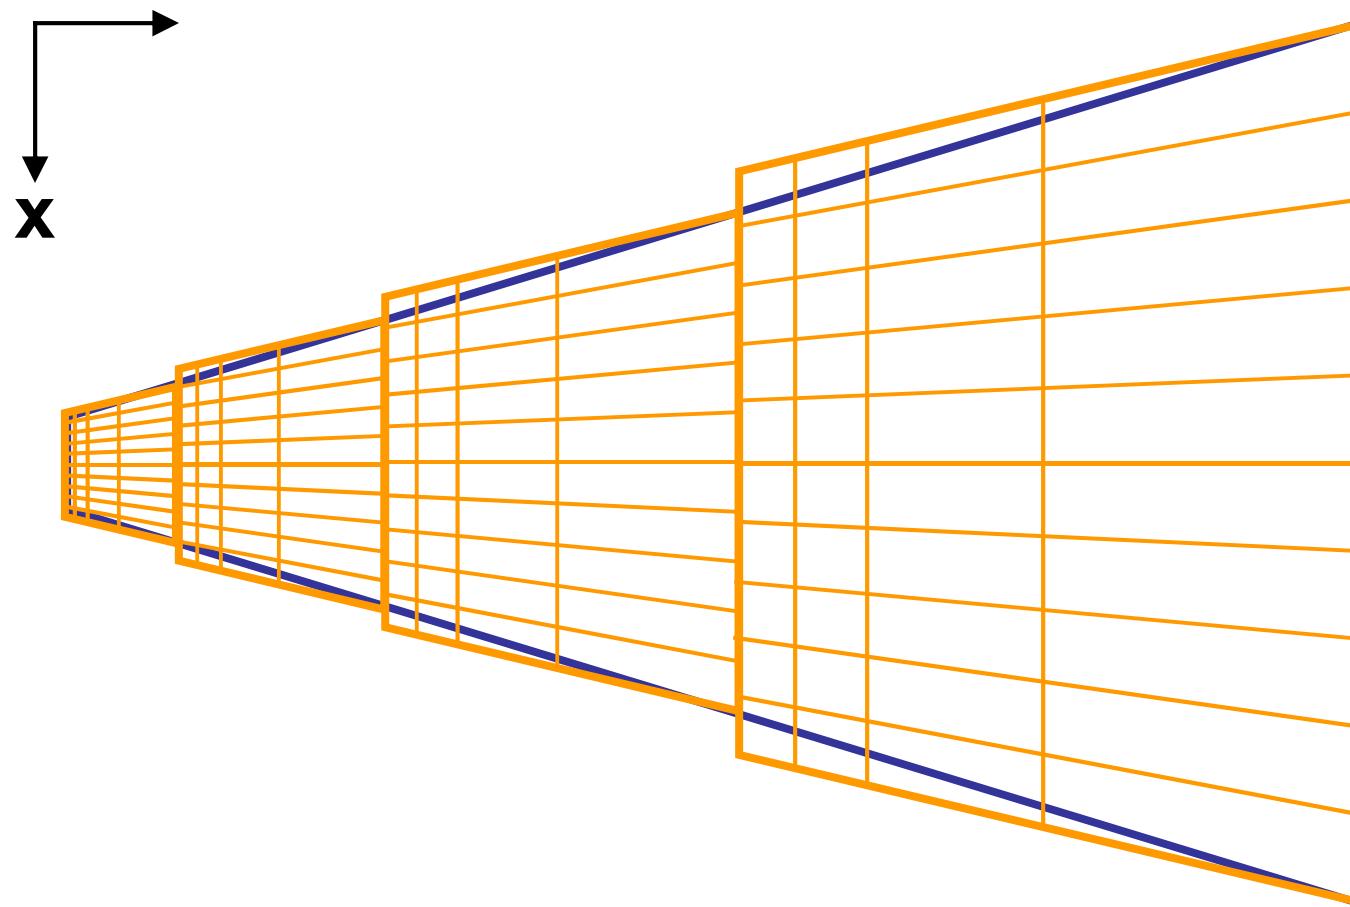
$$C_i = z_n \left(\frac{z_f}{z_n} \right)^{\frac{i}{m}}$$



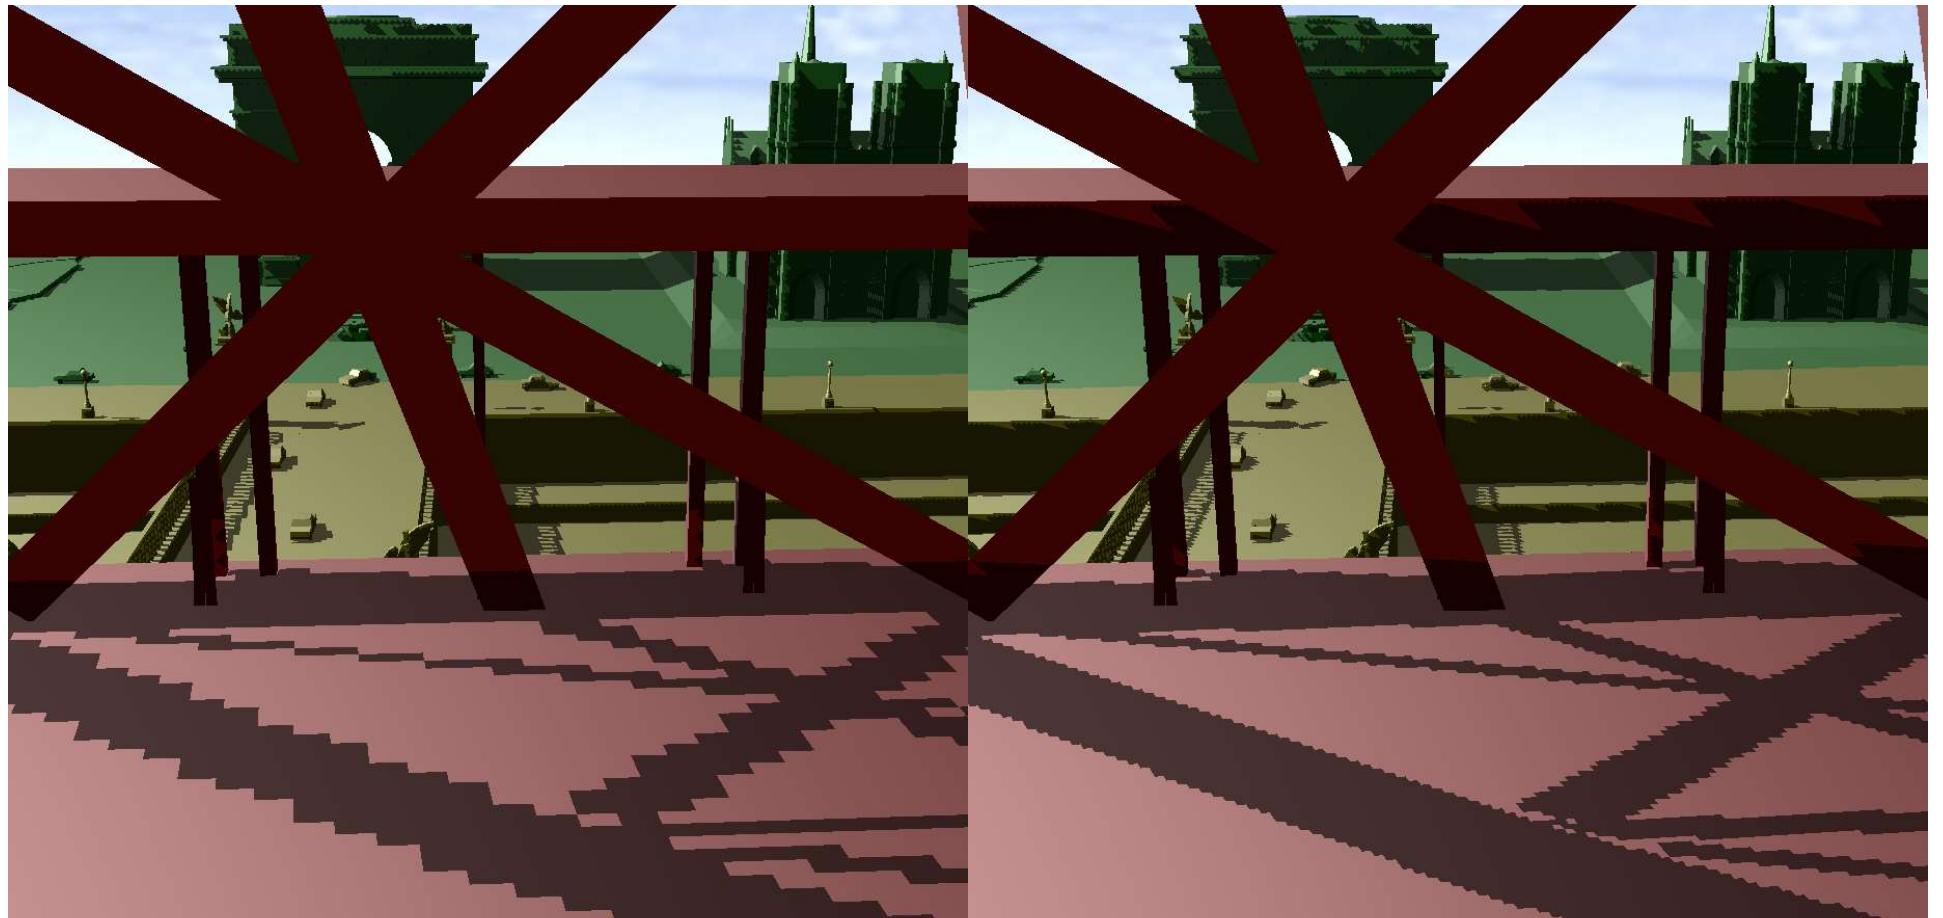
- Linear blend between the two



Z-Partitioning and Warping



Z-Partitioning and Warping

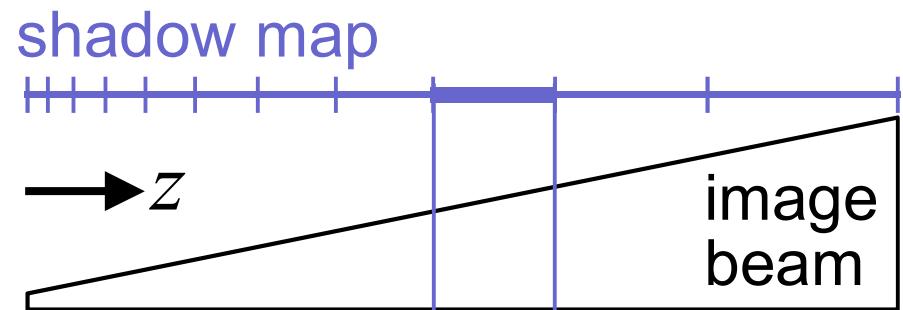


Partitioning

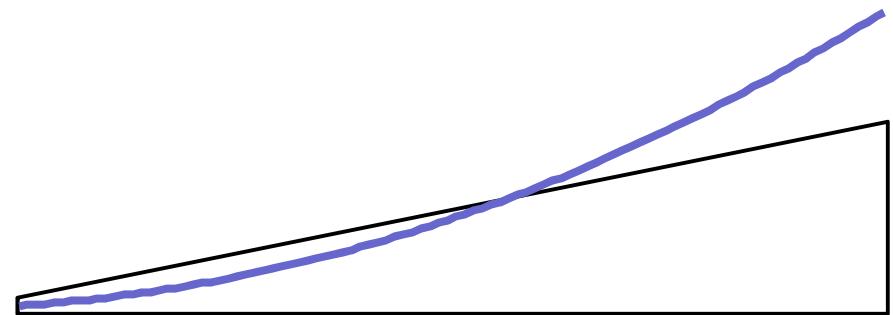
Partitioning + warping

Z-Partitioning

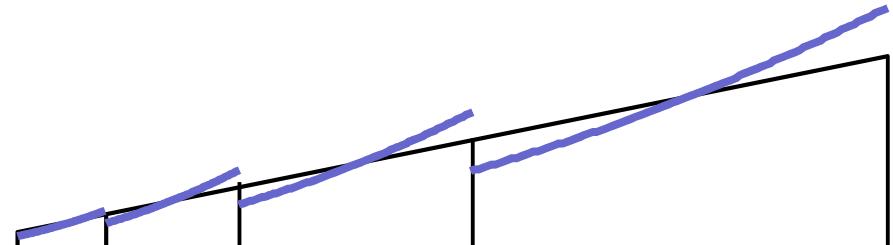
- Optimal shadow map texel spacing



- Spacing from a 4×4 projective transform

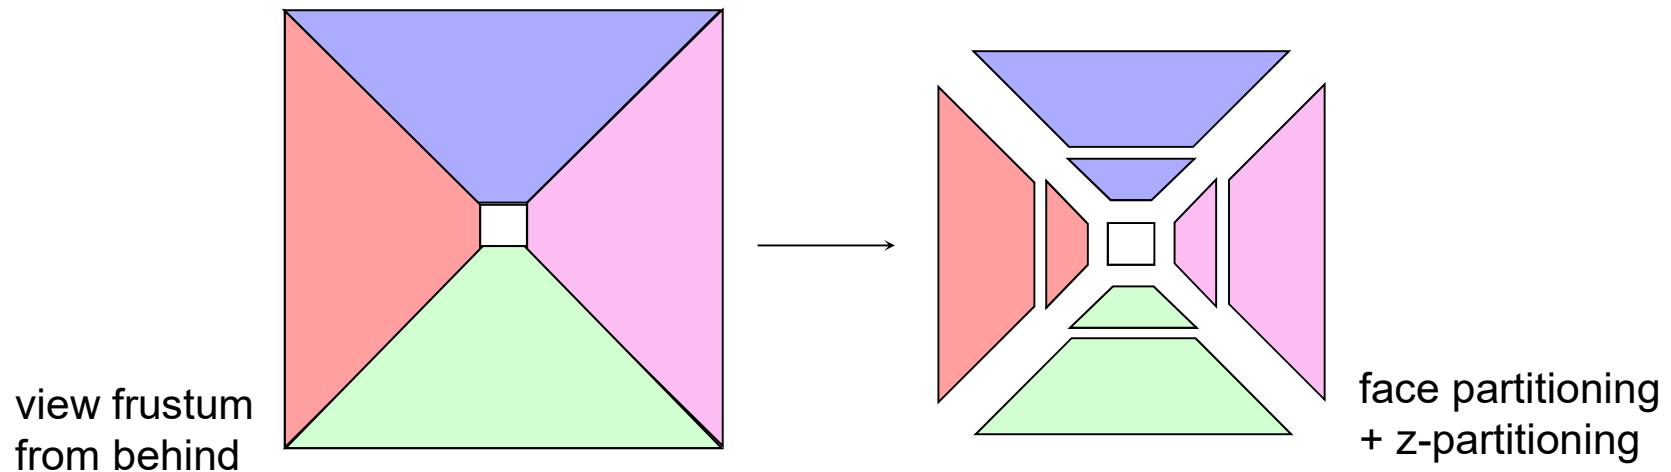


- Partitioning gives a better fit

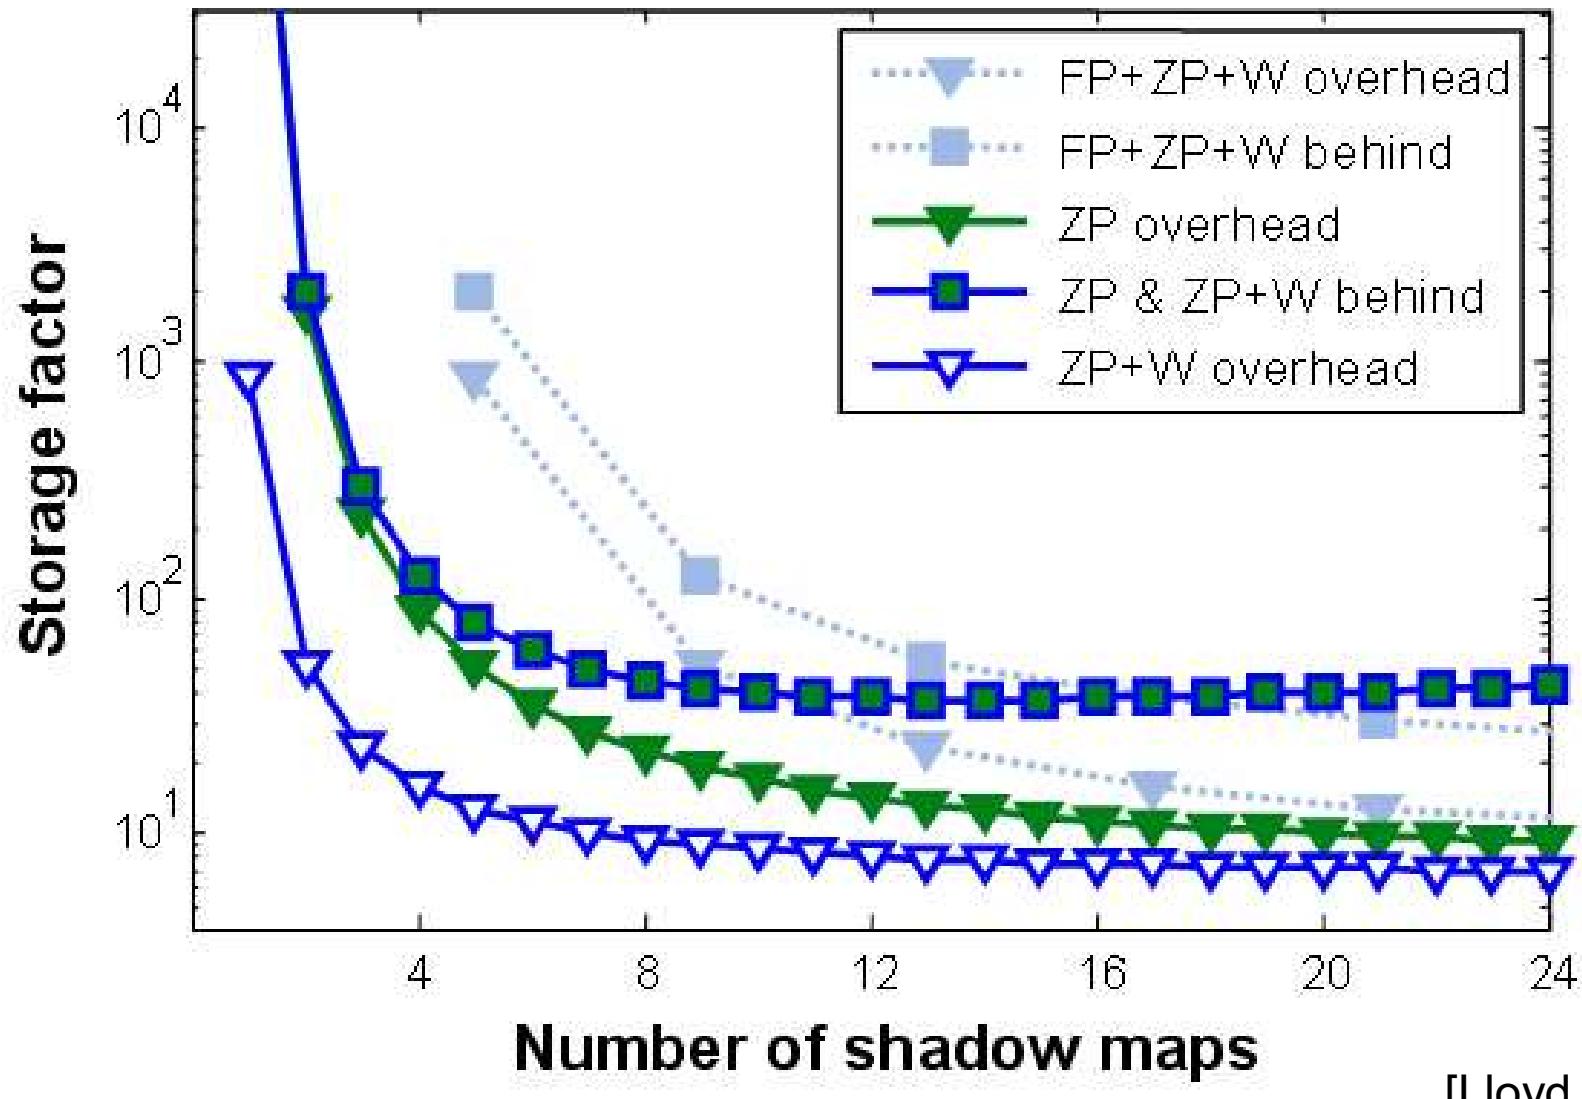


Face partitioning

- Problem: warping not good for all light directions
 - Especially from behind
- Solution: Partition frustum according to faces
- Can be combined with z-partitioning



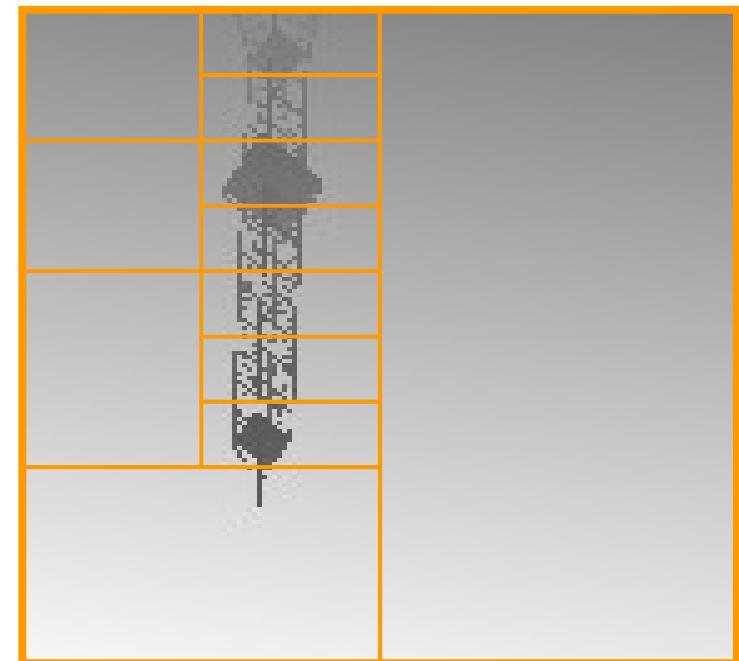
Full Error Analysis



[Lloyd 2006]

Adaptive Partitioning

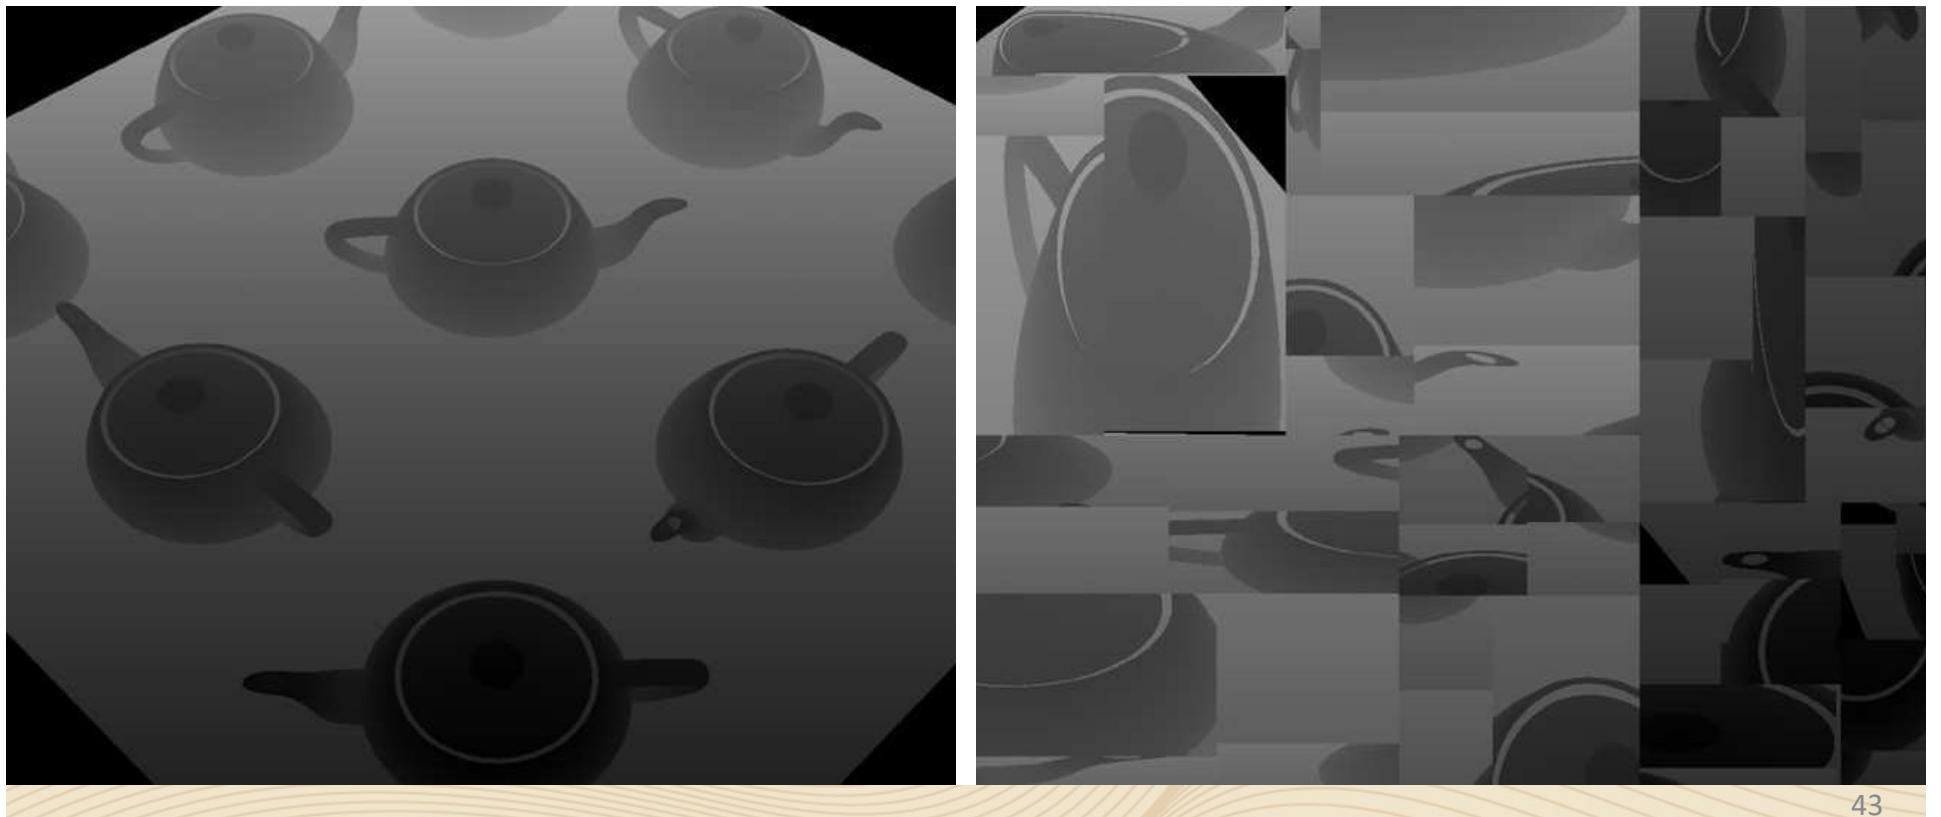
- Warping and z-partitioning are **global** resampling schemes
 - Deal with perspective aliasing
 - Projection aliasing needs local scene adaptive resampling!
- Adaptive partitioning adaptively splits shadow map
 - Usually quad-tree subdivision
 - Algorithms mainly differ in termination criteria



Tiled shadow maps

[Arvo 2004]

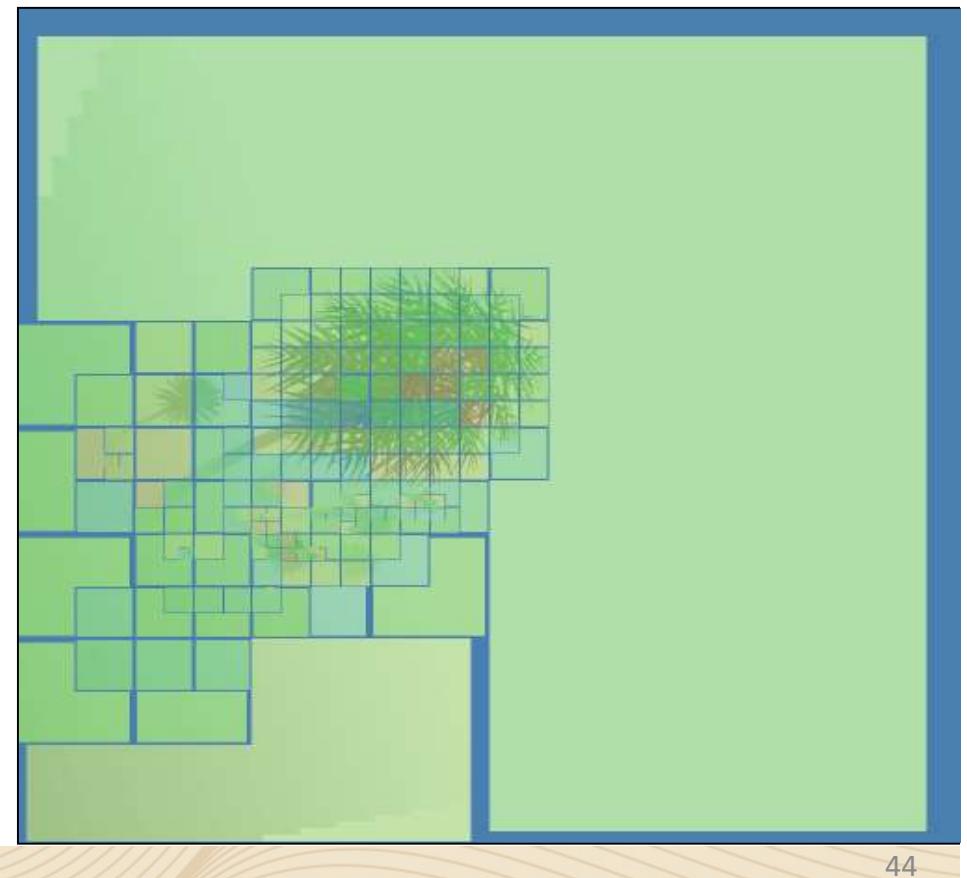
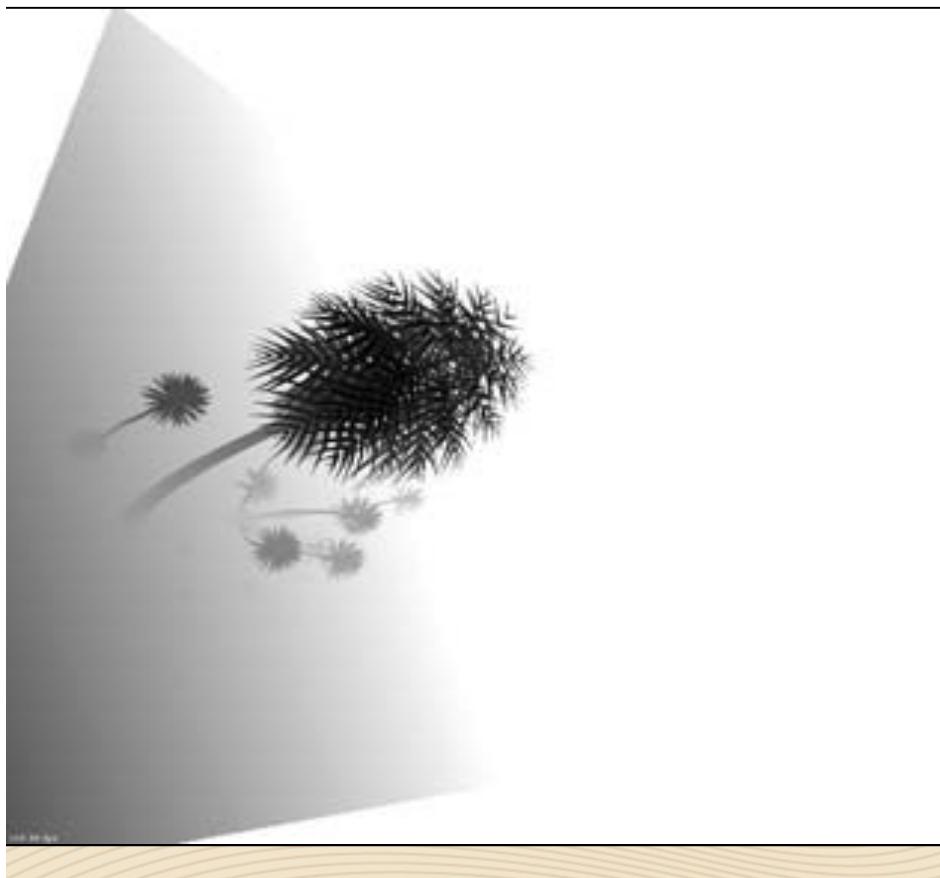
- Split according to heuristics
- Based on depth discontinuities, distances,
- Allows to trade speed against quality



Adaptive shadow maps

[Fernando et al. 2001; Lefohn et al. 2006]

- High resolution only needed at edges
- Search for edge (slow)
 - If edge split

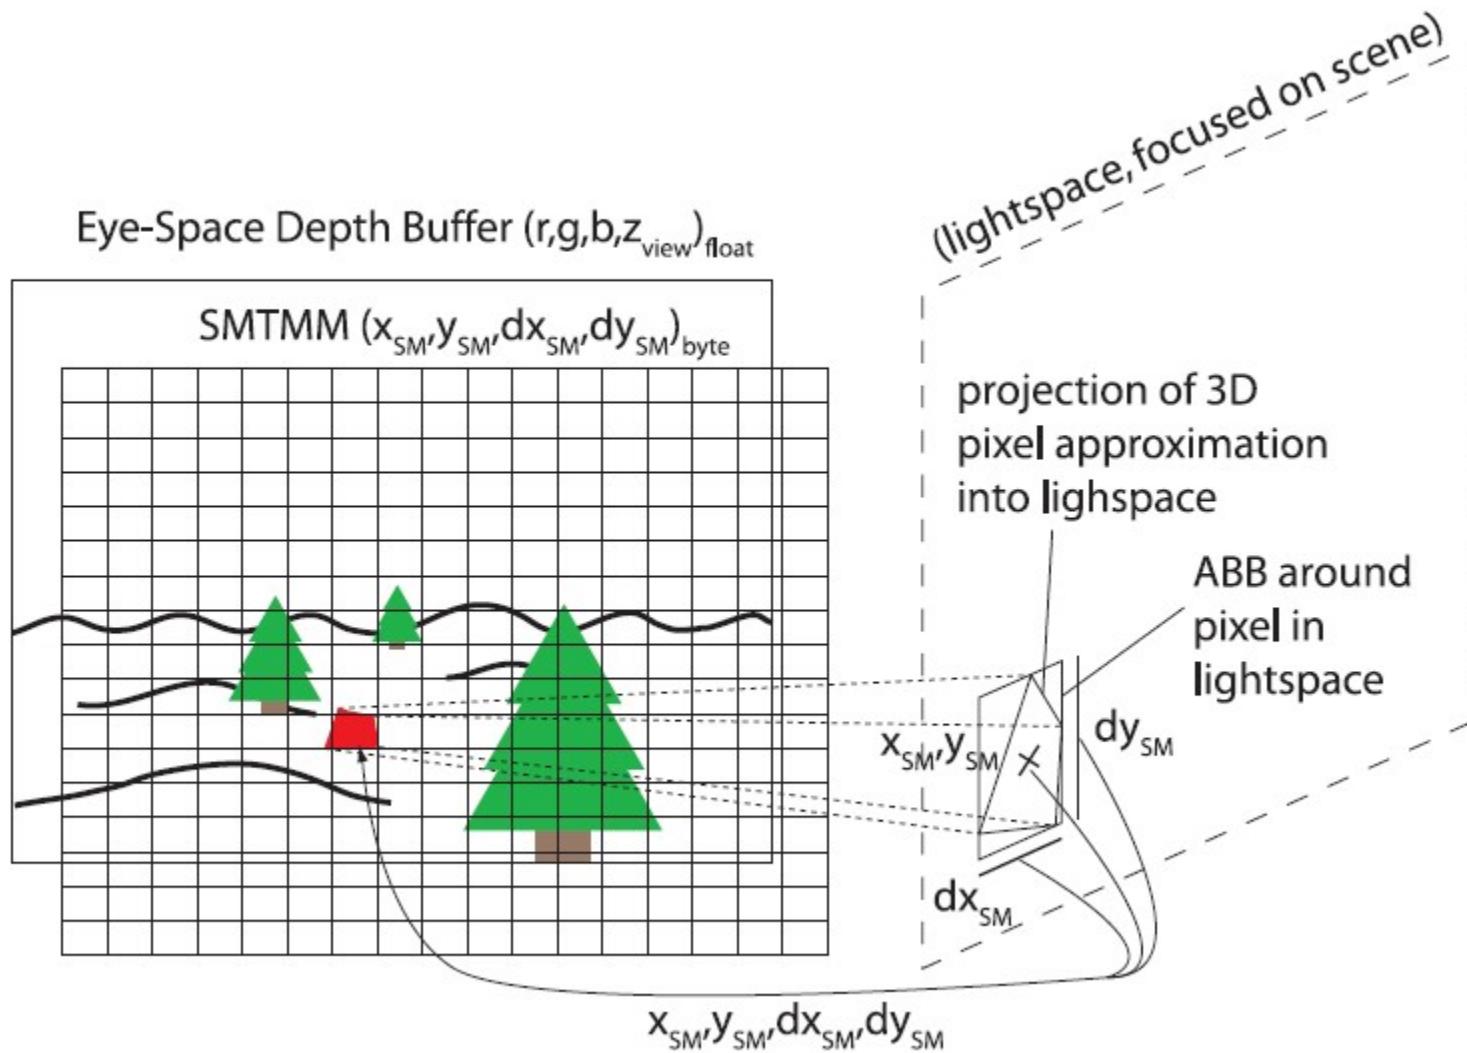


Queried/Fitted Virtual Shadow Maps

- Queried Virtual Shadow Maps [Gieg'l 2006]
 - Render 4 shadow maps of a quad-tree node
 - Check if this was necessary (occlusion query)
- Fitted Virtual Shadow Maps [Gieg'l 2007]
Resolution Matched Shadow Maps [Lefohn 2007]
 - Project pixel footprints into shadow map
 - Generate quad-tree
 - FVSM: CPU, RMSM: GPU

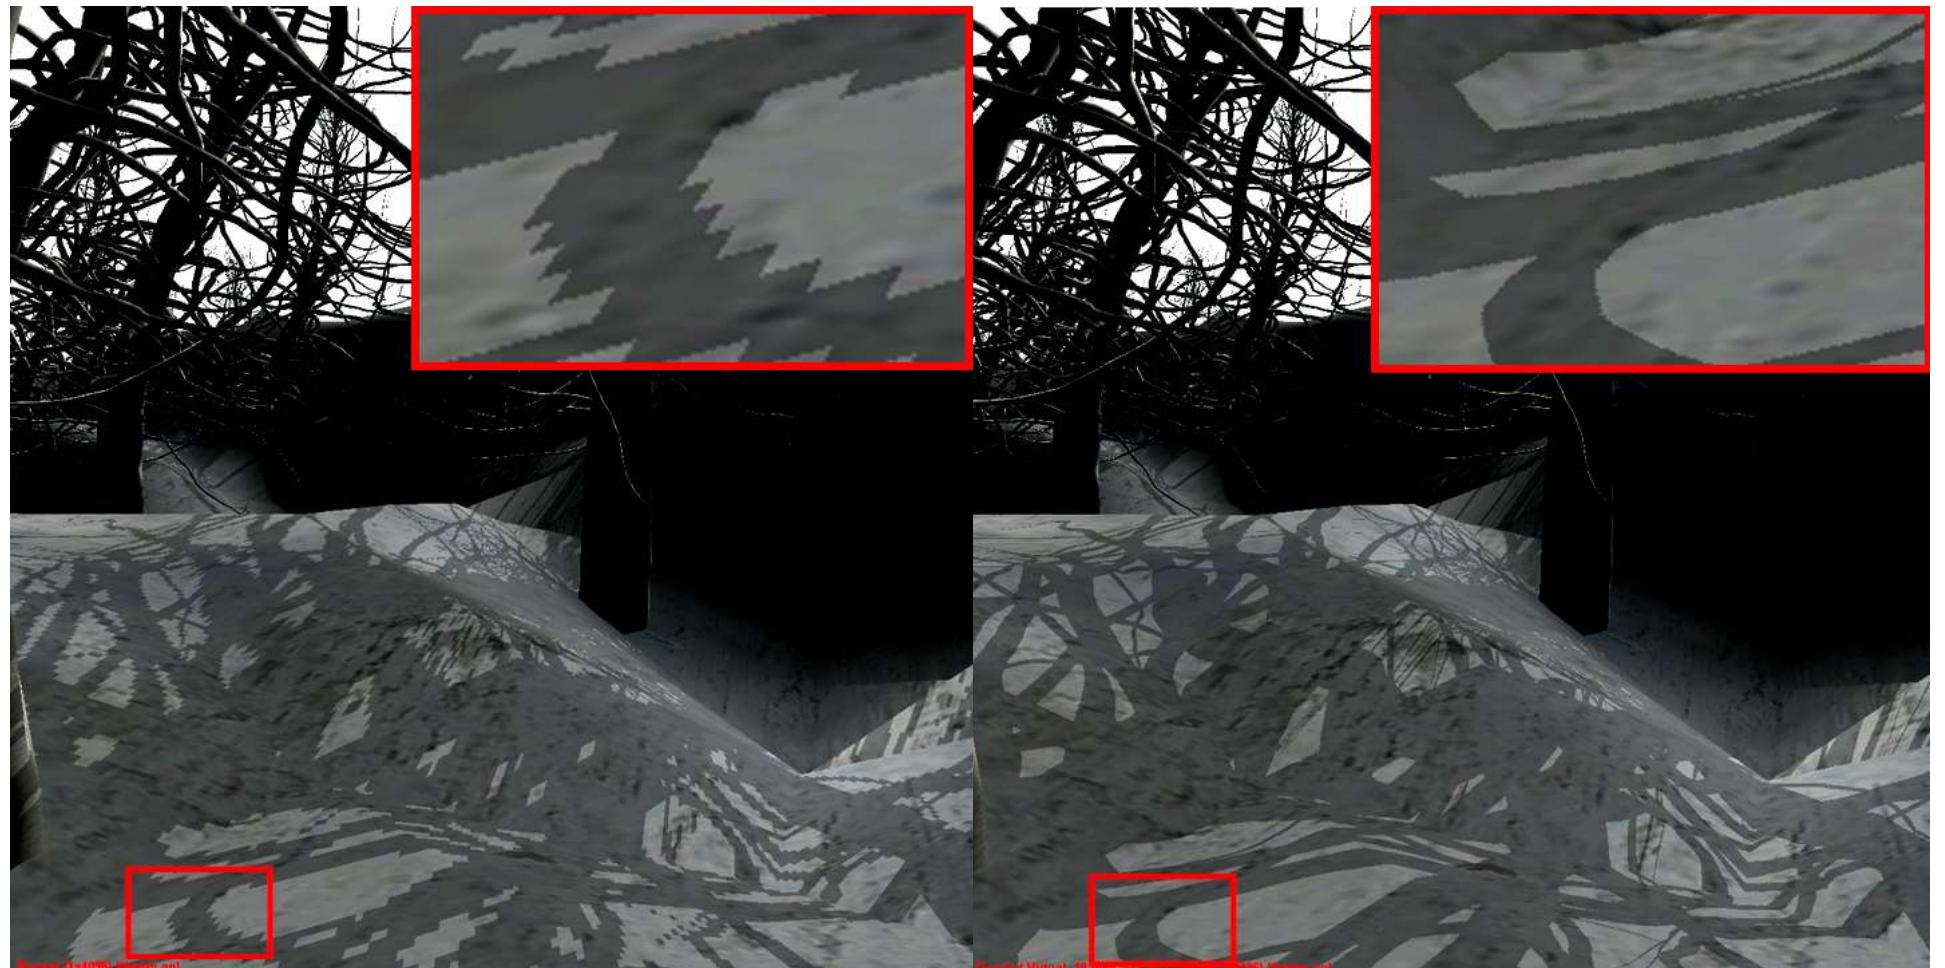
Fitted Virtual Shadow Maps

[Gieg'l & Wimmer 2007]



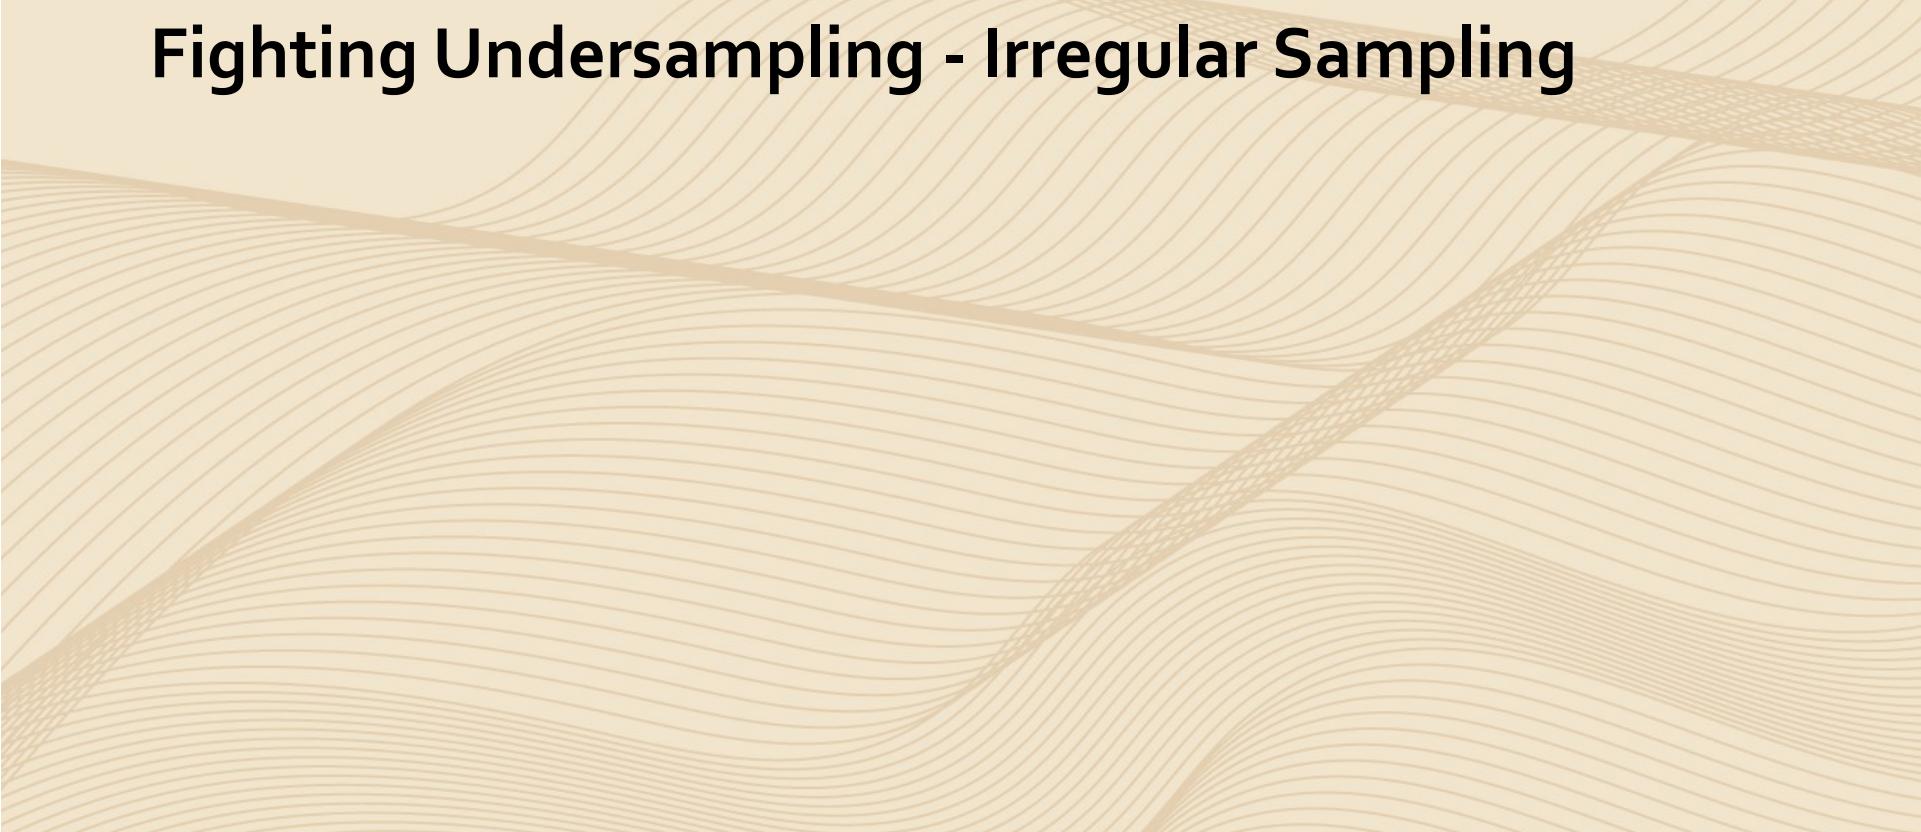
Fitted Virtual Shadow Maps

[Gieg'l & Wimmer 2007]

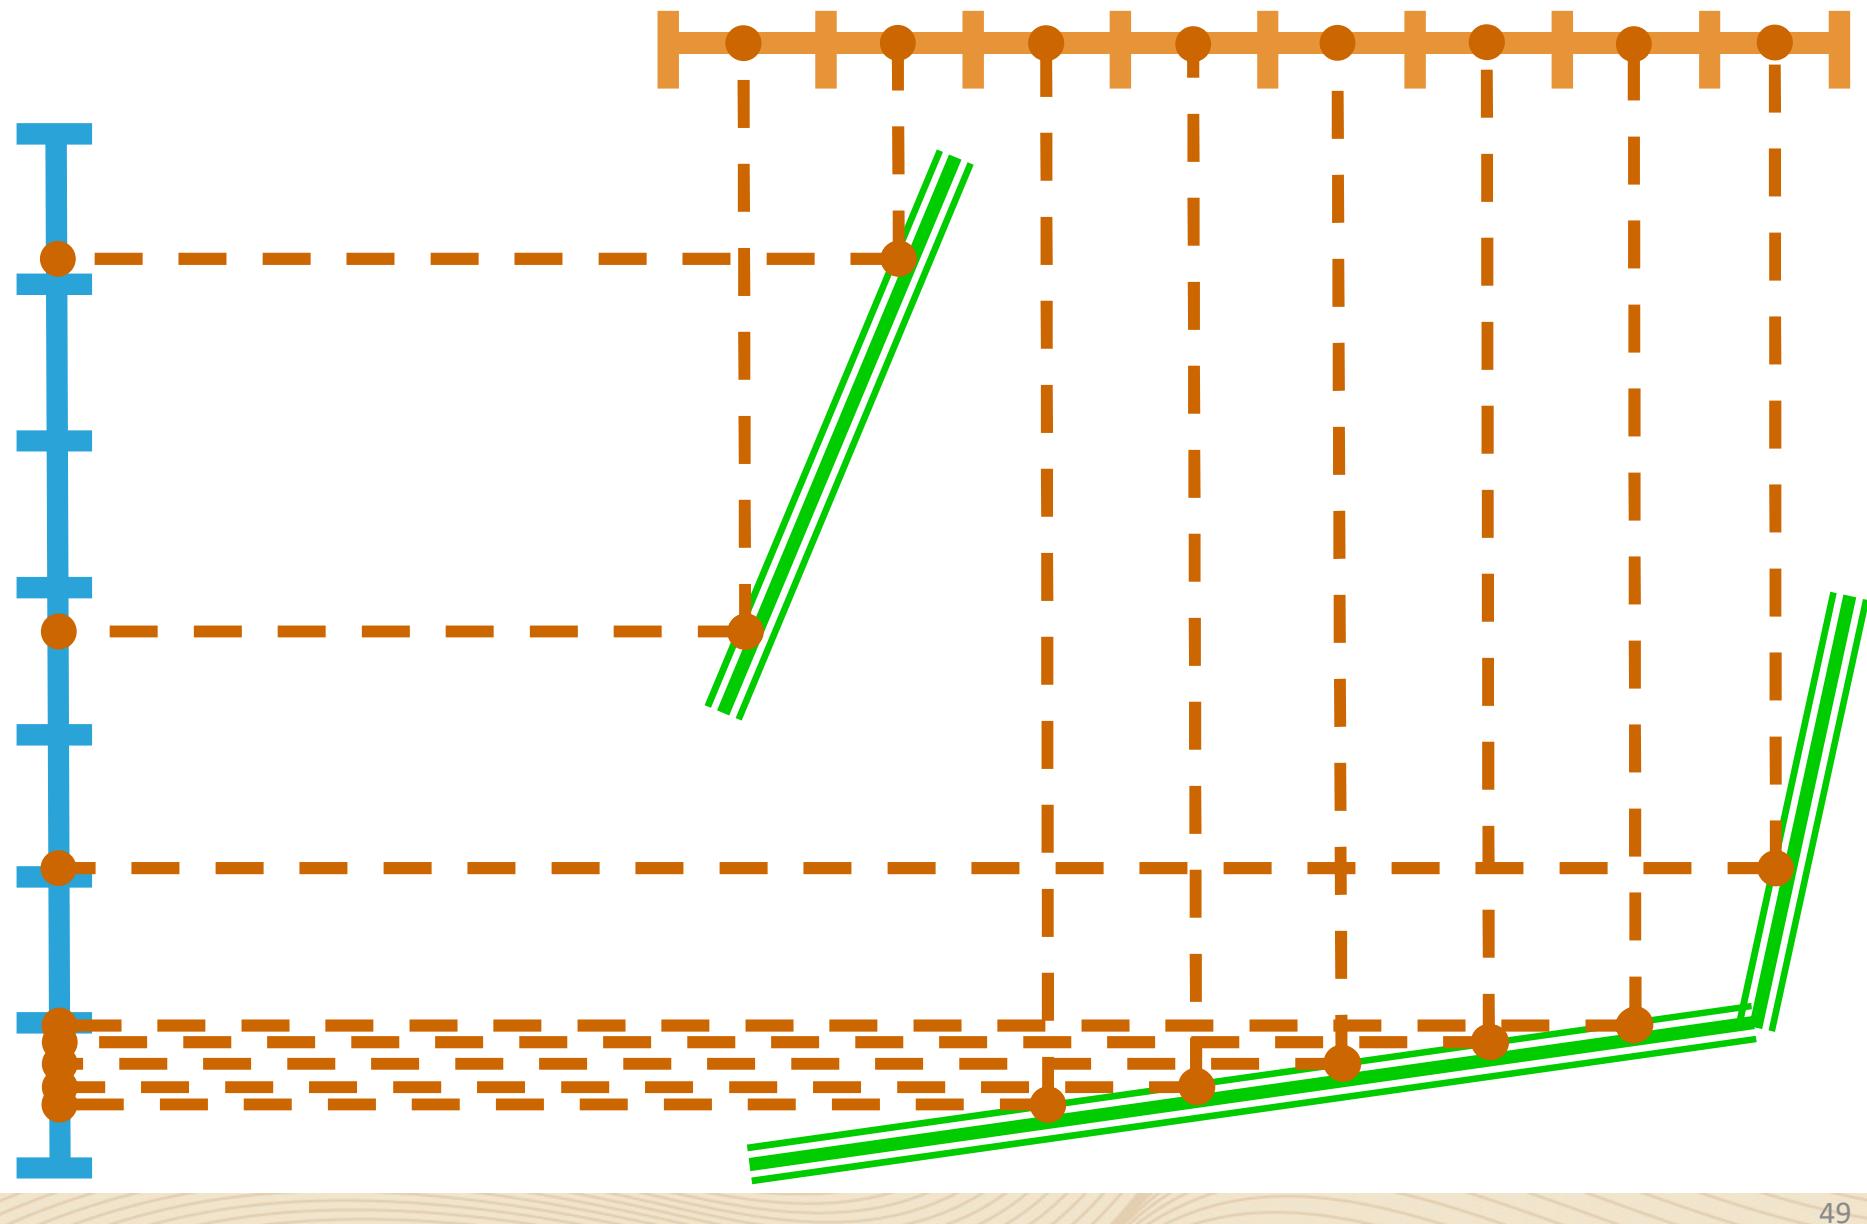


Hard Shadows

Fighting Undersampling - Irregular Sampling

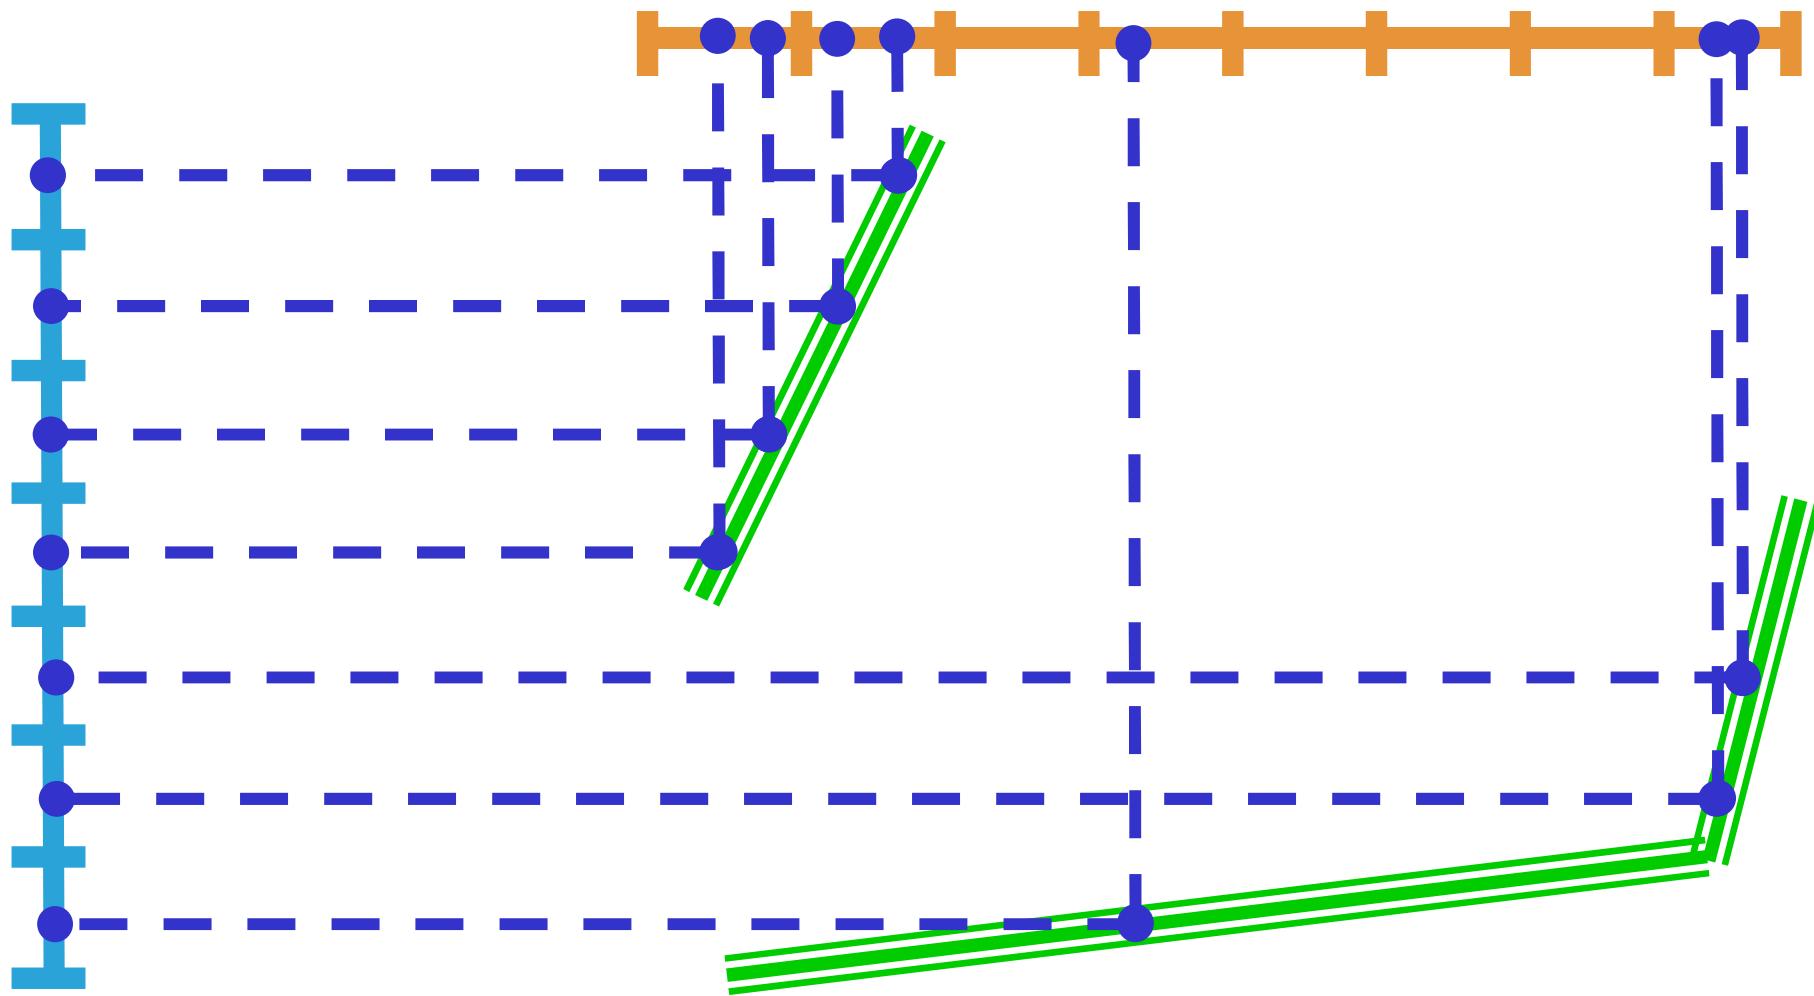


Shadow Mapping Sampling



What Samples Do We Want?

- Idea: use eye space samples to generate shadow map samples



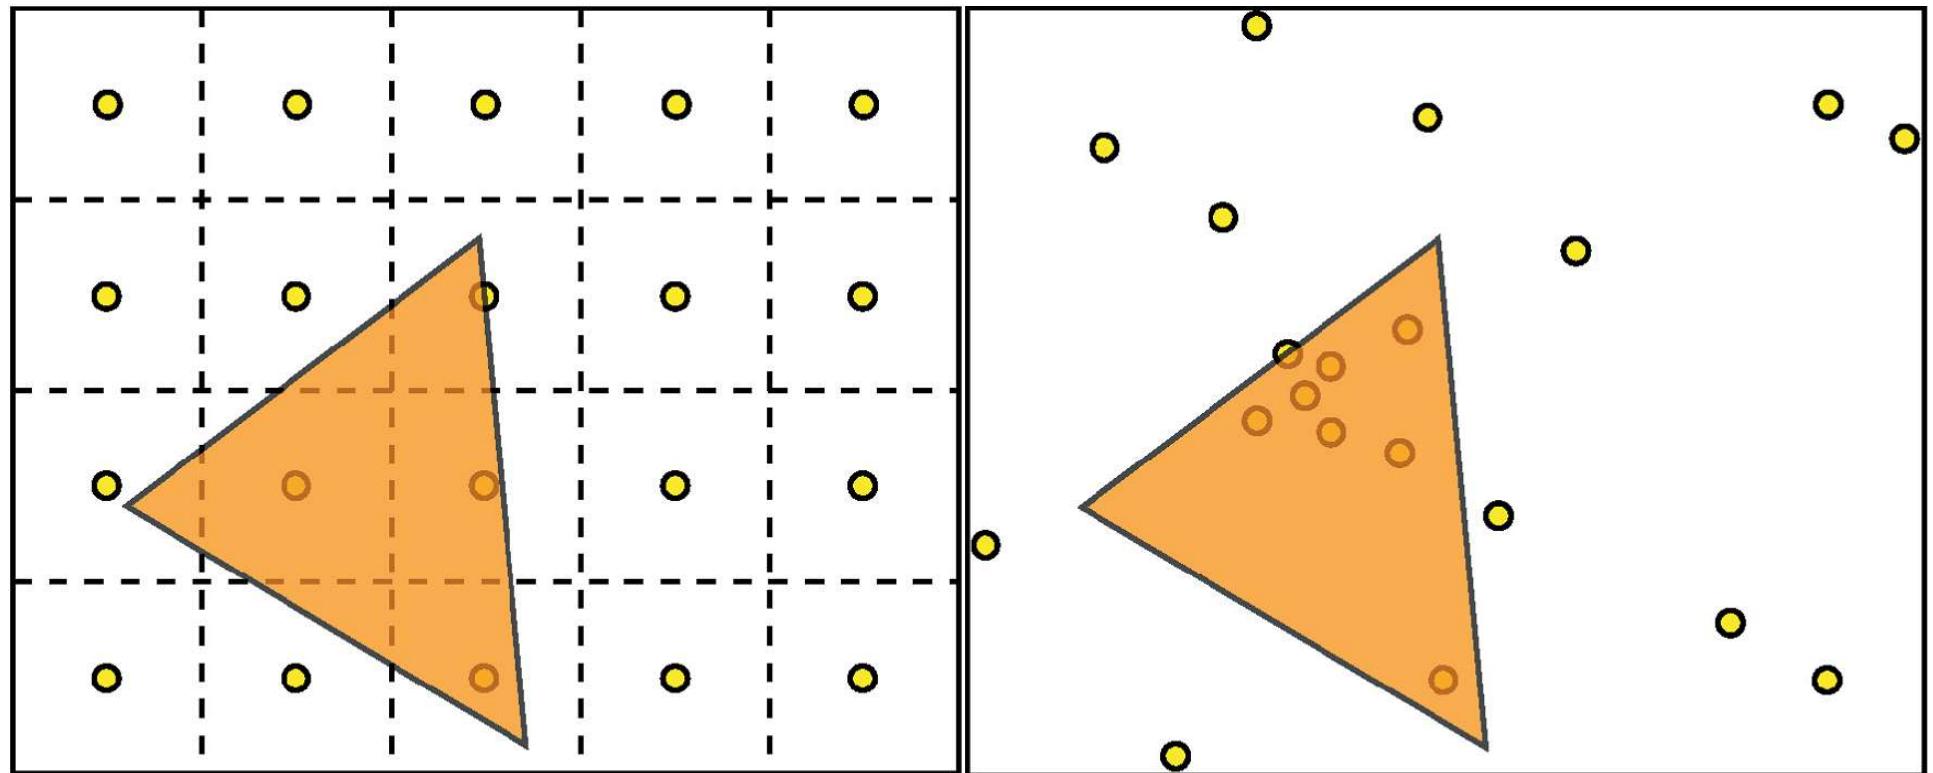
Regular vs Irregular Sampling



Irregular z-Buffer

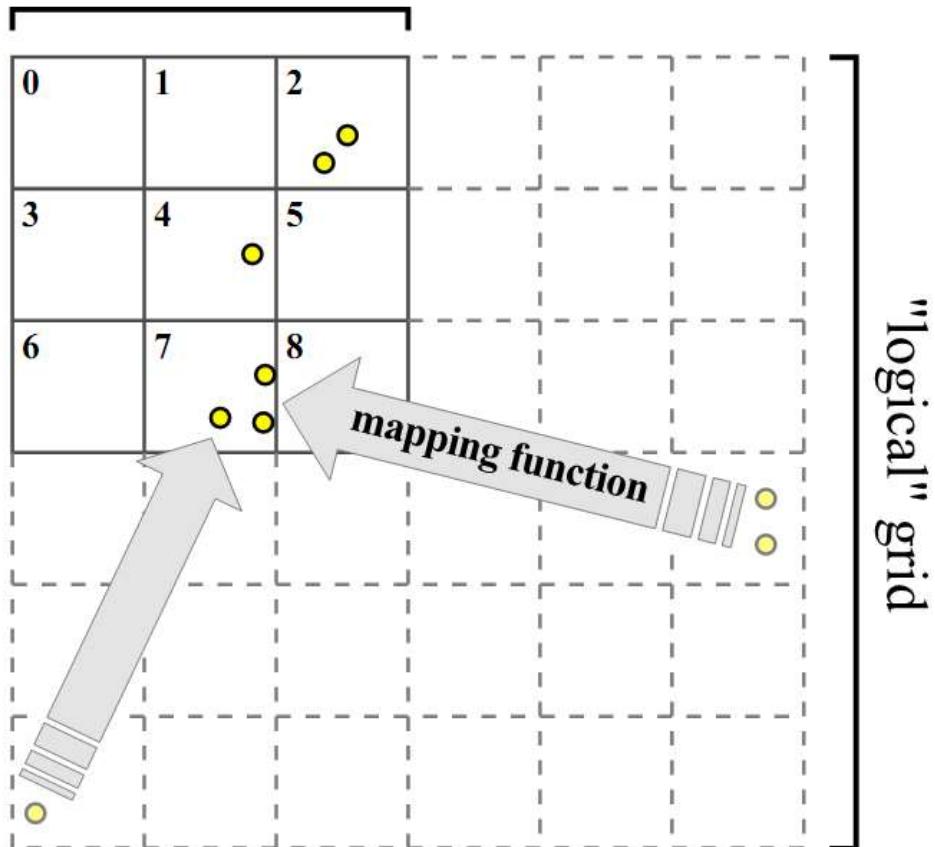
[Johnson et al. 2004]

- Proposes hardware extension

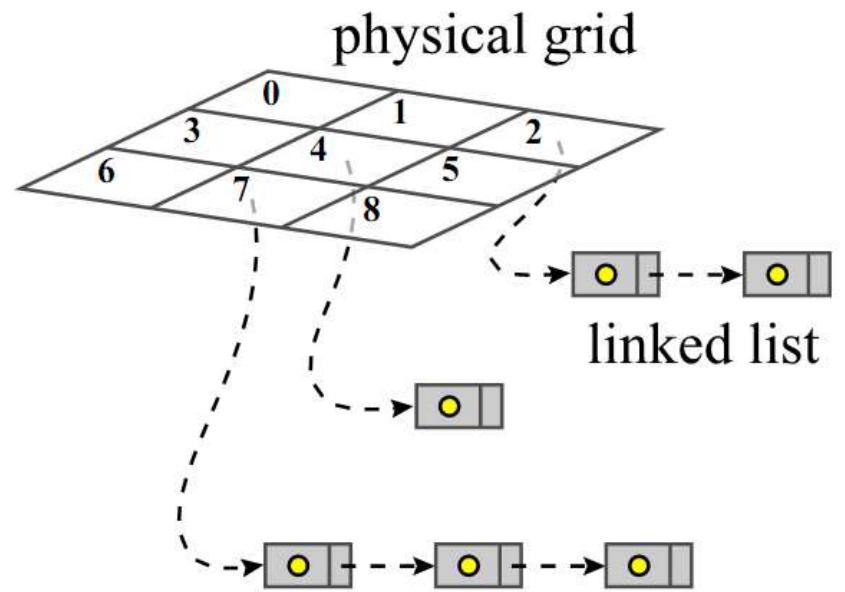


Irregular z-Buffer

"physical" grid



light view image plane



○ - shadow map sample location

Alias Free Shadow Maps

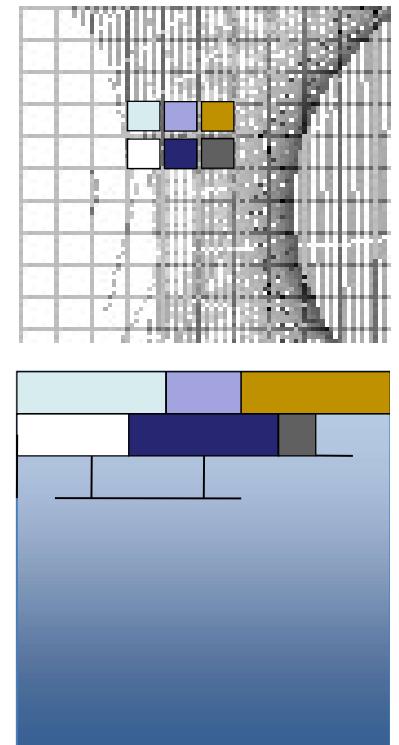
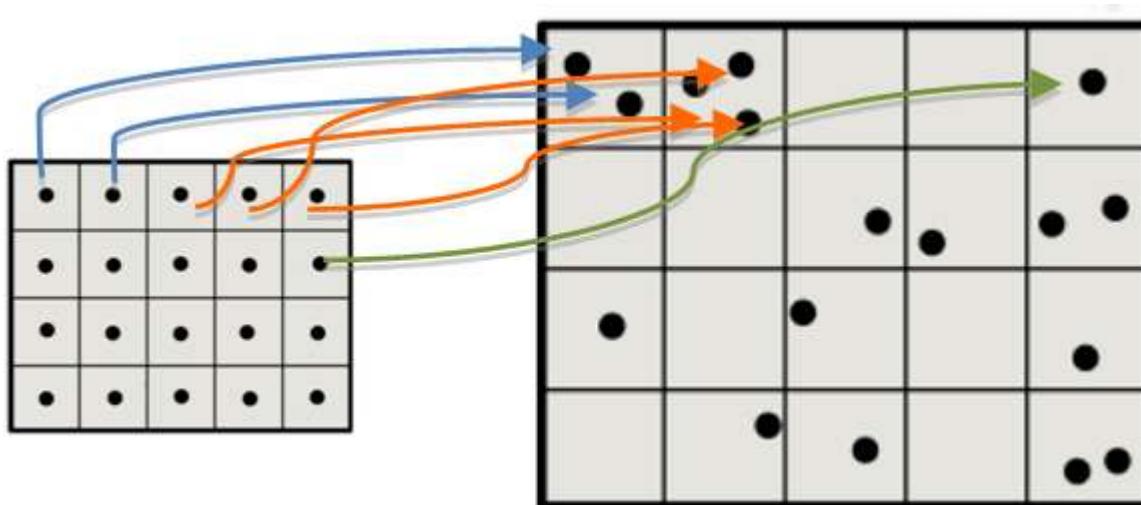
[Aila and Laine 2004]

- = Irregular Z-Buffer [Johnson 2004] in software
- Transform and project view-samples into light-space
- Rasterize blocker geometry using them as sampling points
 - Test if sampling point covered
 - Hierarchical processing of sampling points (axis-aligned 2D BSP)
 - In software

Alias Free Hard Shadows

[Sintorn et al. 2008]

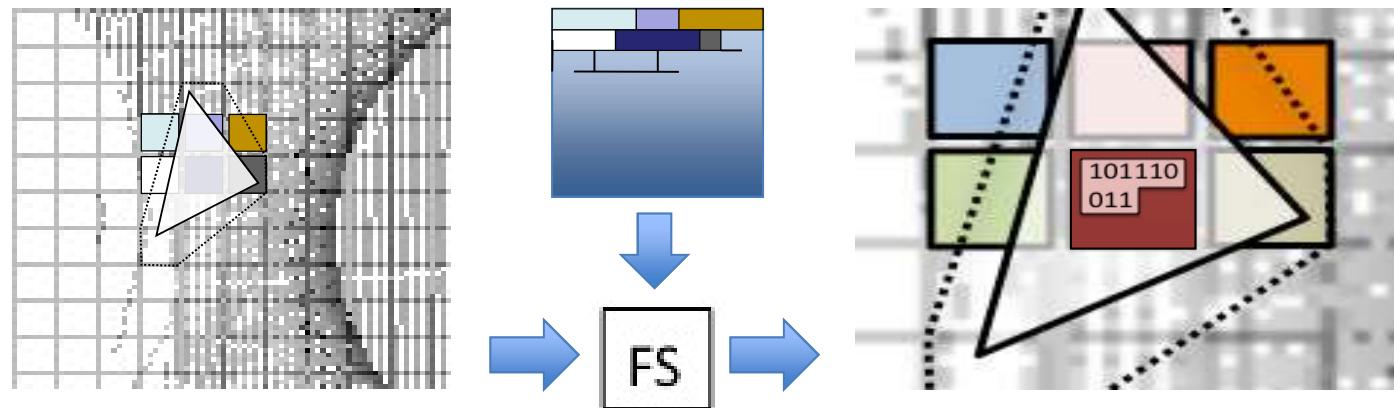
- Transform and project view-samples into light-space
- Store in a compact data structure with a list per light-space texel



Alias Free Hard Shadows

[Sintorn et al. 2008]

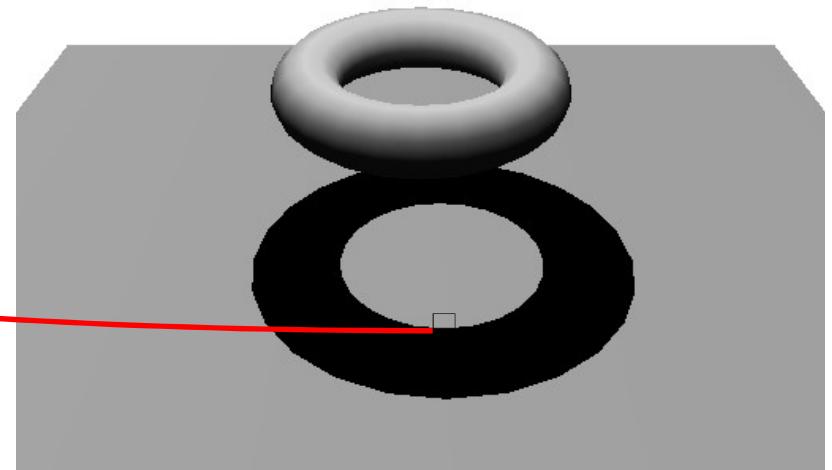
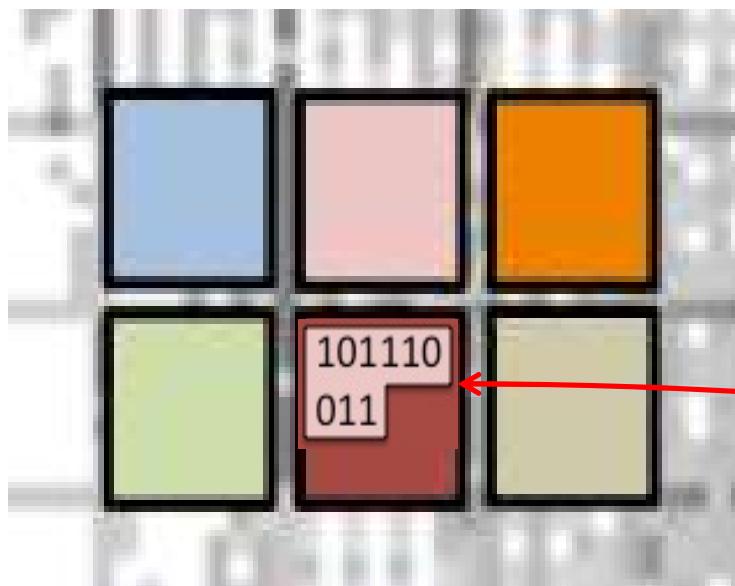
- Render all geometry (conservatively) from lights point of view
- For each generated fragment, test all view-samples in list against triangle
- Set corresponding output bit



Alias Free Hard Shadows

[Sintorn et al. 2008]

- Final screen-space pass
- Use bitmask from previous pass for shadowing



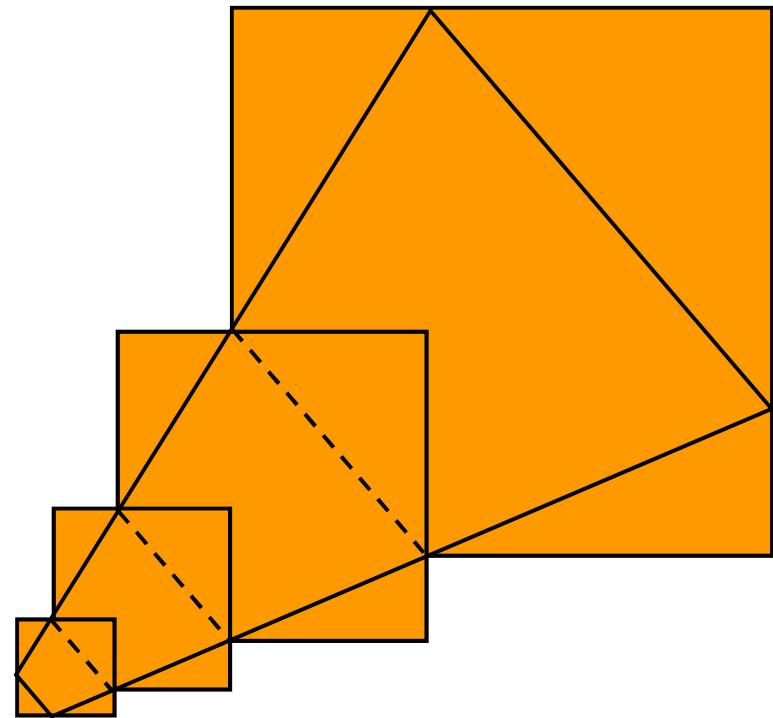
Hard Shadows

Fighting Undersampling - Temporal Reprojection



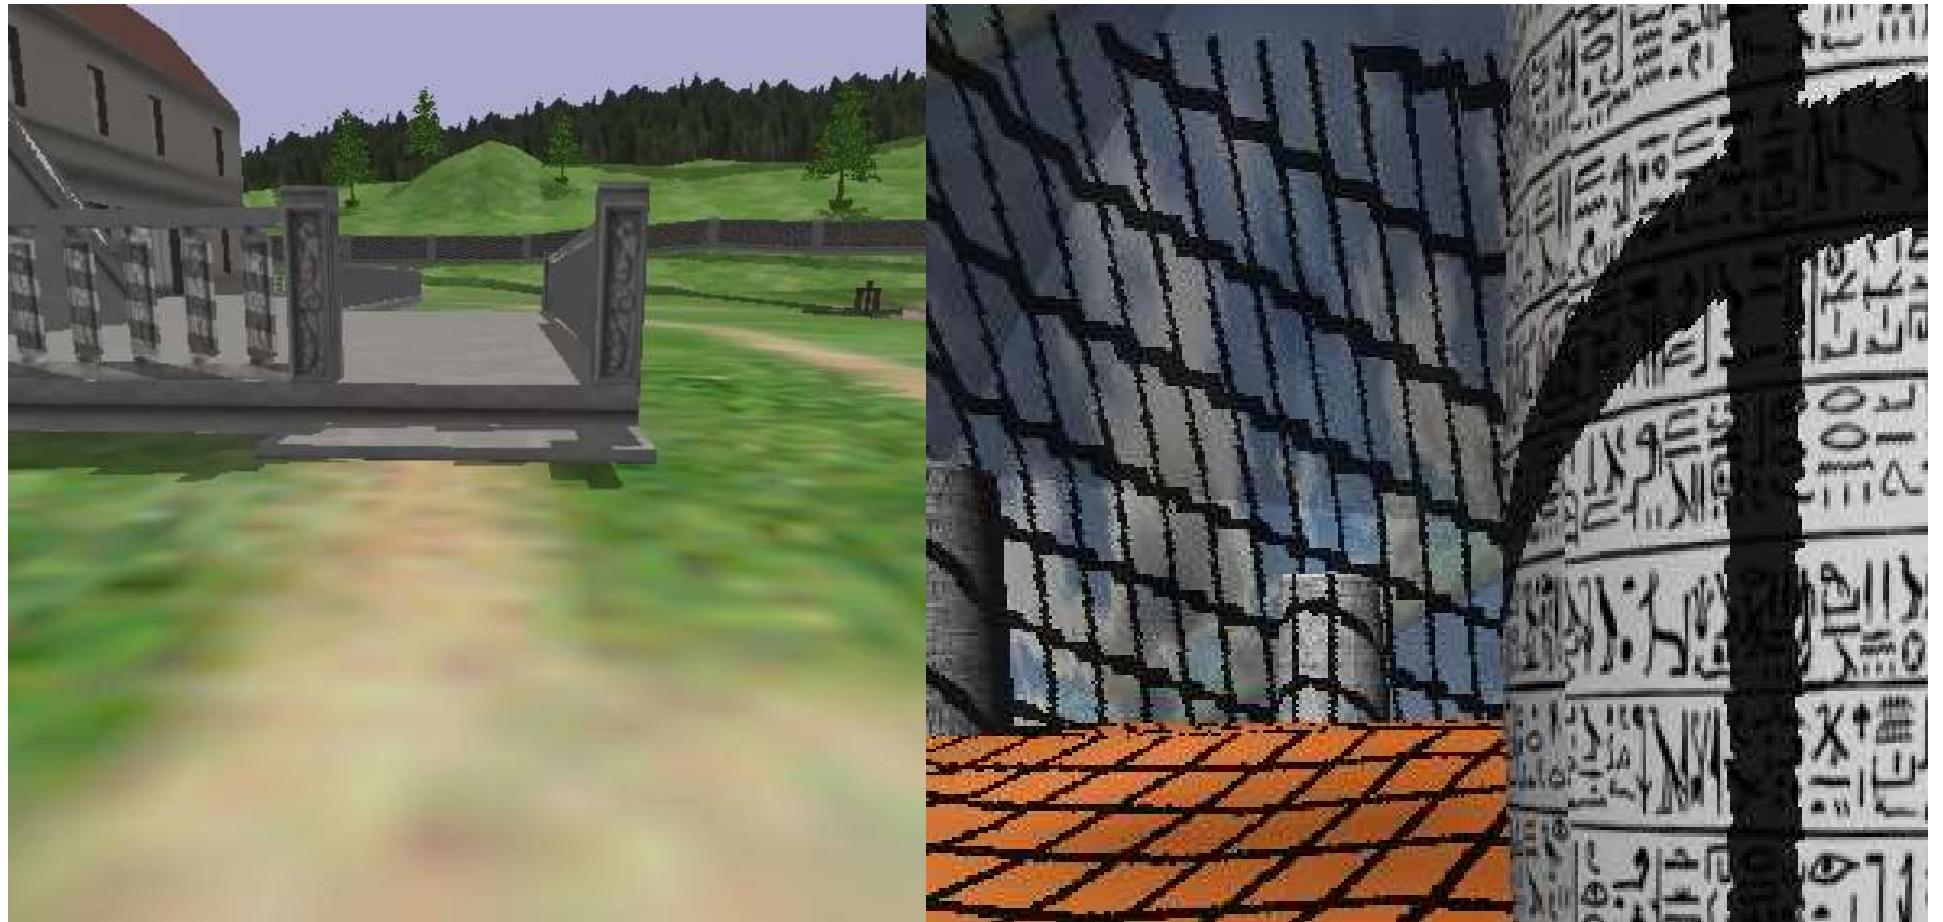
Fighting Temporal Aliasing

- Quick fix: texel snapping
 - Fix coordinate system in world space
 - Shadow maps move at integral multiples of a texel width
 - Gives up warping
 - Popping when zooming



Motivation

- Temporal aliasing
→ temporal smoothing
- Undersampling
→ shadow test confidence



Temporal Smoothing

- Shadowing result of previous frame is stored in *cache*



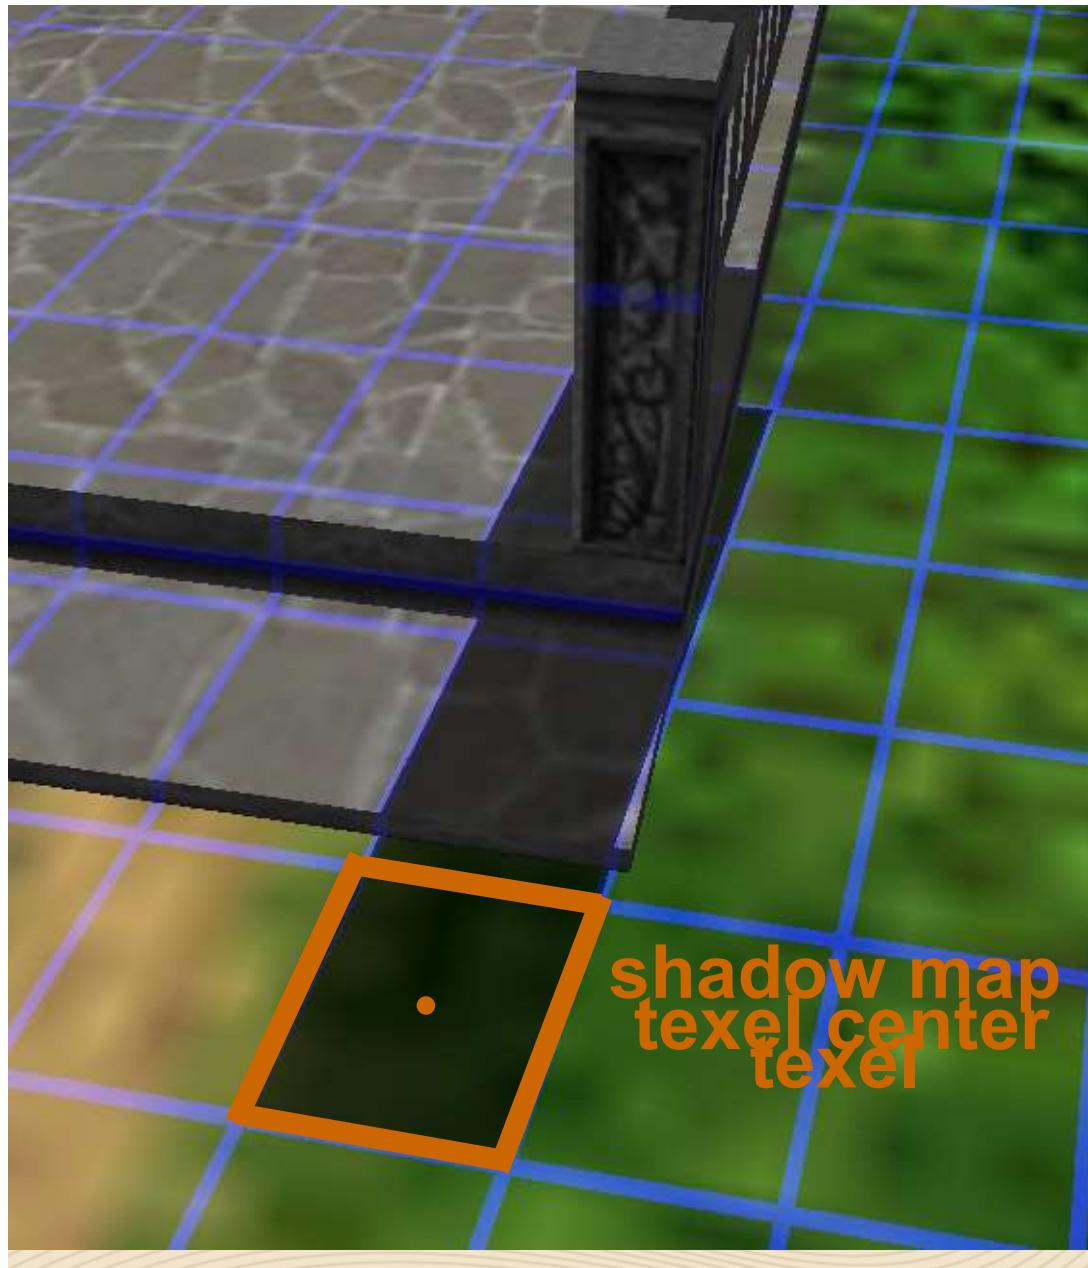
Temporal Smoothing

- Shadowing result for each pixel of previous frame ($n-1$) is stored in *cache*
- Shadowing result of current frame (n)
 $cache(n) := w * s(n) + (1-w) * cache(n-1)$
 - $s(n)$new shadow map test result, binary
 - $cache(n-1)$..shadow results of previous frame
 - $cache(n)$shadow results of current frame
 - wweight, impact of current vs. previous
 - Large w - fast update, weak smoothing
 - Small w - slow update, strong smoothing

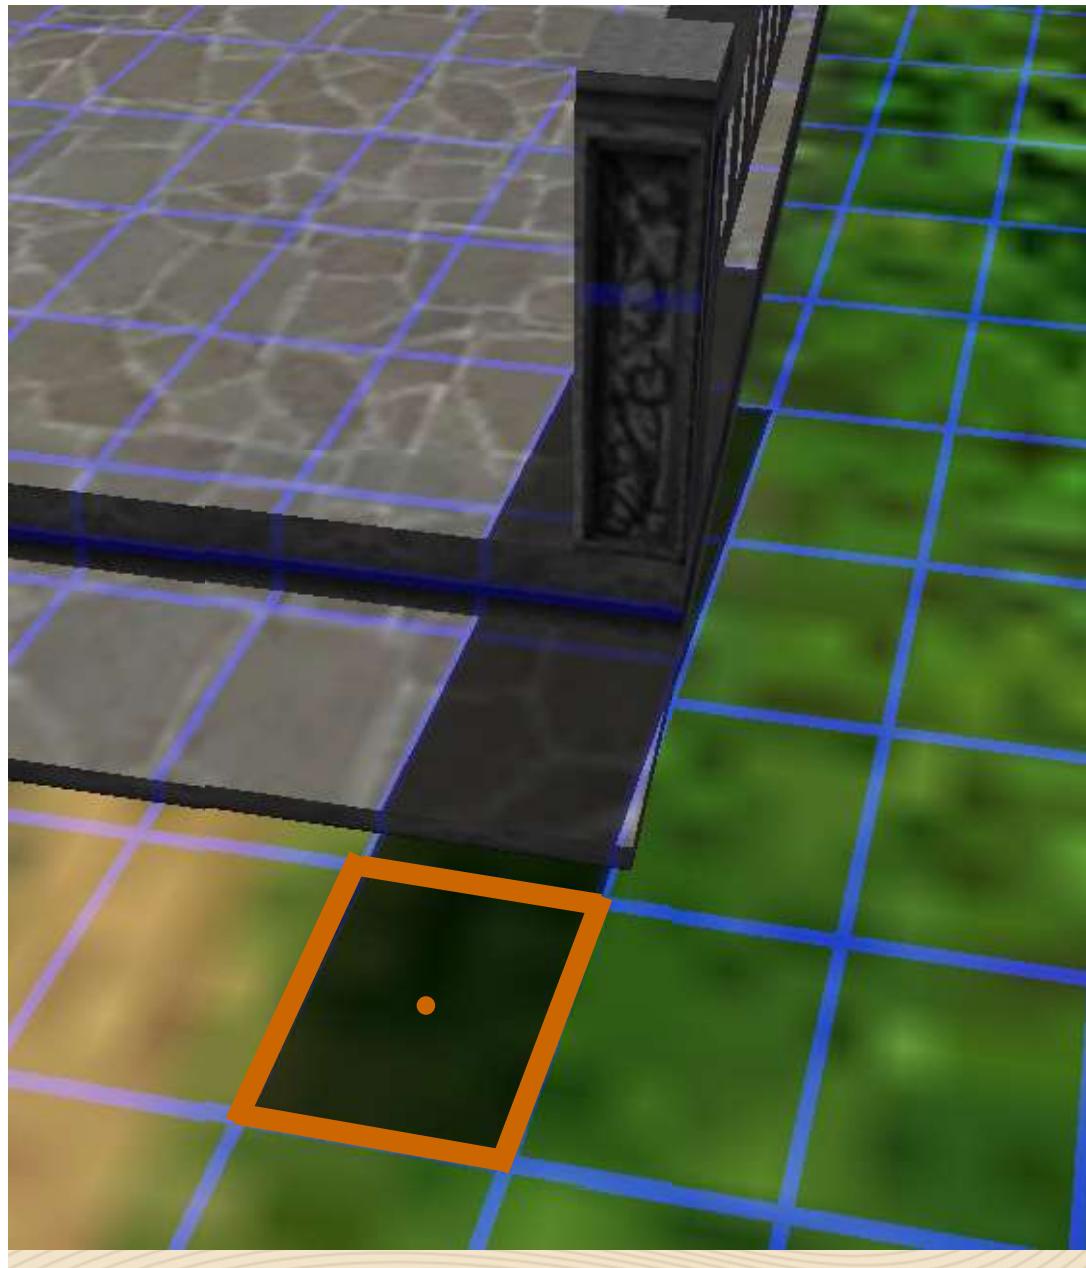
Temporal Smoothing



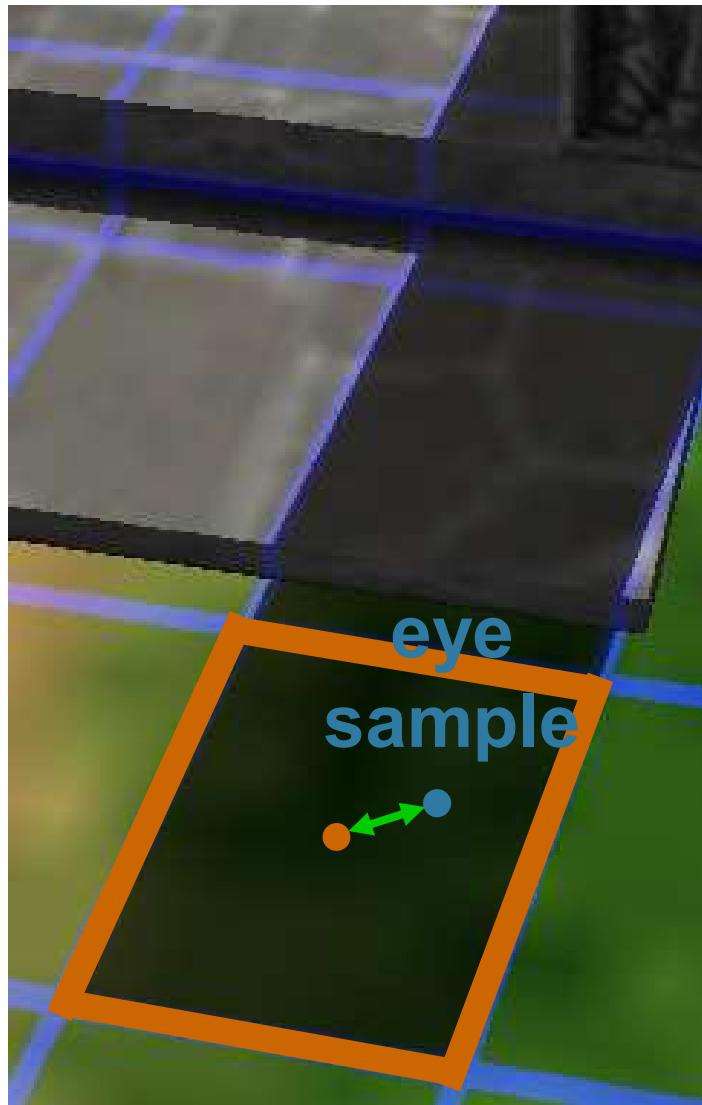
Confidence Estimation



Confidence Estimation



Confidence Estimation

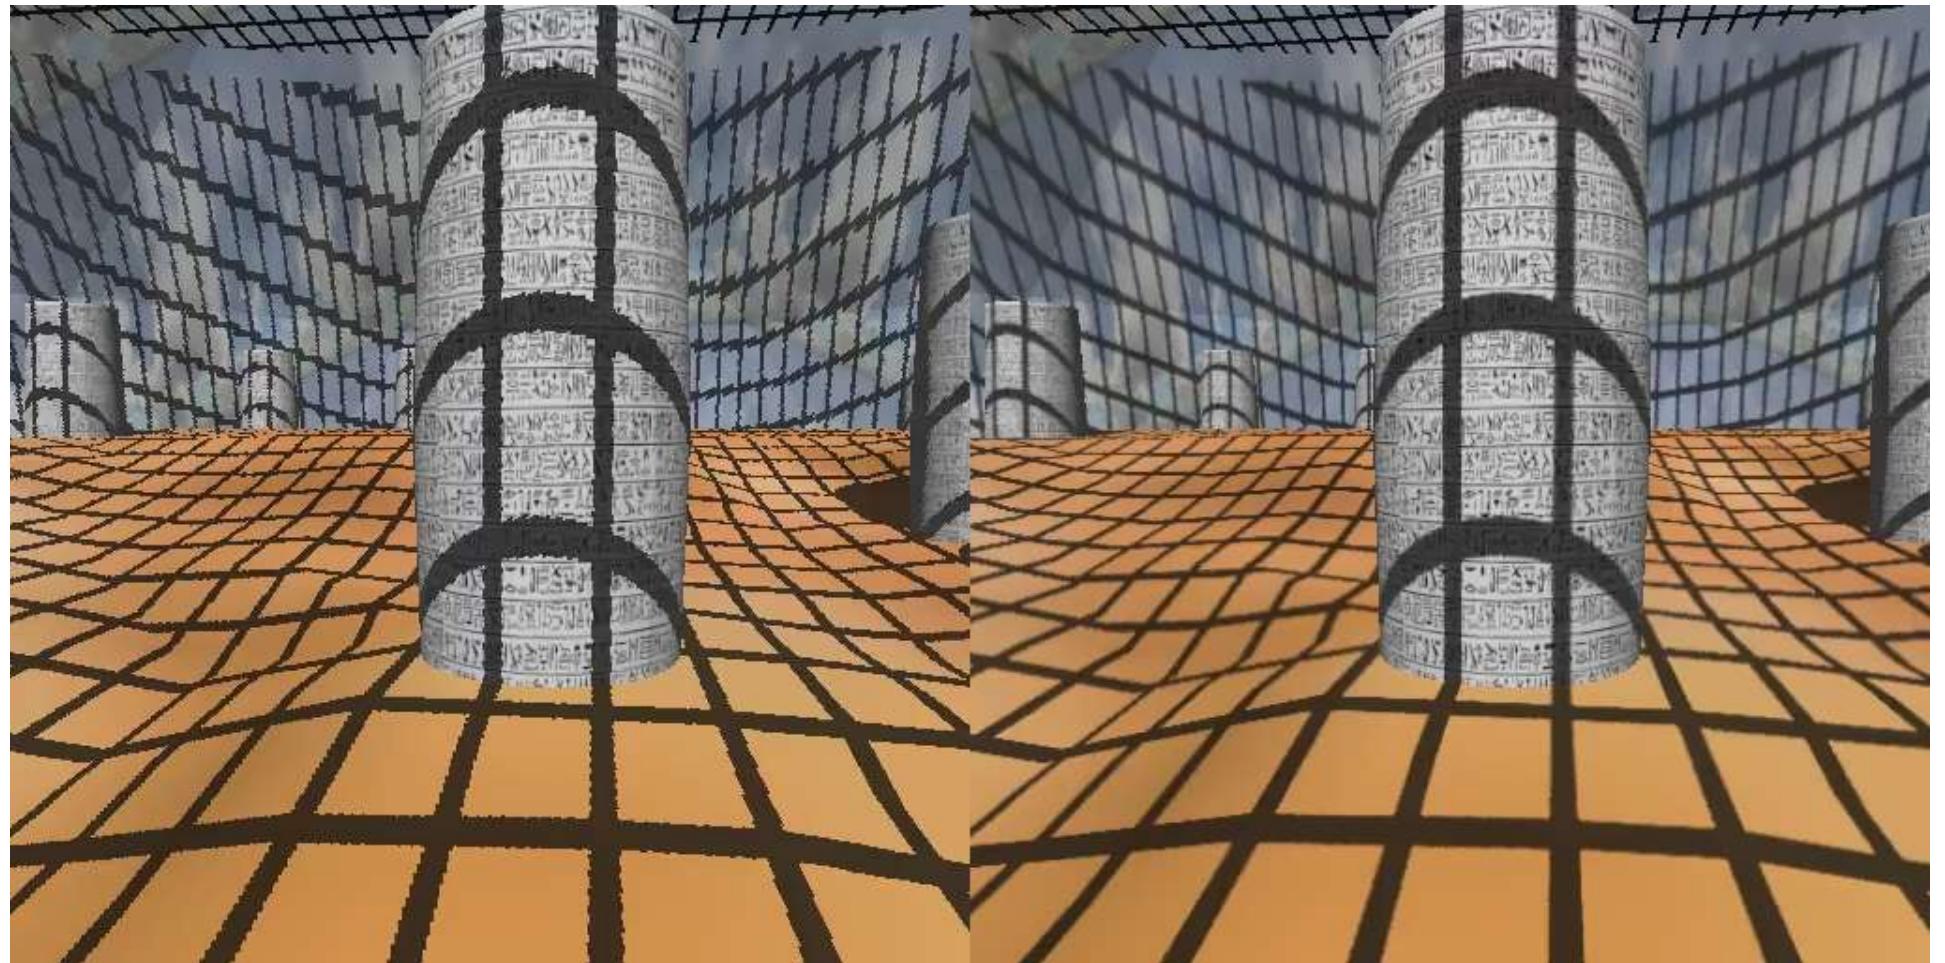


- Shadow map test probably more correct closer to texel center
- Confidence:
 $\text{conf} := 1 - \text{dist}(\text{eye}, \text{texel})$
- Greater confidence
→ greater impact

Confidence and Temporal Smoothing

- Confidence is weight for temporal smoothing

$$cache(n) := \text{conf} * s(n) + (1 - \text{conf}) * cache(n-1)$$



Hard Shadows

Fighting Undersampling - Conclusion



Conclusions

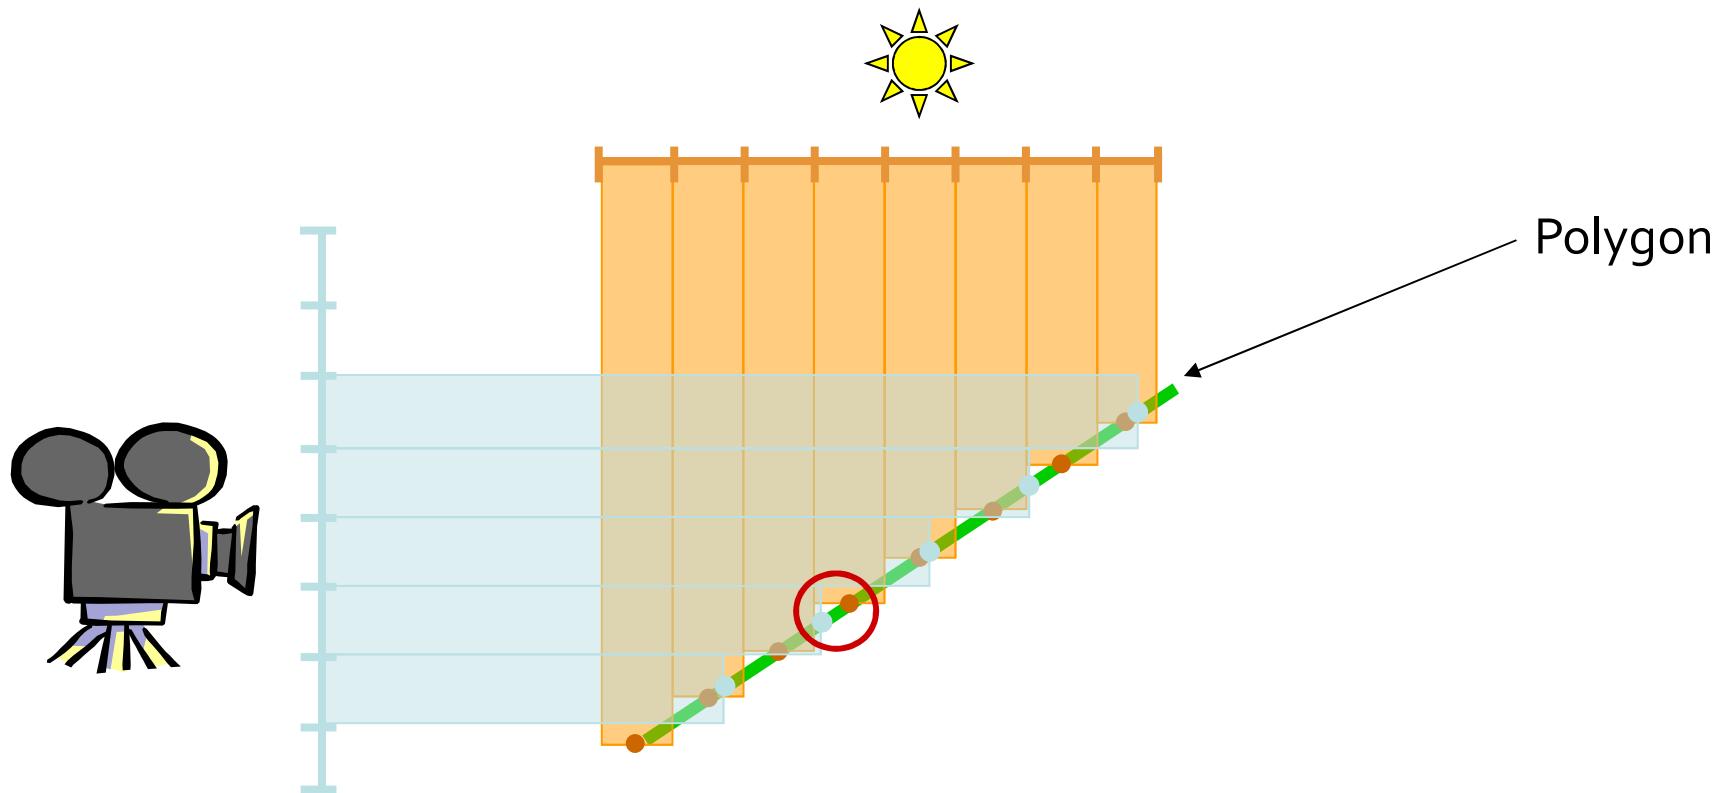
- Fastest speed, single shadow map: **warping**
 - Good for outdoor
- Fast speed, better quality, multiple shadow maps: **z-partitioning**
- High quality, lower speed: **adaptive partitioning**
- Reference quality, even slower: **irregular sampling**

Hard Shadows

Depth Biasing

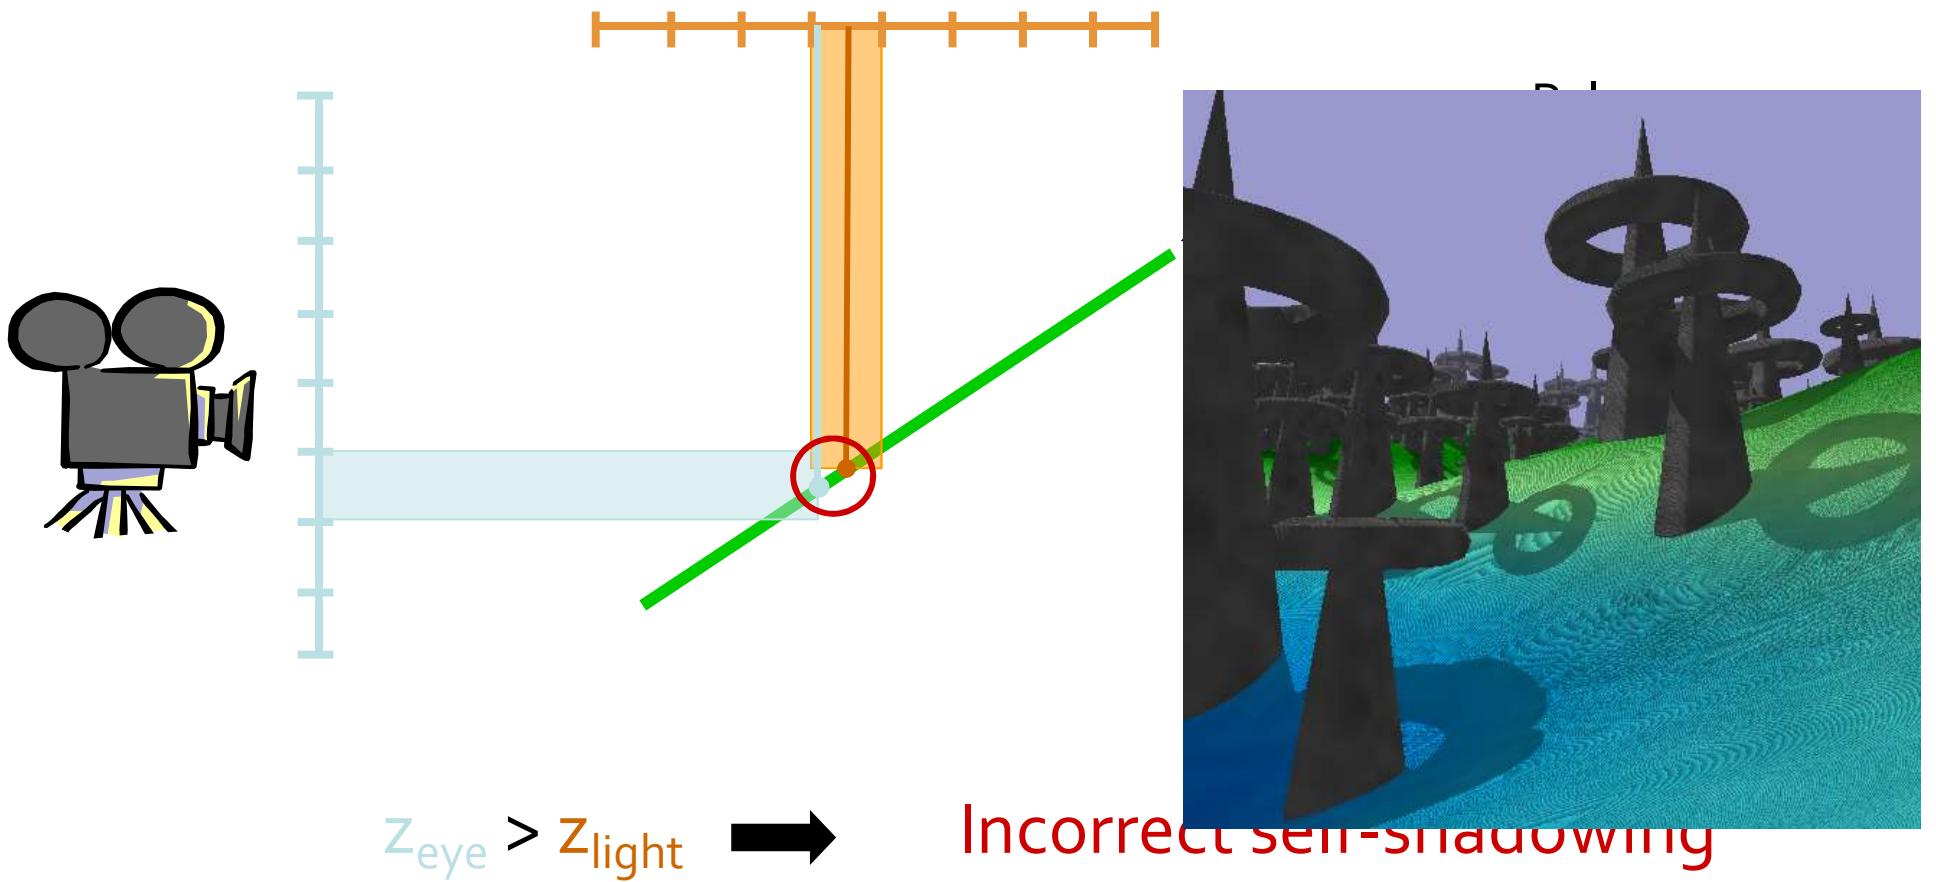
The background features a series of fine, light brown wavy lines that create a sense of depth and motion. These lines are more concentrated in the lower half of the slide, forming a base for the text. In the upper half, they transition into a more sparse, abstract pattern that recedes towards the top right corner.

Depth Biasing – Incorrect Self Shadowing

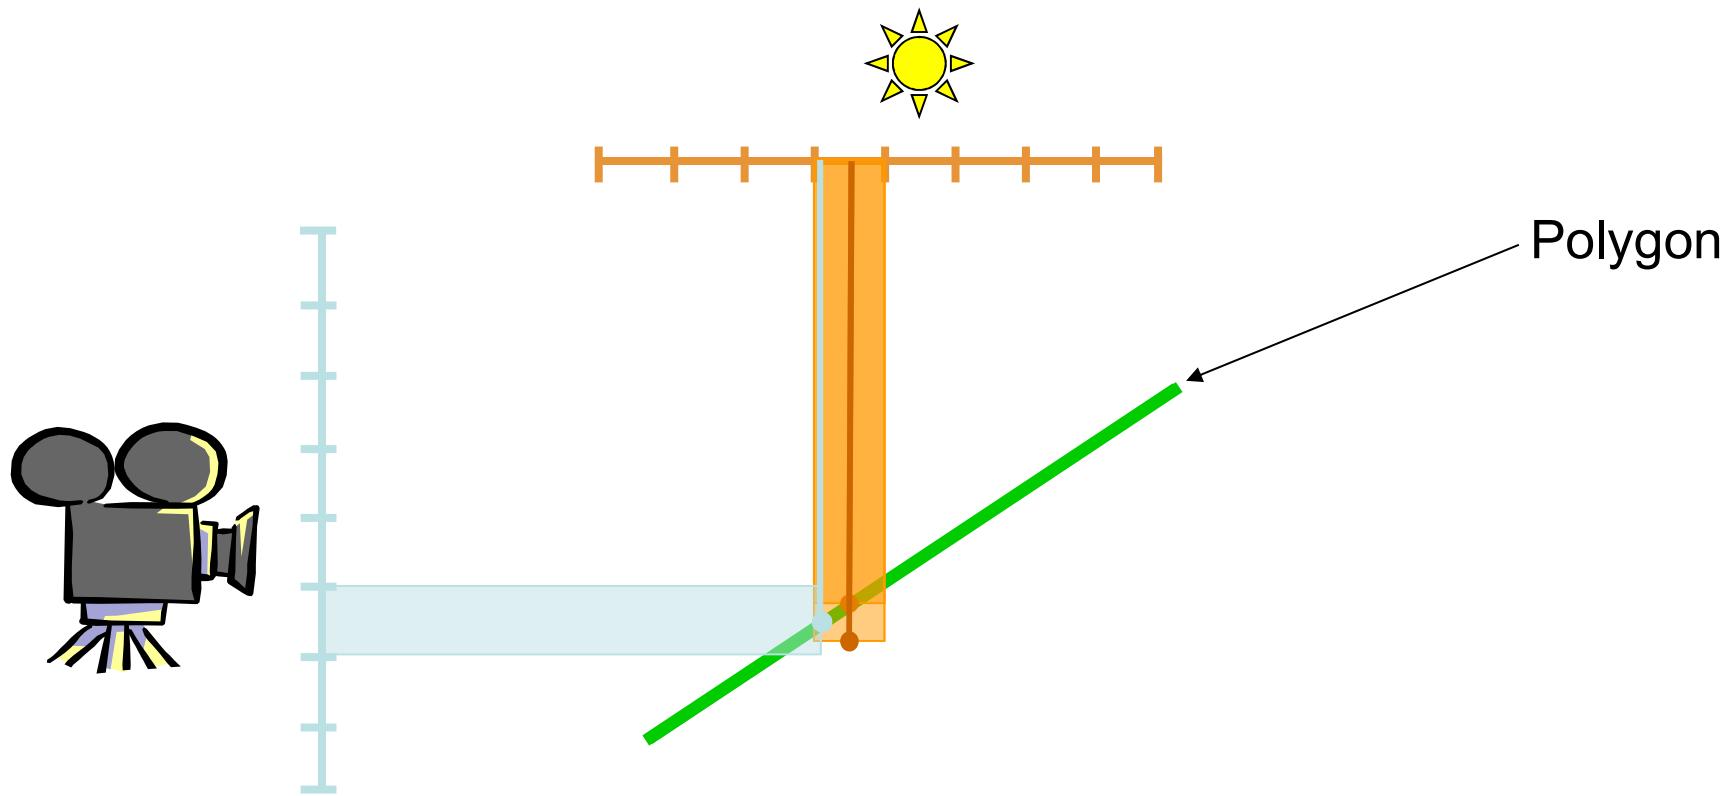


Depth Biasing – Incorrect Self Shadowing

- Numerical instabilities caused by undersampling

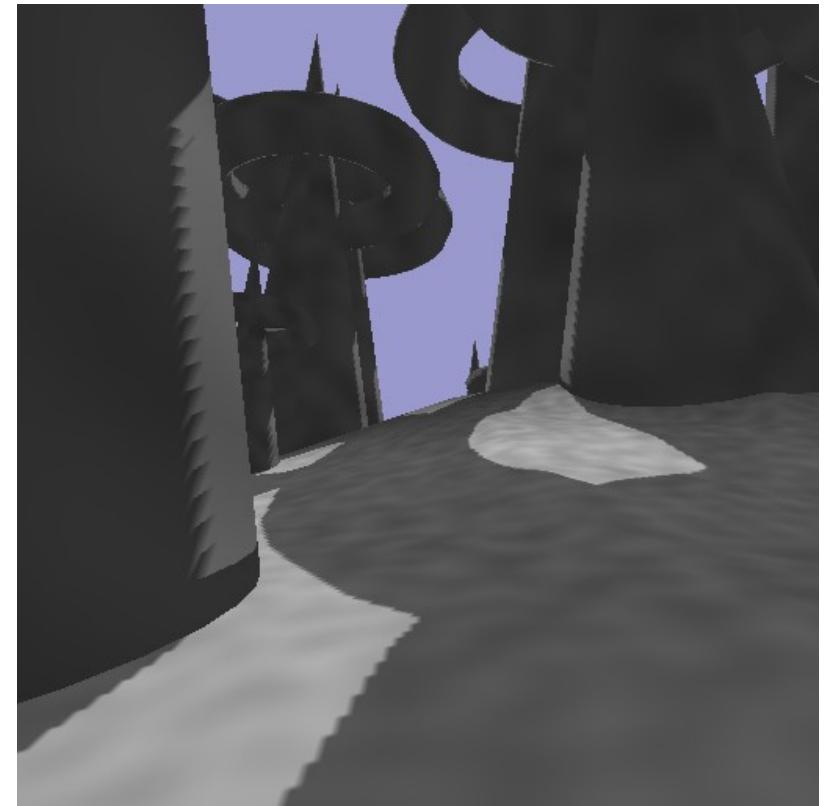
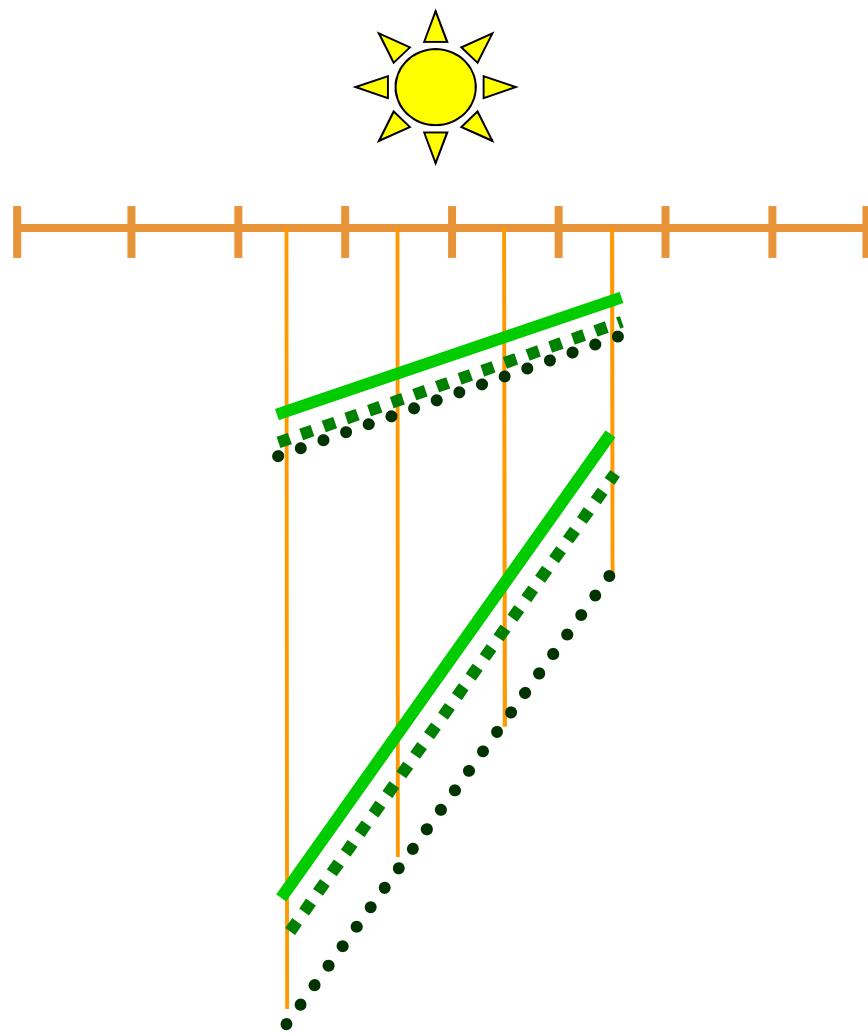


Depth Biasing



$z_{\text{eye}} > z_{\text{light}}$ → Incorrect self-shadowing
 $z_{\text{eye}} < z_{\text{light}}$ → No self-shadowing

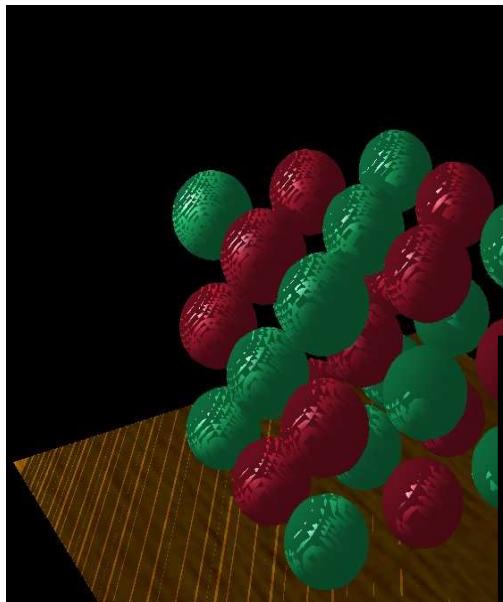
How to Choose the Bias?



- No biasing
- Constant biasing
- Slope-scale biasing

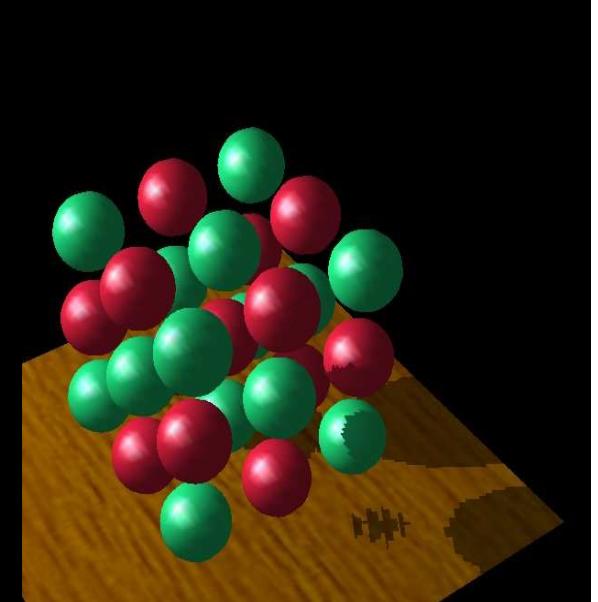
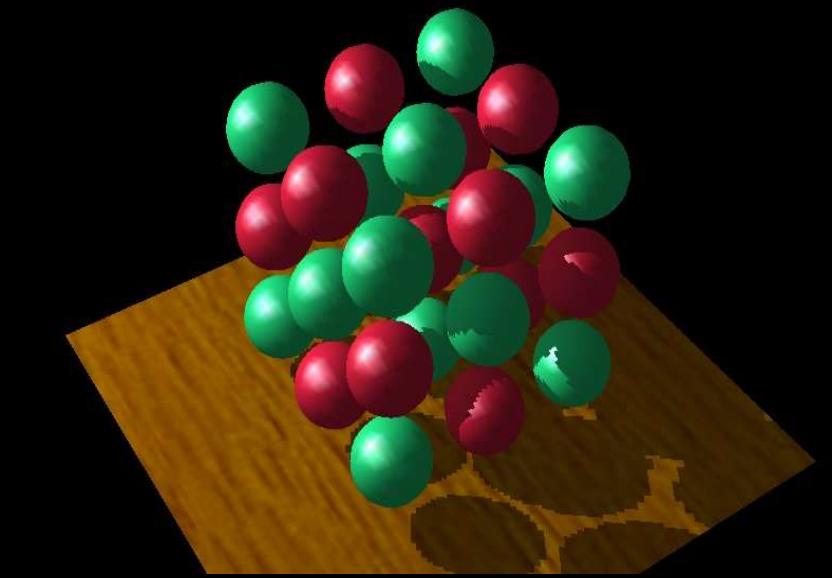
How Much Bias?

- Depends...



Too little bias,
everything begins to
shadow

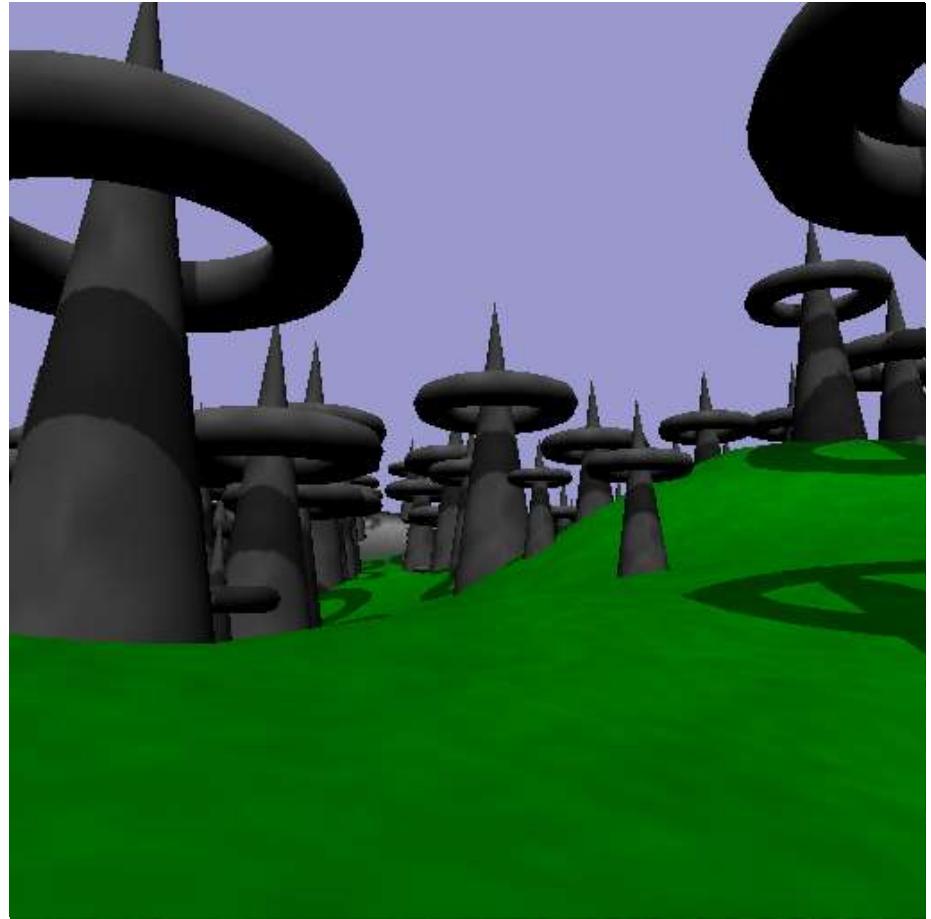
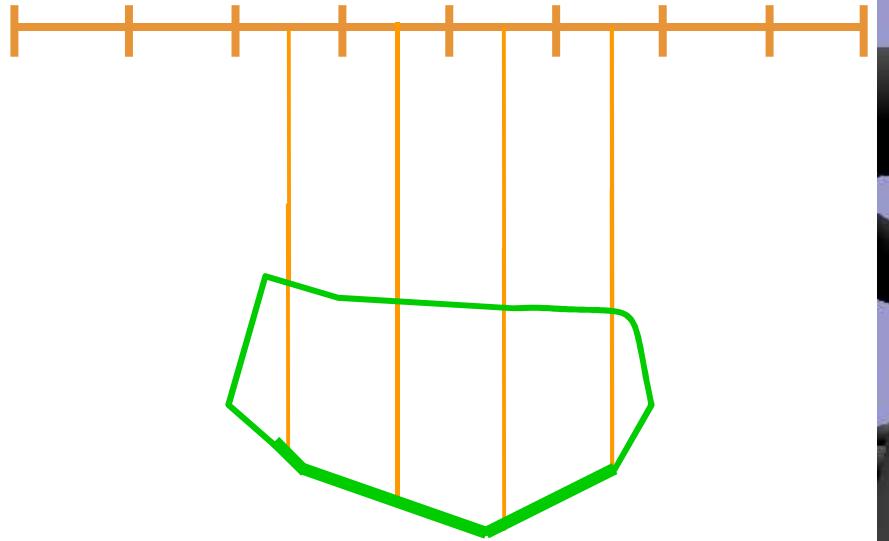
Just right



Too much bias, shadow
starts too far back

Second-Depth Shadow Mapping

[Wang and Molnar 94]



- Normally: render front-side polygons into SM
- Now: render back-side polygons

Error Overview

■ Undersampling

- Improve initial sampling
- Depth Biasing

■ Reconstruction error

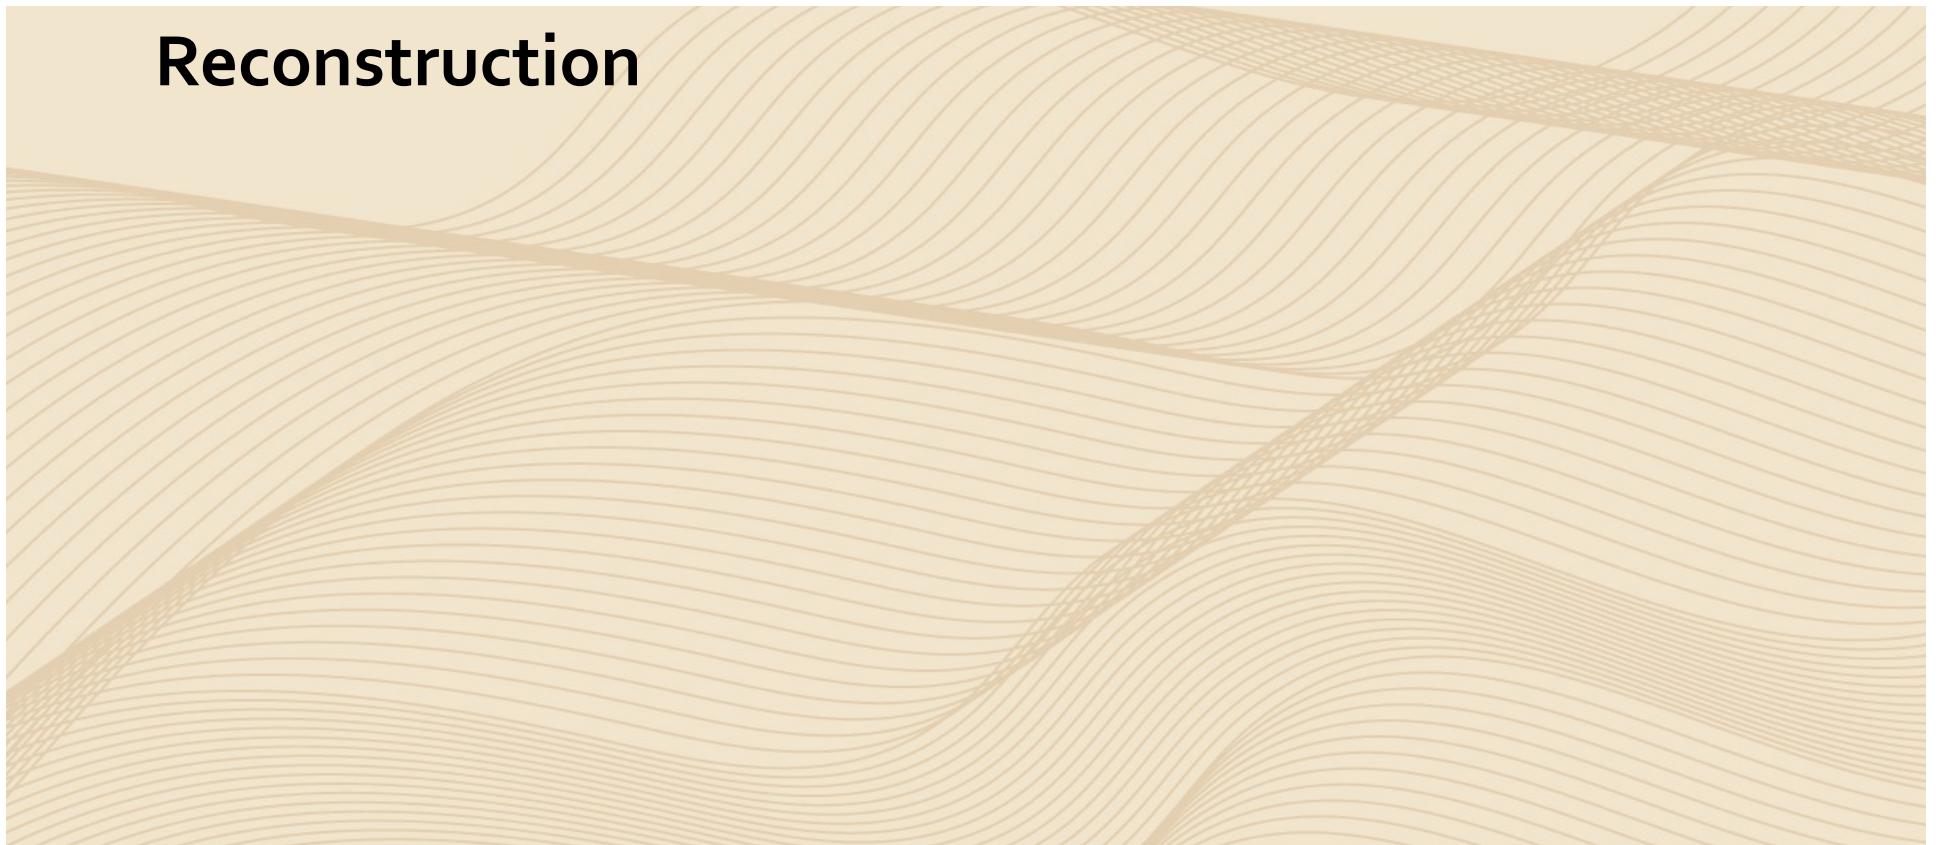
- Use different reconstruction algorithm
(Silhouette Shadow Maps)
- Use better reconstruction filters
(PCF, PCF with Poisson disk sampling)

■ Oversampling

- Use bandlimiting filters
(Convolution Shadow Maps etc., see later)

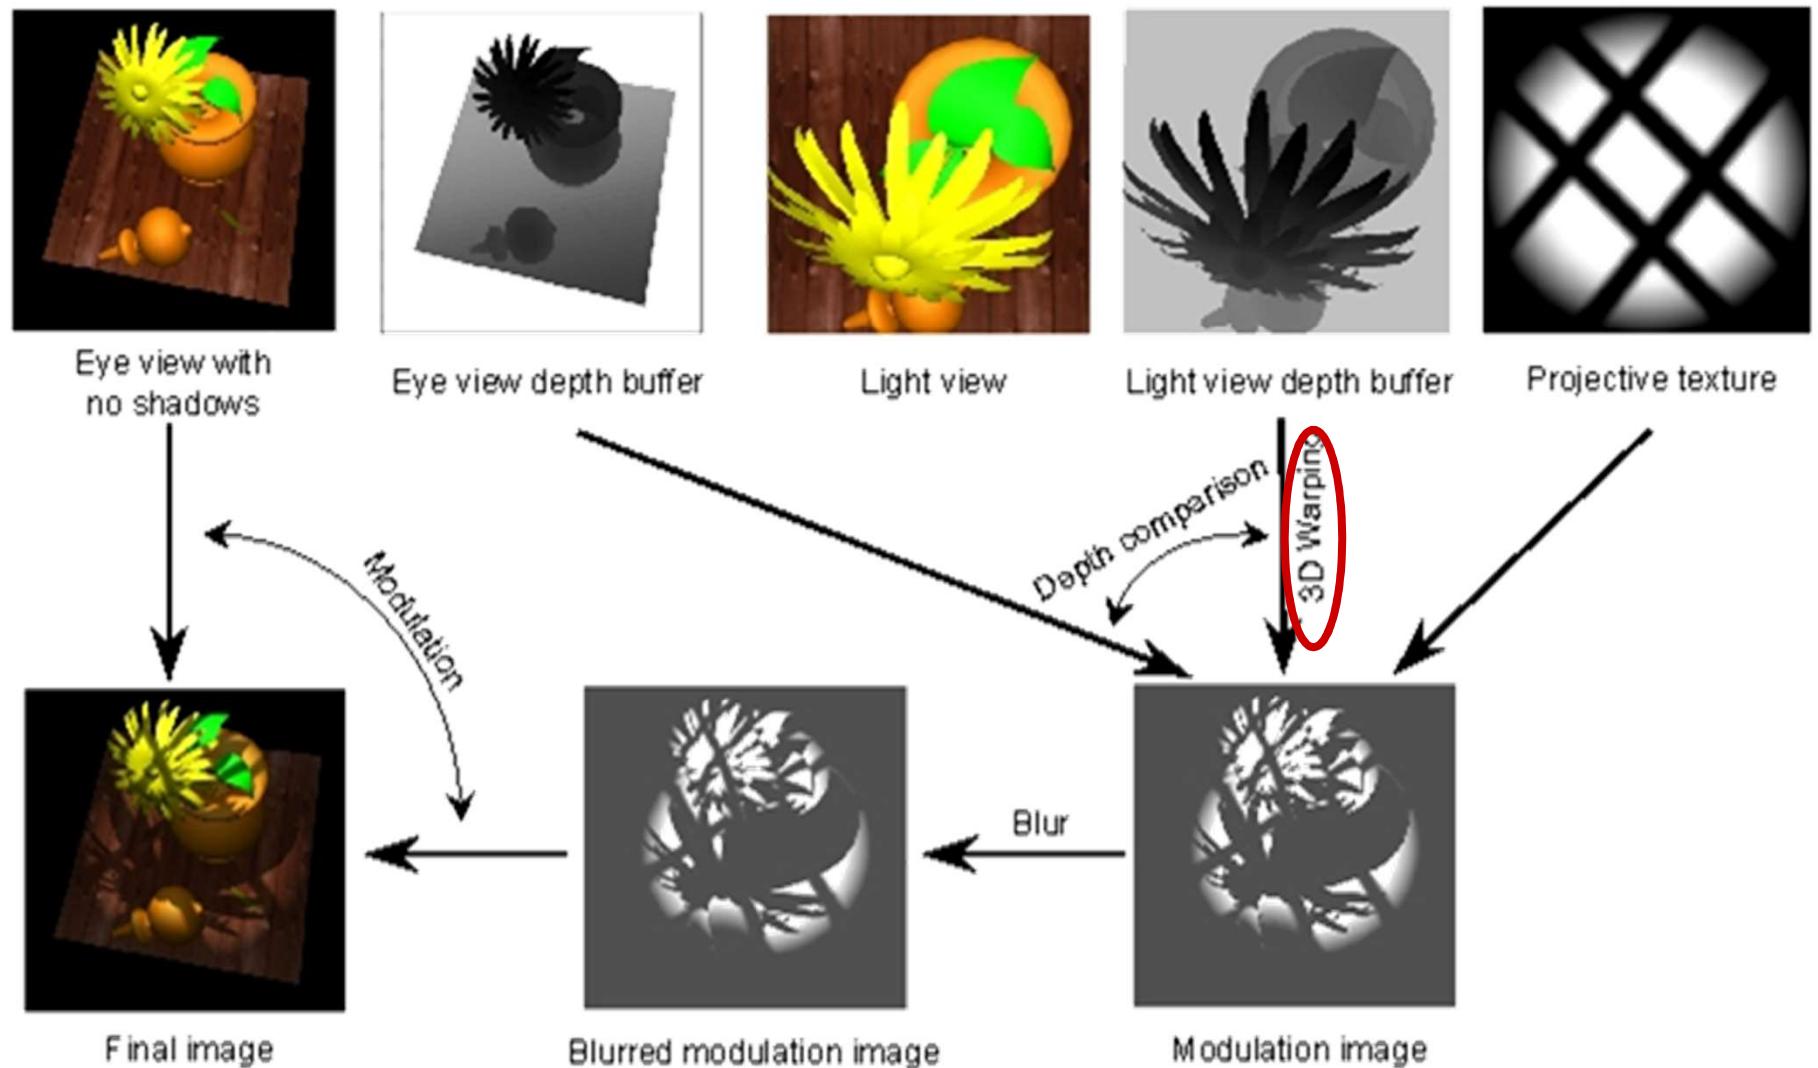
Hard Shadows

Reconstruction



Forward Shadow Mapping

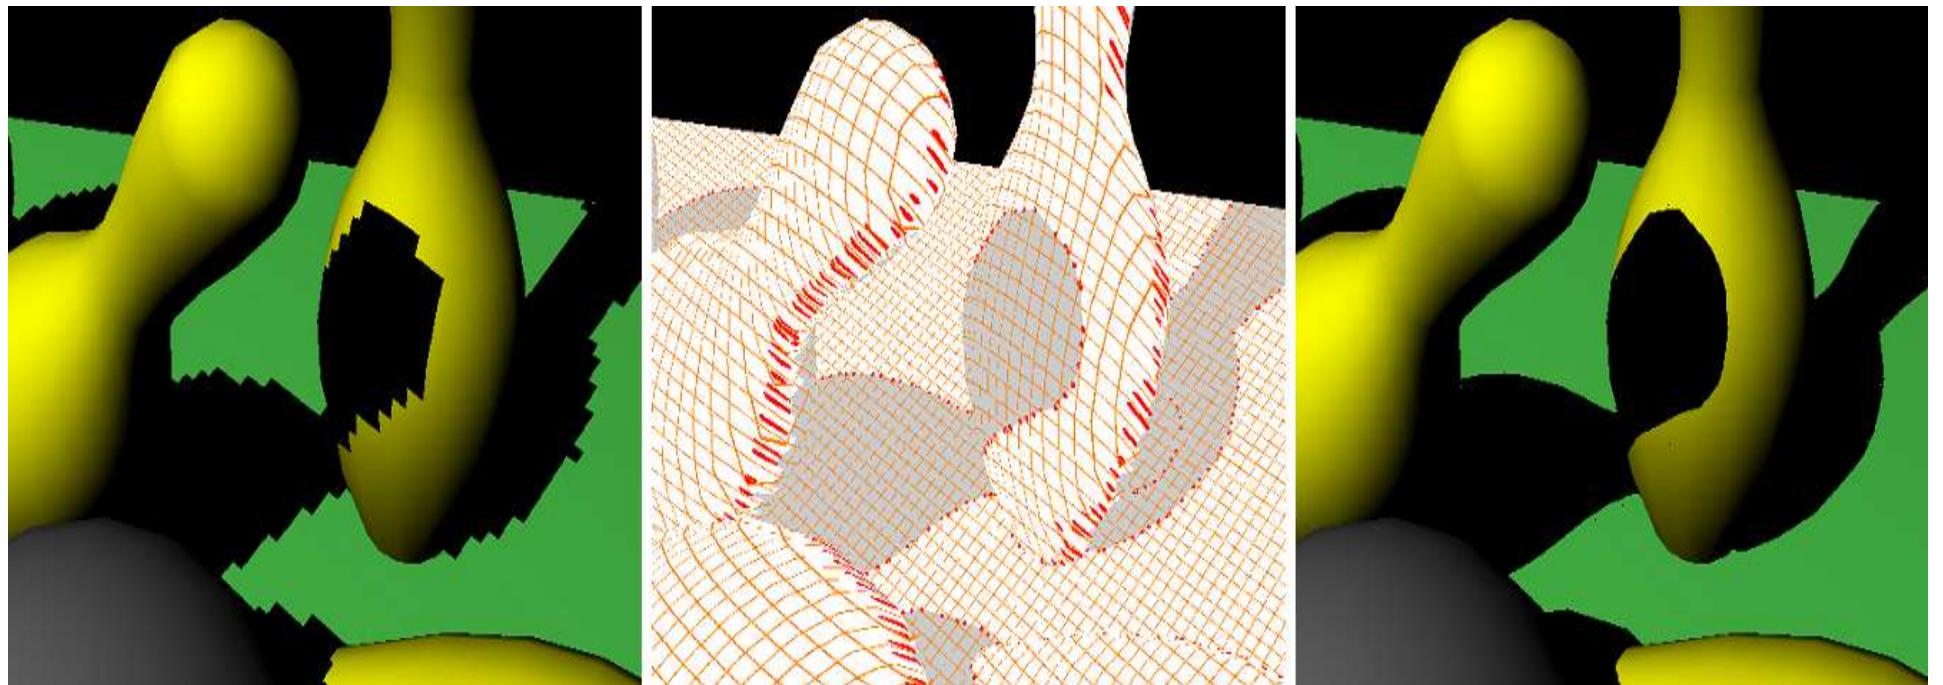
[Zhang 1998]



Shadow Silhouette Maps

[Sen et al. 2003]

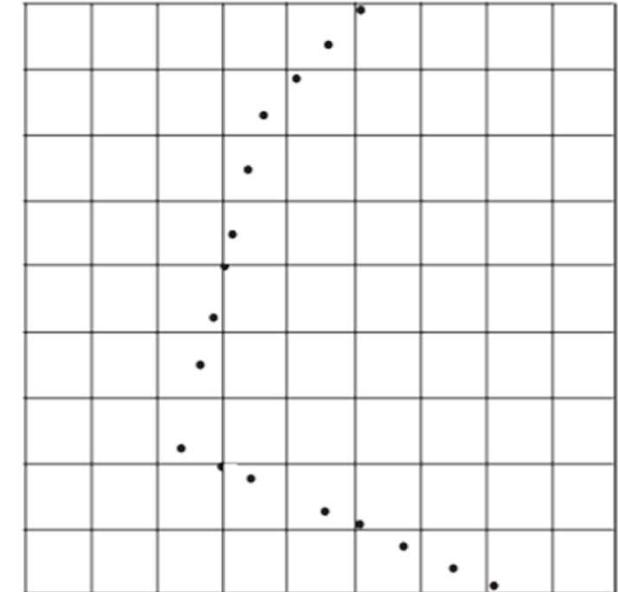
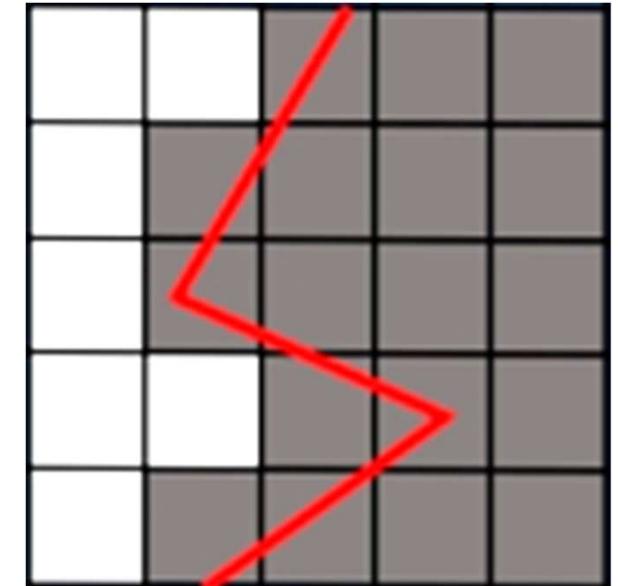
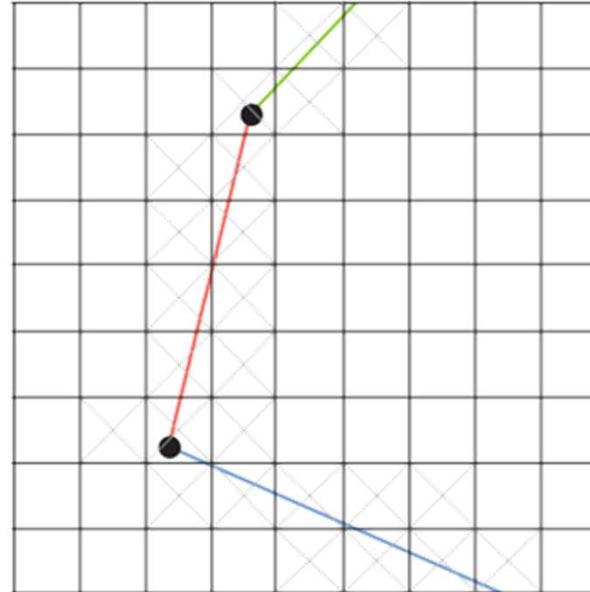
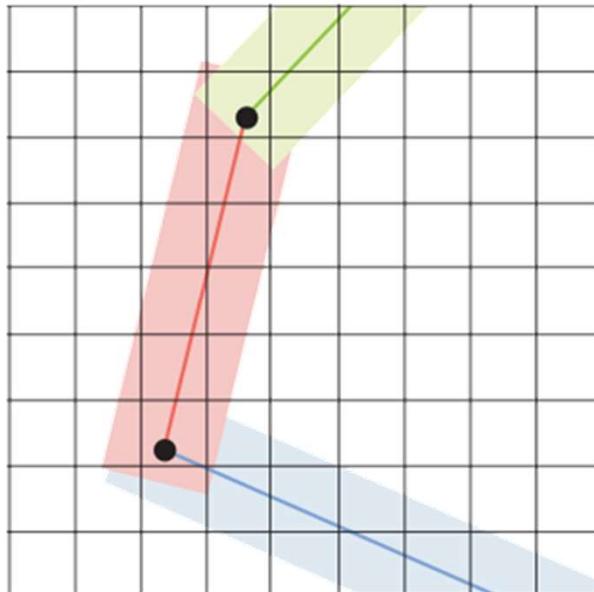
- Use a better silhouette approximation
- Store additional information of shadow edge



Shadow Silhouette Maps

[Sen et al. 2003]

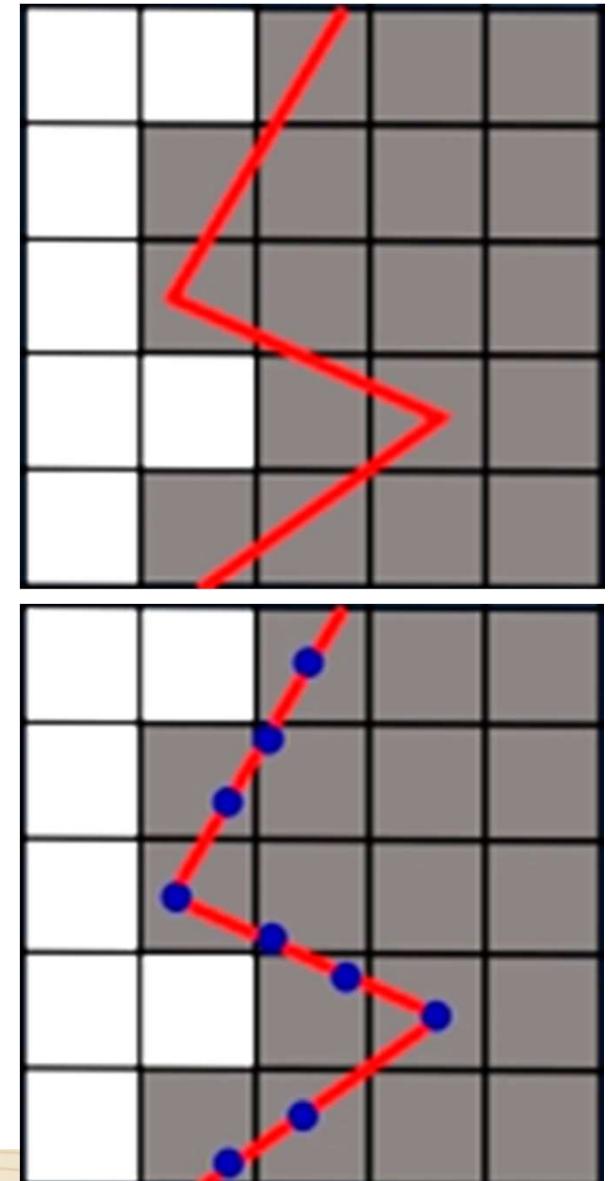
1. Create shadow map
2. Find silhouette edges
3. Rasterize silhouettes
 - a. Find points that lie on silhouette edges
 - b. Store such points into silhouette map



Shadow Silhouette Maps

[Sen et al. 2003]

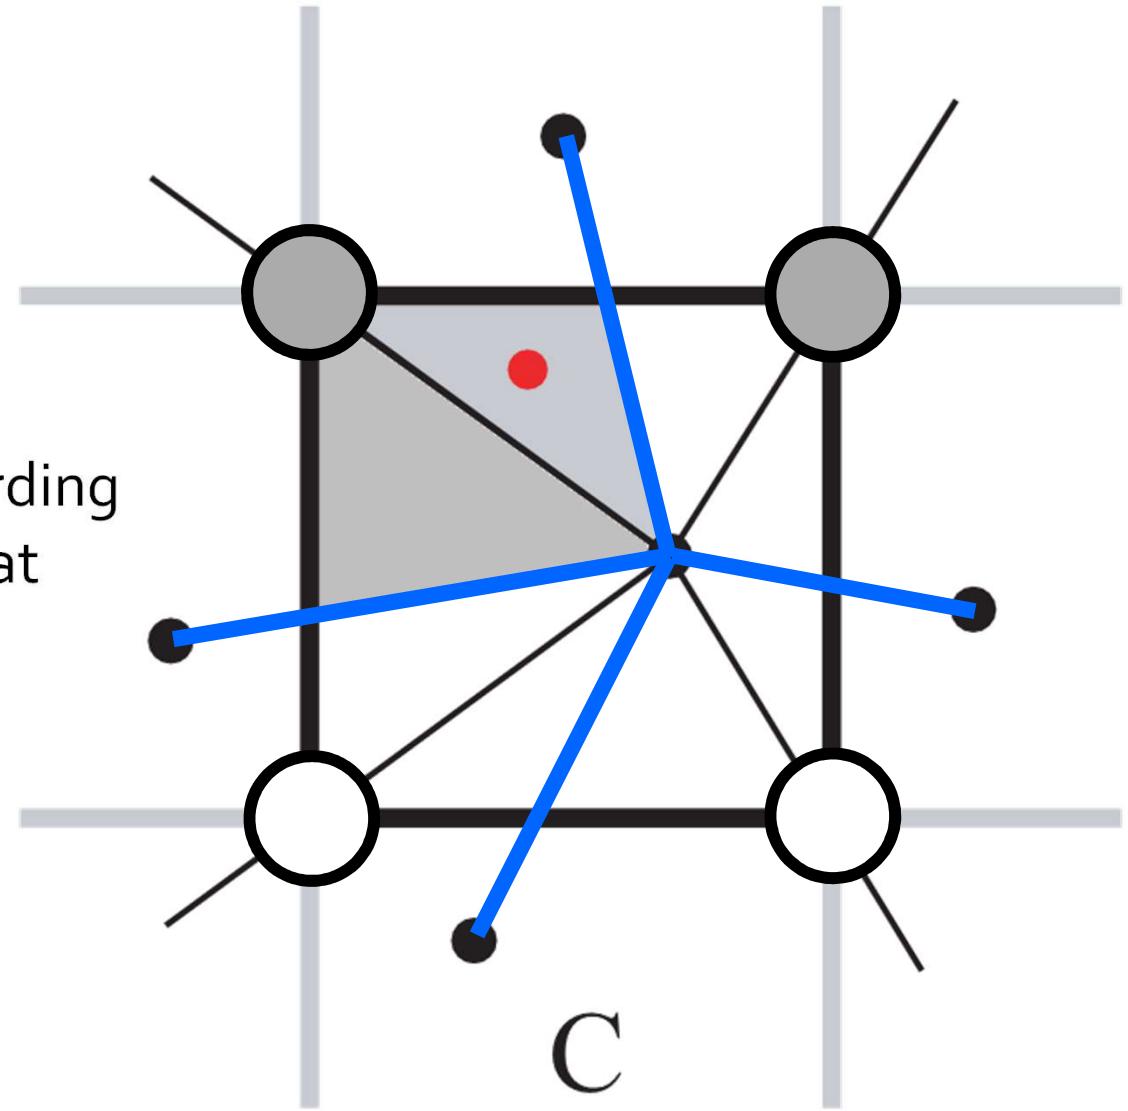
1. Create shadow map
2. Find silhouette edges
3. Rasterize silhouettes
 - a. Find points that lie on silhouette edges
 - b. Store such points into silhouette map
4. Compute shadows
 - a. Non-silhouette pixels use standard shadow map
 - b. Silhouette pixels use silhouette map



Shadow Silhouette Maps

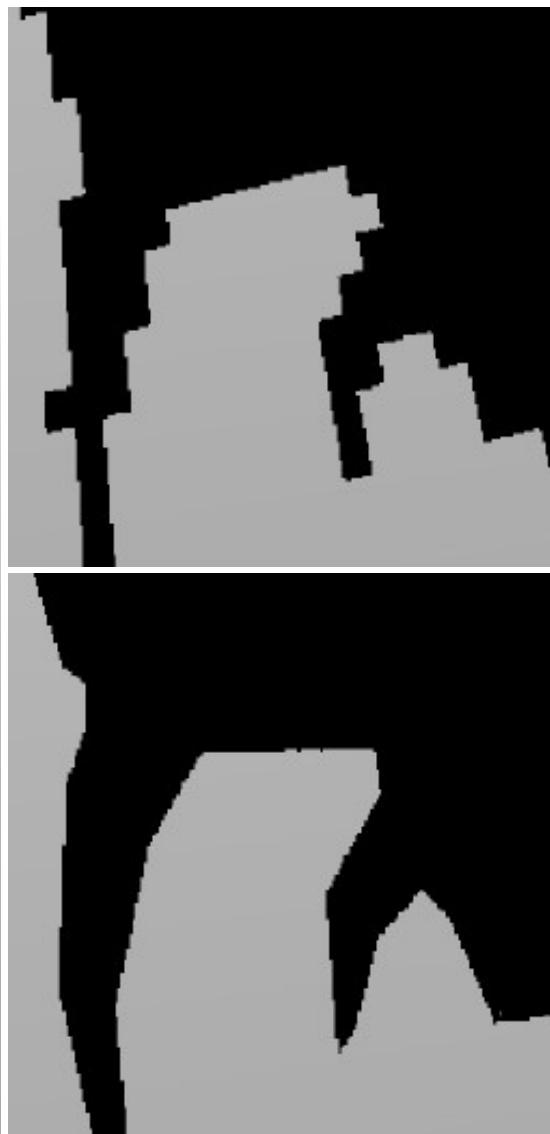
[Sen et al. 2003]

- Fragment
- Silhouette points
 - 1 interior + 4 external
- Create quadrants
- Shade **fragment** according to shadow test result at corner point



Shadow Silhouette Maps

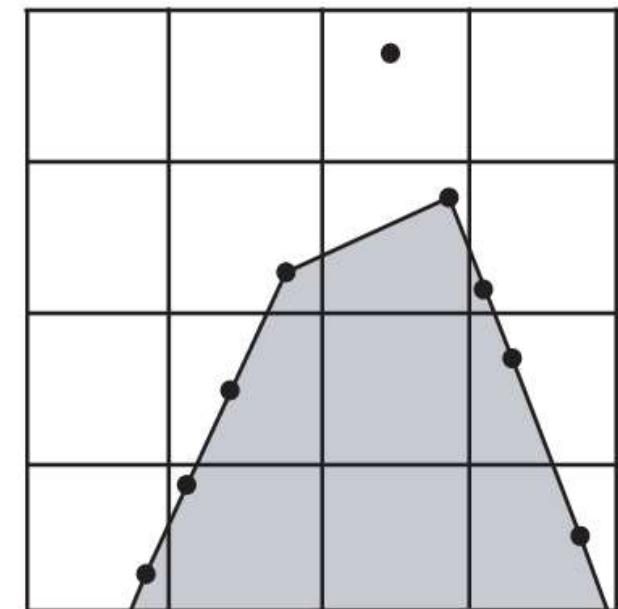
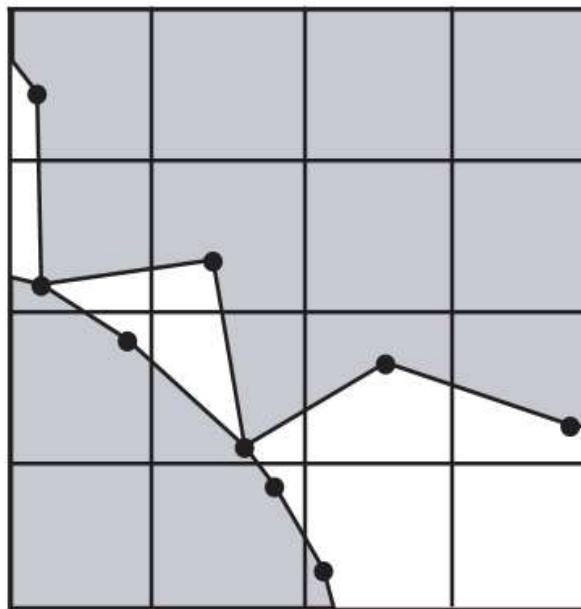
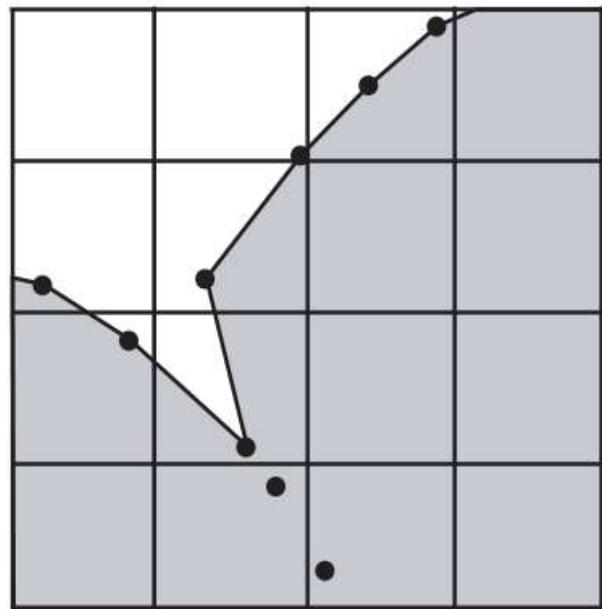
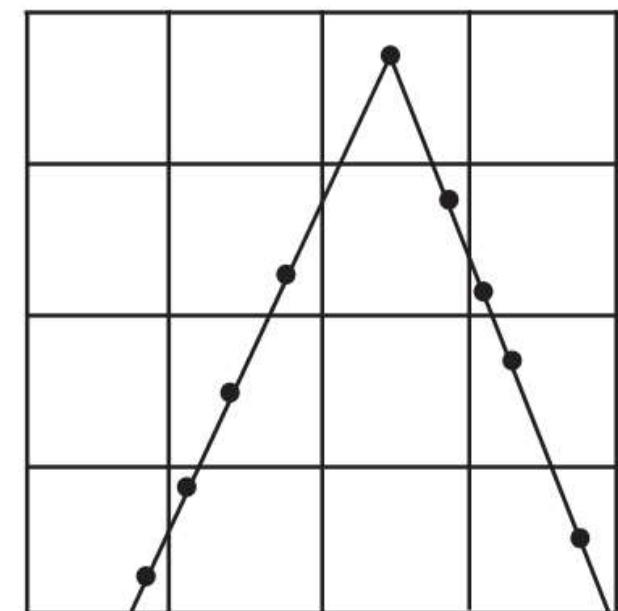
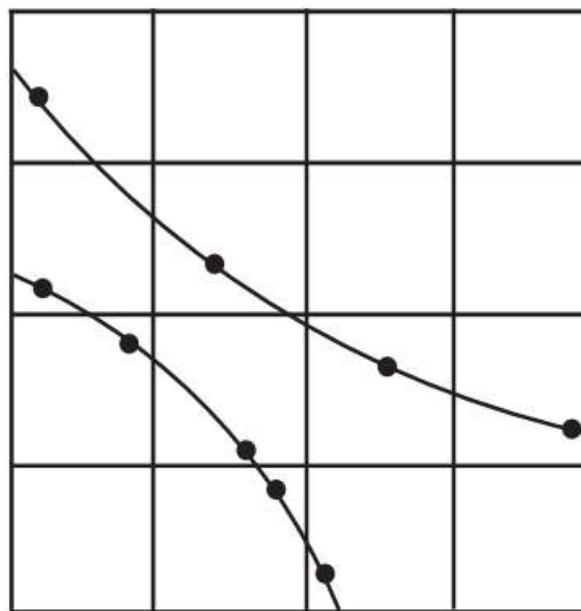
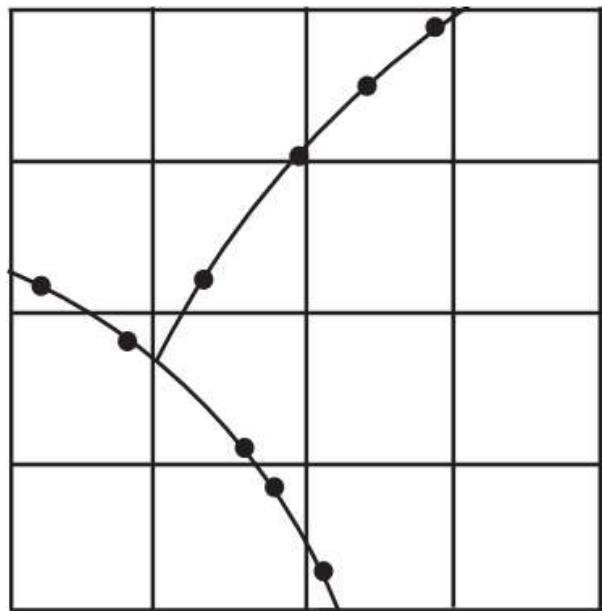
[Sen et al. 2003]

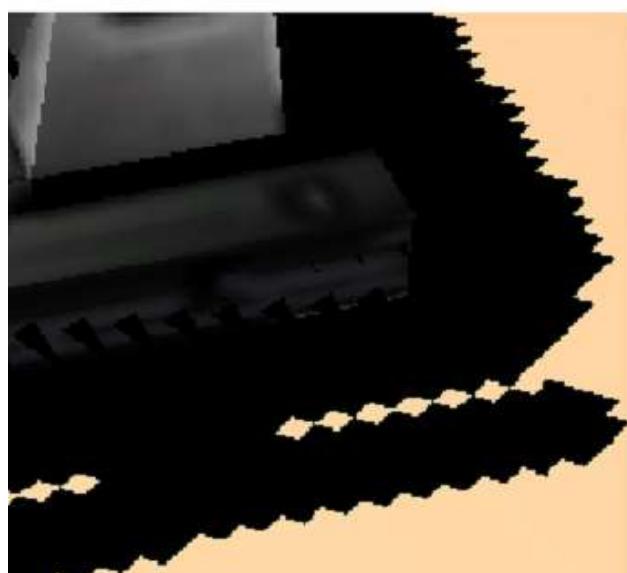
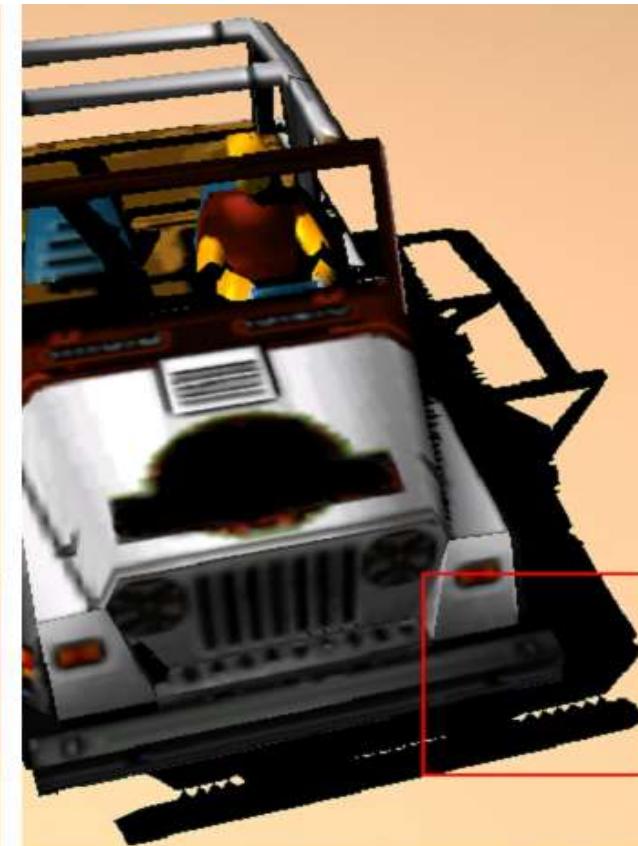


shadow map

silhouette map

Shadow Silhouette Maps

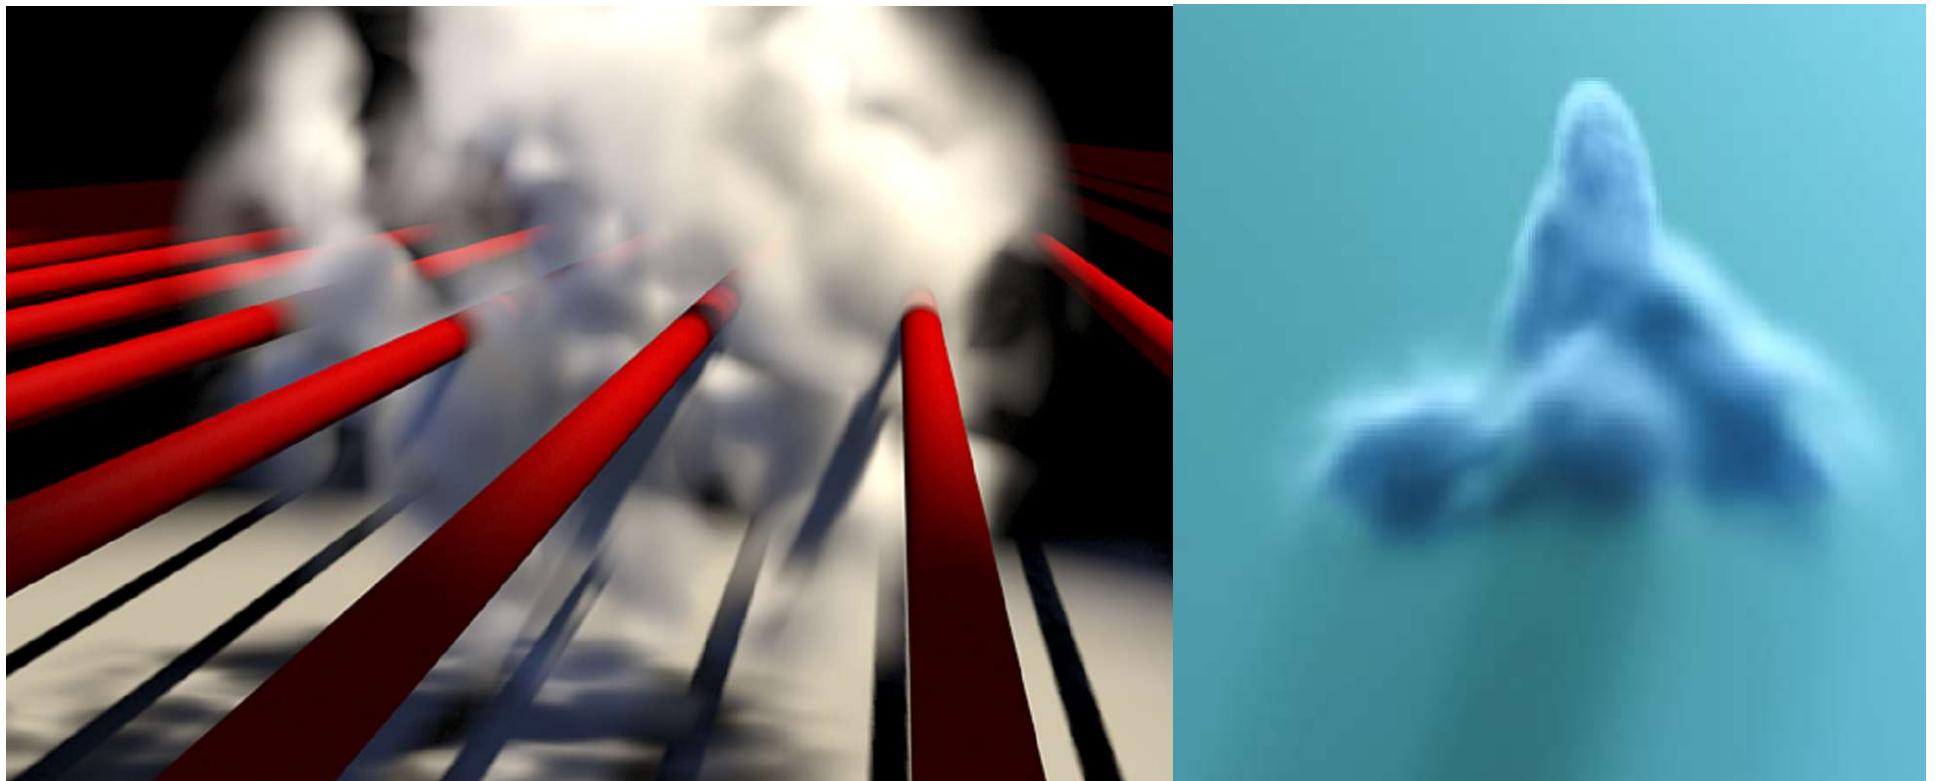




Deep Shadow Maps

[Lokovic and Veach 2000]

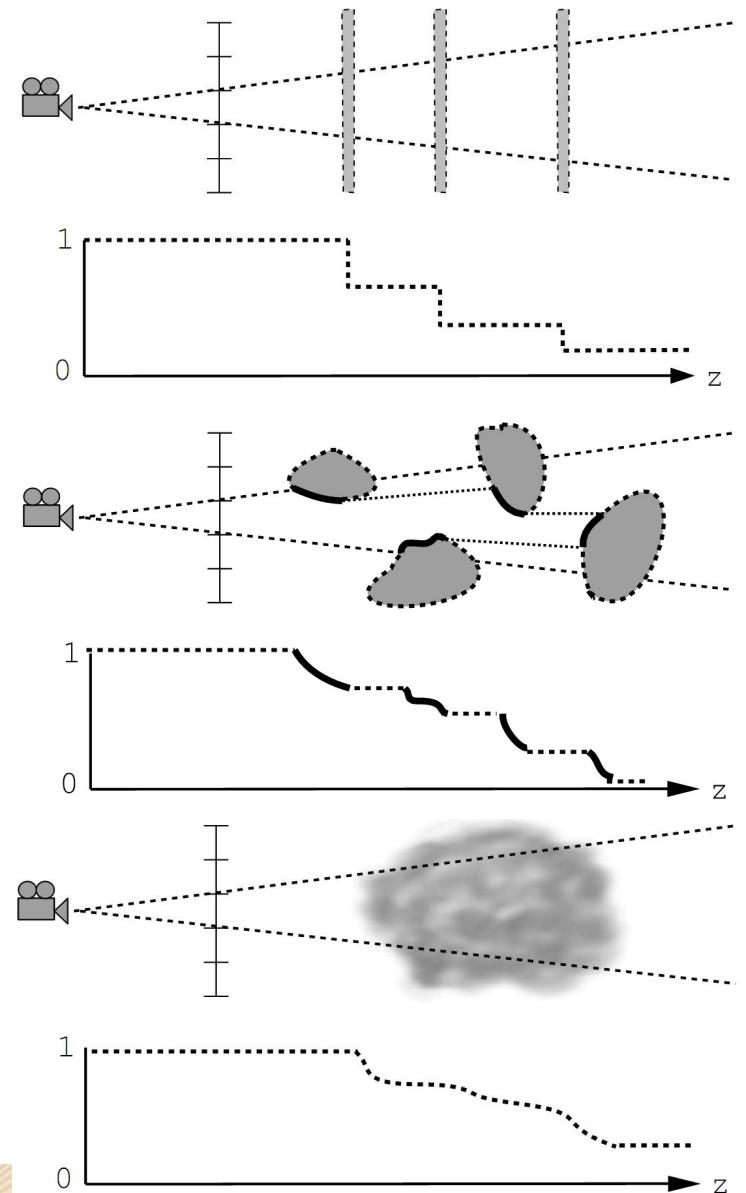
- High quality shadows
- Smoke, hair, fur, ...
- Stores function for each texel



Deep Shadow Maps

[Lokovic and Veach 2000]

- Transmittance function
 - Calculate reduction of light
- Store compressed version
 - Piecewise linear
 - Fixed number of function samples – store in array



Deep Shadow Maps

[Lokovic and Veach 2000]



without self-shadowing



with self-shadowing